

UNIVERSITY OF VERONA

Department of COMPUTER SCIENCE

Master's Degree in
COMPUTER SCIENCE AND ENGINEERING

Master Thesis

**Towards Process Comprehension of Industrial
Control Systems: a Framework for Analyzing
Industrial Systems**

Candidate:

Marco OLIANI
VR457249

Supervisor:

Prof. Massimo MERRO

Co-supervisor:

Prof. Ruggero LANOTTE
University of Insubria

Academic Year 2022/2023

"If you spend more on coffee than on IT security, you will be hacked. What's more, you deserve to be hacked"

(Richard Clarke)

Abstract

Bla bla bla

Contents

1	Introduction	1
1.1	Contribution	1
1.2	Outline	1
2	Background on Industrial Control Systems	3
2.1	Industrial Control Systems Architecture	5
2.2	Operational Technology Networks	7
2.2.1	Field I/O Devices Layer	7
2.2.2	Controller Network Layer	8
2.2.2.1	Programmable Logic Controllers	8
2.2.2.2	Remote Terminal Units	12
2.2.3	Area Control Layer	13
2.2.3.1	Supervisory Control And Data Acquisition	13
2.2.3.2	Human-Machine Interface	13
2.2.4	Operations/Control Layer	14
2.2.5	Demilitarized Zone	14
2.2.6	Industrial Protocols	15
2.2.6.1	Modbus	16
2.2.6.2	EtherNet/IP	19
2.2.6.3	Common Industrial Protocol (CIP)	21
3	State of the Art	23
3.1	Literature on Process Comprehension	24

3.2	Ceccato et al.'s black-box dynamic analysis for water-tank systems	26
3.2.1	Testbed	27
3.2.2	Scanning of the System and Data Pre-processing	31
3.2.3	Graphs and Statistical Analysis	33
3.2.4	Invariants Inference and Analysis	34
3.2.5	Business Process Mining and Analysis	36
3.2.6	Application	38
3.2.7	Limitations	44
4	A Framework to Improve Ceccato et al.'s Work.	55
4.1	The novel Framework	56
4.1.1	Framework Structure	58
4.1.2	Python Libraries and External Tools	59
4.2	Analysis Phases	60
4.2.1	Phase 1: Data Pre-processing	60
4.2.1.1	Subsystem Selection	62
4.2.1.2	Dataset Enrichment	64
4.2.1.3	Datasets Merging	71
4.2.1.4	Brief Analysis of the Obtained Subsystem	72
4.2.2	Phase 2: Graphs and Statistical Analysis	75
4.2.3	Phase 3: Invariant Inference and Analysis	80
4.2.3.1	Revised Daikon Output	81
4.2.3.2	Types of Analysis	85
4.2.4	Phase 4: Business Process Analysis	90
5	Case study: the iTrust SWaT System	91
5.1	Architecture	91
5.1.1	Physical Process	92
5.1.2	Control and Communication Network	94
5.2	Datasets	96
5.2.1	Our Case Study: the 2015 Dataset	97

6 Our framework at work: reverse engineering of the SWaT system	101
6.1 Pre-processing	101
6.2 Graph Analysis	101
6.2.1 Conjectures About the System	101
6.3 Invariants Analysis	101
6.3.1 Actuators Detection	101
6.3.2 Daikon Analysis and Results Comparing	101
6.4 Extra information on the Physics	101
6.5 Business Process Analysis	101
7 Conclusions	103
7.1 Discussions	103
7.2 Guidelines	103
7.3 Future work	103
List of Figures	105
List of Tables	107
Listings	109
References	111

Introduction

LOREM ipsum dolor bla bla bla. Ma dove metto l'abstract? Prova di interlinea che direi posso anche andare bene, ma bisogna poi vedere il tutto come si incasca alla fine, in modo da ottenere un bel risultato alla vista.

1.1 Contribution

Lorem Ipsum is simply dummy text of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s, when an unknown printer took a galley of type and scrambled it to make a type specimen book. It has survived not only five centuries, but also the leap into electronic typesetting, remaining essentially unchanged. It was popularised in the 1960s with the release of Letraset sheets containing Lorem Ipsum passages, and more recently with desktop publishing software like Aldus PageMaker including versions of Lorem Ipsum.

1.2 Outline

The thesis is structured as follows:

Chapter 2: provides background on the topics covered in this thesis: Industrial Control Systems (ICSs), Supervisory Control And Data Acquisition (SCADA), Programmable Logic Controllers (PLCs) and other devices, industrial communication protocols.

Chapter 3:

Chapter 4:

Chapter 5:

Chapter 6:

Chapter 7:

Background on Industrial Control Systems

¹ INDUSTRIAL CONTROL SYSTEMS (ICSs) are information systems used to
² control industrial processes such as manufacturing, product handling,
³ production, and distribution [1]. ICSs are often found in critical infrastruc-
⁴ ture facilities such as power plants, oil and gas refineries, and chemical
⁵ plant ICSs are different from traditional IT systems in several key ways.

⁶ Firstly, ICSs are designed to control physical processes, whereas IT sys-
⁷ tems are designed to process and store data. This means that ICSs have dif-
⁸ ferent requirements for availability, reliability, and performance. Secondly,
⁹ ICSs are typically deployed in environments that are harsh and have lim-
¹⁰ ited resources, such as extreme temperatures and limited power. Thirdly,
¹¹ the protocols hardware and software used in ICSs are often proprietary.

¹² ICSs are becoming increasingly connected to the internet and other net-
¹³ works, which has led to increased concerns about their security. Industrial
¹⁴ systems were not originally designed with security in mind, and many of
¹⁵ them have known vulnerabilities that could be exploited by attackers. Ad-
¹⁶ ditionally, the use of legacy systems and equipment can make it difficult
¹⁷ to implement security measures. As a result, ICSs are increasingly seen
¹⁸ as a potential target for cyber attacks, which could have serious conse-
¹⁹ quences for the safe and reliable operation of critical infrastructure: some
²⁰ notorious examples of cyber attacks are (*i*) the **STUXnet** worm [2], which

²¹ purpose was to sabotage the nuclear centrifuges of the enrichment plant
²² at the Natanz nuclear facility in Iran; (ii) **Industroyer** [3], also referred
²³ as *Crashoverride*, responsible for the attack on the Ukrainian power grid
²⁴ on December 17, 2016; (iii) the attack on February, 2021 to a water treat-
²⁵ ment plant in Oldsmar, Florida [4], where the level of sodium hydroxide
²⁶ was intentionally increased to a level approximately 100 times higher than
²⁷ normal.

²⁸

²⁹ The increasing connectivity of ICSs and the associated security risks have
³⁰ led to a growing interest in the field of ICS security. Researchers and prac-
³¹ titioners are working to develop new security technologies, standards, and
³² best practices to protect ICSs from cyber attacks. This includes efforts to
³³ improve the security of ICS networks and devices, as well as the develop-
³⁴ ment of new monitoring and detection techniques to identify and respond
³⁵ to cyber attacks.

³⁶

³⁷ Table 2.1 summarizes the differences between traditional IT and ICSs [5]:

	Traditional IT	ICSs
Focus	Data	Asset
Update Frequency	High	Low
Priority	Confidentiality Integrity Availability	Availability Integrity Confidentiality
Operating System	Standardized	Proprietary
Protocols	Standardized	Proprietary
Attacker Motivation	Monetization	Disruption

Table 2.1: differences between Information Technology (IT) and Industrial Control Systems (ICSs)

³⁸ 2.1 Industrial Control Systems Architecture

³⁹ In the past, there has been a clear division between *Information Technology* (IT) and *Operational Technology* (OT), both at the technical and organizational levels. Each domain has maintained its own distinct technology stacks, protocols, and standards. However, with the emergence of Industry 4.0 and the rapid expansion of industrial automation, which heavily relies on IT tools for monitoring and controlling critical infrastructures, the boundary between IT and OT has started to blur. This trend has paved the way for greater integration between these two domains, thus improving productivity and process quality.

⁴⁸

⁴⁹ General ICS architecture consists in **six levels** each representing a functionality: this architecture comprising the OT and IT parts is represented in Figure 2.1 [6][5], according to the *Purdue Enterprise Reference Architecture* (PERA), or simply **Purdue Model**:

- ⁵³ • Level 0 (**Processes, or Field I/O Devices**): contains **field devices**.
- ⁵⁴ • Level 1 (**Intelligent Devices, or Controller Network**): includes **local or remote controllers** that sense, monitor and control the physical process, such as **PLCs** (2.2.2.1) and **RTUs** (2.2.2.2). Controllers interface directly to the field devices reading data from sensors and sending commands to actuators.
- ⁵⁹ • Level 2 (**Control Systems, or Area Control**): contains computer systems used to supervising and monitoring the physical process: they provide a **Human-Machine Interface** (HMI, 2.2.3.2) and *Engineering Workstations* (EW) for operator control.
- ⁶³ • Level 3 (**Manufacturing/Site Operations, or Operations/Control**): comprises systems used to manage the production workflow for plant-wide control: they collate informations from the previous levels and store them in Data Historian servers.

2.1 Industrial Control Systems Architecture

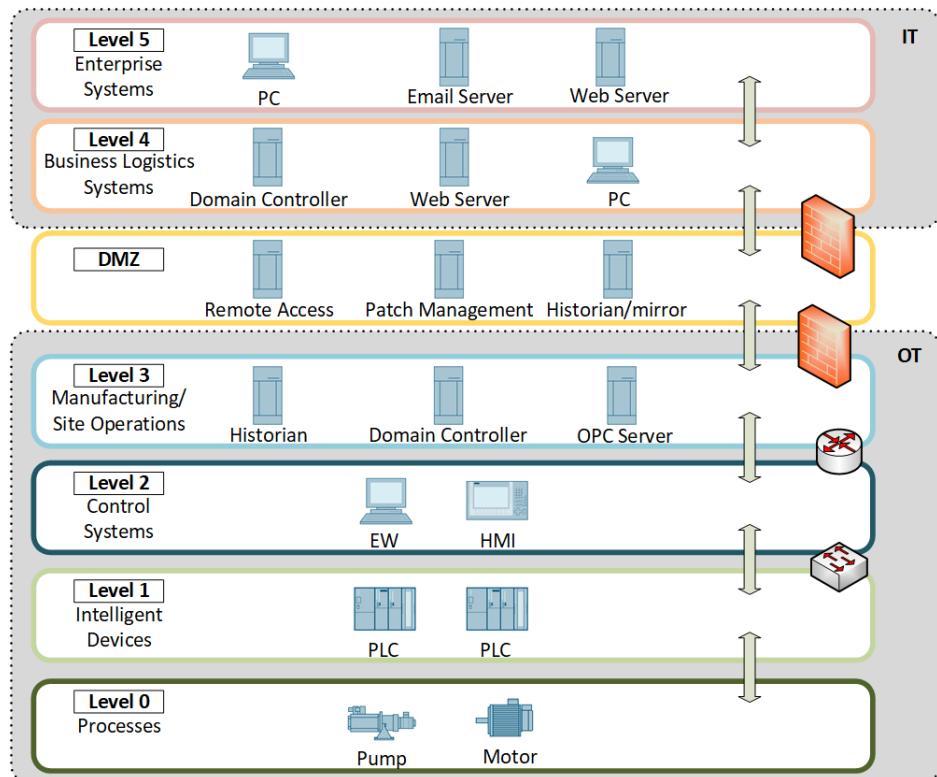


Figure 2.1: ICS architecture schema

- **Industrial Demilitarized Zone (DMZ):** intermediate level that connects the *Operational Technology* (OT) part (levels 0-3) with the *Information Technology* (IT) part of the system (levels 4 and 5). Communication takes place indirectly through services such as *proxy servers* and *remote access servers*, which act as intermediaries between the two environments.
- Level 4 (**Business Logistics Systems**, or **Business Planning/Logistics**): collect and aggregates data from the Manufacturing/Site Operations level overseeing the IT-related activities to generate **reporting** to the Enterprise System layer. At this layer we can find application and e-mail servers, and *Enterprise Resource Planning* (ERP) systems.
- Level 5 (**Enterprise Systems**): represents the enterprise network, used for the business-to-business activities and for business-to-client pur-

pose services. At Enterprise Systems level are typical IT services such as mail servers, web servers and all the systems used to manage the ongoing process.

As previously discussed, the gap between IT and OT is steadily narrowing. Nowadays, it is increasingly common to encounter IT elements within the OT realm. For example, desktop PCs are now frequently found in OT environments, and industrial devices are interconnected using standard IT communication protocols like TCP and UDP.

2.2 Operational Technology Networks

Operational Technology primarily encompasses the **tangible aspects** of Industrial Control Systems and directly interfaces with the physical processes of the monitored systems. Its main purpose is to **manage and control the procedures** involved in creating and correcting physical value in various equipment.

This section will focus on the key aspects and components of Operational Technology network, with specific reference to the first four levels of the Purdue model previously seen.

2.2.1 Field I/O Devices Layer

This level concerns all aspects related to the physical environment and the physical elements that are part of it, which have the ability to actively influence the environment.

These physical elements are represented by **Field Devices**, i.e., **sensors** and **actuators** used to collect data from the process and control it: sensors are the elements responsible for reading specific values related to the physical environment (e.g., the level of a liquid), while actuators change its behavior and characteristics (e.g., opening or closing a valve to make the liquid flow).

Examples of field devices include temperature sensors, pressure sensors, valves and pumps.

¹⁰⁹ 2.2.2 Controller Network Layer

¹¹⁰ *Controller Network* layer includes devices that handle data from and to
¹¹¹ the *Field I/O Devices* layer. This kind of device is capable of gathering data
¹¹² from sensors, updating its internal state, and activating actuators (for ex-
¹¹³ ample opening or close a pump that controls the level of a tank), making
¹¹⁴ decisions based on a customized program, known as its control logic.
¹¹⁵ Commonly found within this layer are *Programmable Logic Controllers* (PLCs)
¹¹⁶ and *Remote Terminal Units* (RTUs): in the upcoming sections, we will ex-
¹¹⁷ amine these elements in detail.

¹¹⁸ 2.2.2.1 Programmable Logic Controllers

¹¹⁹ A *Programmable Logic Controller* (PLC) is a **small and specialized in-**
¹²⁰ **dustrial computer** having the capability of controlling complex industrial
¹²¹ and manufacturing processes [7].

¹²² Compared to relay systems and personal computers, PLCs are opti-
¹²³ mized for control tasks and industrial environments: they are rugged and
¹²⁴ designed to withdraw harsh conditions such as dust, vibrations, humid-
¹²⁵ ity and temperature: they have more reliability than personal computers,
¹²⁶ which are more prone to crash, and they are more compact a require less
¹²⁷ maintenance than a relay system.

¹²⁸ Furthermore, I/O interfaces are already on the controller, so PLCs are eas-
¹²⁹ ier to expand with additional I/O modules (if in a rack format) to manage
¹³⁰ more inputs and ouputs, without reconfiguring hardware as in relay sys-
¹³¹ tems when a reconfiguration occurs.

¹³² PLCs are more *user-friendly*: they are not intended (only) for computer
¹³³ programmers, but designed for engineers with a limited knowledge in
¹³⁴ programming languages: control program can be entered with a simple
¹³⁵ and intuitive language based on logic and switching operations instead of
¹³⁶ a general-purpose programming language (*i.e.* C, C++, ...).

137 **PLC Architecture** The basic hardware architecture of a PLC consists of
138 these elements [8]:

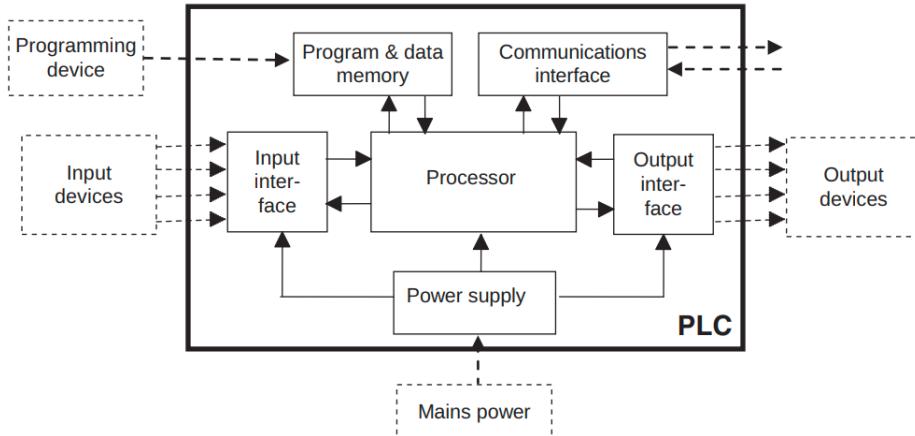


Figure 2.2: PLC architecture

- 139 • **Processor unit (CPU):** contains the microprocessor. This unit interprets the input signals from I/O modules, executes the control program stored in the Memory Unit and sends the output signals to the I/O Modules. The processor unit also sends data to the Communication interface, for the communication with additional devices.
- 140 • **Power supply unit:** converts AC voltage to low DC voltage.
- 141 • **Programming device:** is used to store the required program into the memory unit.
- 142 • **Memory Unit:** consists in RAM memory and ROM memory. RAM memory is used for storing data from inputs, ROM memory for storing operating system, firmware and user program to be executed by the CPU.
- 143 • **I/O modules:** provide interface between sensors and final control elements (actuators).
- 144 • **Communications interface:** used to send and receive data on a network from/to other PLCs.

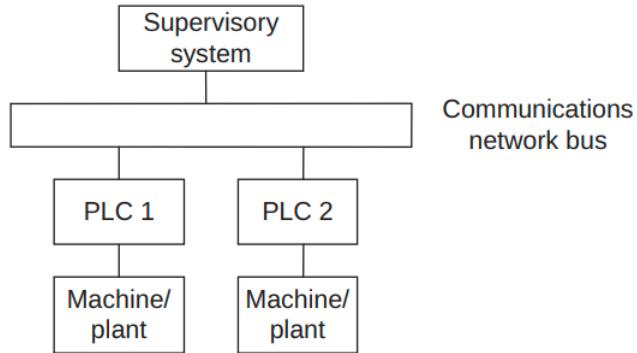


Figure 2.3: PLC communication schema

₁₅₅ **PLC Programming** Two different programs are executed in a PLC: the ₁₅₆ **operating system** and the **user program**.

₁₅₇ The operating system tasks include executing the user program, managing memory areas and the *process image table* (memory registers where ₁₅₉ inputs from sensors and outputs for actuators are stored).

₁₆₀ The user program needs to be uploaded on the PLC via the programming device and runs on the process image table in *scan cycles*: each scan ₁₆₂ is made up of three phases [9]:

- ₁₆₃ 1. reading inputs from the process images table
- ₁₆₄ 2. execution of the control code and computing the physical process evolution
- ₁₆₆ 3. writing output to the process image table to have an effect on the physical process. At the end of the cycle, the process image table is ₁₆₈ refreshed by the CPU

₁₆₉ Standard PLCs **programming languages** are basically of two types: ₁₇₀ **textuals** and **graphicals**. Textual languages include languages such as ₁₇₁ *Instruction List* (IL) and *Structured Text* (ST), while *Ladder Diagrams* (LD), ₁₇₂ *Function Block Diagram* (FBD) and *Sequential Function Chart* (SFC) belong ₁₇₃ to the graphical languages.

174 Graphical languages are more simple and immediate comparing to the
 175 textual ones and are preferred by programmers because of their features
 176 and simplicity, in particular the **Ladder Logic programming** (see Figure
 177 2.4 for a comparison).

```

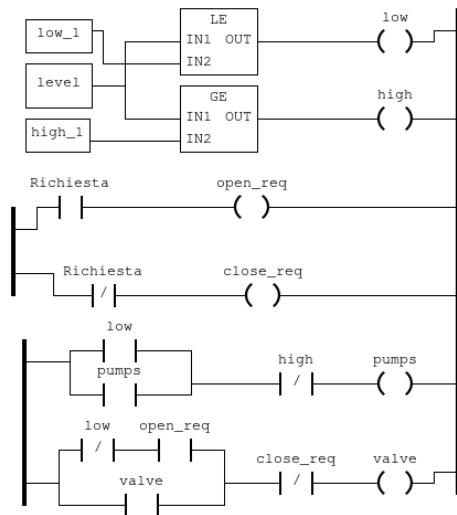
PROGRAM PLC1
VAR
  level AT %IW0 : INT;
  Richiesta AT %QX0..2 : BOOL;
  request AT %IW1 : INT;
  pumps AT %QX0..0 : BOOL;
  valve AT %QX0..1 : BOOL;
  low AT %MW0..0 : BOOL;
  high AT %MW0..1 : BOOL;
  open_req AT %MW0..3 : BOOL;
  close_req AT %MW0..4 : BOOL;
  low_1 AT %MW0 : INT := 40;
  high_1 AT %MW1 : INT := 80;
END_VAR
VAR
  LE3_OUT : BOOL;
  GE7_OUT : BOOL;
END_VAR

LE3_OUT := LE(level, low_1);
low := LE3_OUT;
GE7_OUT := GE(level, high_1);
high := GE7_OUT;
open_req := Richiesta;
close_req := NOT(Richiesta);
pumps := NOT(high) AND (low OR pumps);
valve := NOT(close_req) AND (open_req AND NOT(low) OR valve);

END_PROGRAM

CONFIGURATION Config0
  RESOURCE Res0 ON PLC
    TASK task0(INTERVAL := T#20ms,PRIORITY := 0);
    PROGRAM instance0 WITH task0 : PLC1;
  END_RESOURCE
END_CONFIGURATION
  
```

(a) Example of ST programming



(b) Example of Ladder Logic

Figure 2.4: Comparison between ST language and Ladder Logic

178 **PLC Security** PLCs were originally designed to operate as closed sys-
 179 tems, not connected and exposed to the outside world via communication
 180 networks: the question of the safety of these systems, therefore, was not
 181 a primary aspect. The advent of Internet has brought undoubted advan-
 182 tages, but has introduced problems relating to the safety and protection of
 183 PLCs from external attacks and vulnerabilities.

184 Indeed, a variety of different communication protocols used in ICSs are
 185 designed to be efficient in communications, but do not provide any secu-
 186 rity measure i.e. confidentiality, authentication and data integrity, which
 187 makes these protocols vulnerable against many of the IT classic attacks
 188 such as *Replay Attack* or *Man in the Middle Attack*.

189 Countermeasures to enhance security in PLC systems may include [10]:

- 190 • protocol modifications implementing **data integrity, authentication**
191 and **protection** against *Replay Attacks*
- 192 • use of *Intrusion Detection and Prevention Systems* (IDP)
- 193 • creation of *Demilitarized Zones* (DMZ) on the network

194 In addition to this, keeping the process network and Internet sepa-
195 rated, limiting the use of USB devices among users to reduce the risks of
196 infections, and using strong account management and maintenance poli-
197 cies are best practices to prevent attacks and threats and to avoid potential
198 damages.

199 **2.2.2.2 Remote Terminal Units**

200 *Remote Terminal Units* (RTUs) are computers with radio interfacing sim-
201 ilar to PLCs: they transmit telemetry data to the control center or to the
202 PLCs and use messages from the master supervisory system to control
203 connected objects [11].

204 The purpose of RTUs is to operate efficiently in remote and isolated
205 locations by utilizing wireless connections. In contrast, PLCs are designed
206 for local use and rely on high-speed wired connections. This key difference
207 allows RTUs to conserve energy by operating in low-power mode for ex-
208 tended periods using batteries or solar panels. As a result, RTUs consume
209 less energy than PLCs, making them a more sustainable and cost-effective
210 option for remote operations.

211 Industries that require RTUs often operate in areas without reliable ac-
212 cess to the power grid or require monitoring and control substations in re-
213 mote locations. These include telecommunications, railways, and utilities
214 that manage critical infrastructure such as power grids, pipelines, and wa-
215 ter treatment facilities. The advanced technology of RTUs allows these in-
216 dustries to maintain essential services, even in challenging environments
217 or under adverse weather conditions.

²¹⁸ 2.2.3 Area Control Layer

²¹⁹ The Area Control layer encompasses hardware and software systems
²²⁰ useful for supervising, monitoring and controlling the physical process,
²²¹ driving the behavior of the entire infrastructure. The layer includes sys-
²²² tems such as *Supervisory Control and Data Acquisition* (SCADA), *Distributed*
²²³ *Control Systems* (DCSs), that perform SCADA functions but are usually de-
²²⁴ ployed locally, and engineer workstations.

²²⁵ 2.2.3.1 Supervisory Control And Data Acquisition

²²⁶ *Supervisory Control And Data Acquisition (SCADA)* is a system of soft-
²²⁷ ware and hardware elements that allows industrial organizations to [12]:

- ²²⁸ • Control industrial processes locally or at remote locations;
- ²²⁹ • Monitor, gather, and process real-time data;
- ²³⁰ • Directly interact with devices such as sensors, valves, pumps, mo-
²³¹ tors, and more through human-machine interface (HMI) software;
- ²³² • Record and aggregate events to send to historian server.

²³³ The SCADA software processes, distributes, and displays the data, help-
²³⁴ ing operators and other employees analyze the data and make important
²³⁵ decisions.

²³⁶ 2.2.3.2 Human-Machine Interface

²³⁷ The *Human-Machine Interface* (HMI) is the hardware and software in-
²³⁸ terface that operators use to monitor the processes and interact with the
²³⁹ ICS. A HMI shows the operator and authorized users information about
²⁴⁰ system status and history; it also allows them to configure parameters on
²⁴¹ the ICS such as set points and, send commands and make control deci-
²⁴² sions [13].

²⁴³ The HMI can be in the form of a physical panel, with buttons and indicator
²⁴⁴ lights, or PC software as shown in Figure 2.5.

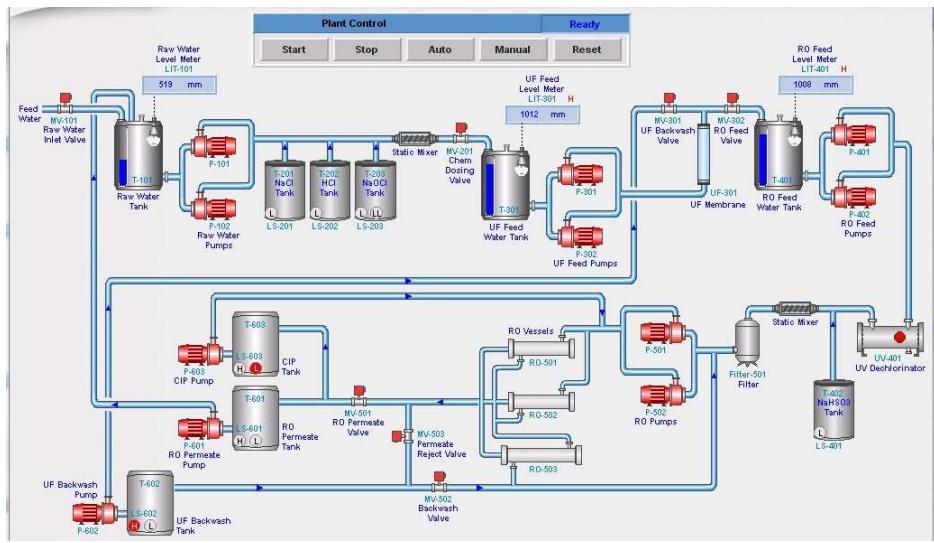


Figure 2.5: Example of HMI for a water treatment plant

2.2.4 Operations/Control Layer

Within this zone, there are specialized OT devices that are utilized to manage production workflows on the shop floor [14]. These devices include:

- Manufacturing Operations Management (MOM)* systems, which are responsible for overseeing production operations.
- Manufacturing Execution Systems (MES)*, which collect real-time data to optimize production processes.
- Data Historians*, which store process data and, in modern solutions, analyze it within its contextual framework.

2.2.5 Demilitarized Zone

This zone comprises security systems like firewalls, proxies, *Intrusion Detection and Prevention systems* (IDP) and *Security Information and Event Management* (SIEM) systems which are implemented to mitigate the risk of lateral threat movement between IT and OT domains. With the rise

260 of automation, the need for bidirectional data flows between OT and IT
261 systems has increased. The convergence of IT and OT in this layer can offer
262 organizations a competitive edge. However, it's important to note that
263 adopting a flat network approach in this context can potentially heighten
264 cyber risks for the organization.

265 2.2.6 Industrial Protocols

266 *Industrial Protocols* are the networks that are used to connect the dif-
267 ferent components of the ICS and allow them to communicate with each
268 other. Industrial Protocols can include wired and wireless networks, such
269 as Ethernet/IP, Modbus, DNP3, Profinet and others.

270 As mentioned at the beginning of this Chapter, industrial systems differ
271 from classical IT systems in the purpose for which they are designed: con-
272 trolling physical processes the former, processing and storing data the lat-
273 ter. For this reason, ICSs require different communication protocols than
274 traditional IT systems for real time communications and data transfer.

275 A wide variety of industrial protocols exists: this is because originally
276 each vendor developed and used its own proprietary protocol. How-
277 ever, these protocols were often incompatible with each other, resulting
278 in devices from different vendors being unable to communicate with each
279 other.

280 To solve this problem, standards were defined with a view to allowing
281 these otherwise incompatible device to intercommunicates.

282 Among all the various protocols, some have risen to prominence as widely
283 accepted standards. These *de facto* protocols are commonly utilized in in-
284 dustrial systems due to their proven reliability and effectiveness. In the
285 following sections, we will provide a brief overview of some of the most
286 prevalent and widely used protocols in the industry.

²⁸⁷ **2.2.6.1 Modbus**

²⁸⁸ *Modbus* is a serial communication protocol developed by Modicon (now
²⁸⁹ Schneider Electric) in 1979 for use with its PLCs [15] and designed ex-
²⁹⁰ pressly for industrial use: it facilitates interoperability of different devices
²⁹¹ connected to the same network (sensors, PLCs, HMIs, ...) and it is also
²⁹² often used to connect RTUs to SCADA acquisition systems.

²⁹³ Modbus is the most widely used communication protocol among in-
²⁹⁴ dustrial systems because it has several advantages:

- ²⁹⁵ • simplicity of implementation and debugging
- ²⁹⁶ • it moves raw bits and words, letting the individual vendor to repre-
²⁹⁷ sent the data as it prefers
- ²⁹⁸ • it is, nowadays, an **open** and *royalty-free* protocol: there is no need
²⁹⁹ to sustain licensing costs for implementation and use by industrial
³⁰⁰ device vendors

³⁰¹ Modbus is a **request/response** (or *master/slave*) protocol: this makes it
³⁰² independent of the transport layer used.

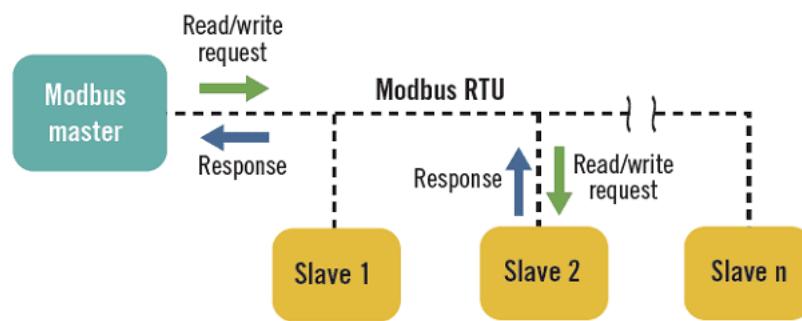


Figure 2.6: Modbus Request/Response schema

³⁰³ In this kind of architecture, a single device (master) can send requests
³⁰⁴ to other devices (slaves), either individually or in broadcast: these slave
³⁰⁵ devices (usually peripherals such as actuators) will respond to the master

306 by providing data or performing the action requested by the master using
 307 the Modbus protocol. Slave devices cannot generate requests to the master
 308 [16].

309 There are several variants of Modbus, of which the most popular and
 310 widely used are Modbus RTU (used in serial port connections) and Mod-
 311 bus TCP (which instead uses TCP/IP as the transport layer). Modbus TCP
 312 embeds a standard Modbus frame in a TCP frame (see Figure 2.7): both
 313 masters and slaves listen and receive data via TCP port 502.

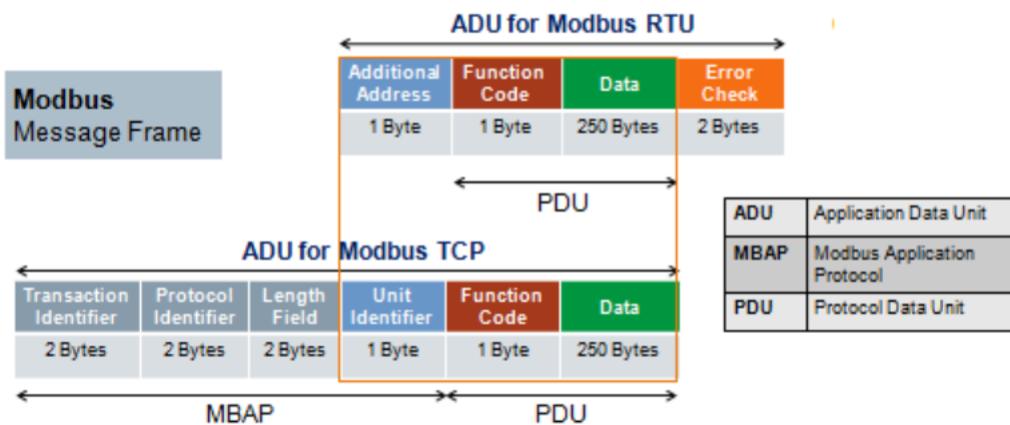


Figure 2.7: Modbus RTU frame and Modbus TCP frame

314 **Modbus registers** Modbus provides four object types, which map the
 315 data accessed by master and slave to the PLC memory:

- 316 • *Coil*: binary type, read/write accessible by both masters and slaves
- 317 • *Discrete Input*: binary type, accessible in read-only mode by masters
 318 and in read/write mode by slaves
- 319 • *Analog Input*: 16 bits in size (word), are accessible in read-only mode
 320 by masters and in read/write mode by slaves
- 321 • *Holding Register*: 16 bits in size (word), accessible in read/write mode
 322 by both masters and slaves. Holding Registers are the most com-
 323 monly used registers for output and as general memory registers.

³²⁴ **Modbus Function Codes** *Modbus Function Codes* are specific codes used
³²⁵ by the Modbus master within a request frame (see Figure 2.7) to tell the
³²⁶ Modbus slave device which register type to access and which action to
³²⁷ perform on it.

³²⁸ Two types of Function Codes exists: for data access and for diagnostic
³²⁹ Function Codes list for data access are listed in Table 2.2:

Function Code	Description
FC01	Read Coils
FC02	Read Discrete Input
FC03	Read Holding Registers
FC04	Read Analog Input Registers
FC05	Write/Force Single Coil
FC06	Write/Force Single Holding Register
FC15	Write/Force Multiple Coils
FC16	Write/Force Multiple Holding Registers

Table 2.2: Modbus Function Codes list

³³⁰ **Modbus Security Issues** Despite its simplicity and widespread use, the
³³¹ Modbus protocol does not have any security feature, which exposes it to
³³² vulnerabilities and attacks.

³³³ Data in Modbus are transmitted unencrypted (*lack of confidentiality*),
³³⁴ with no data integrity controls (*lack of integrity*) and authentication checks
³³⁵ (*lack of authentication*), in addition to the *lack of session*. Hence, the protocol
³³⁶ is vulnerable to a variety of attacks, such as Denial of Services (DoS), buffer
³³⁷ overflows and reconnaissance activities.

³³⁸ The easiest attack to bring to the Modbus protocol, however, is **packet**
³³⁹ **sniffing**: since, as mentioned earlier, network traffic is unencrypted and
³⁴⁰ the data transmitted is in cleartext, it is sufficient to use a packet sniffer to
³⁴¹ capture the network traffic, read the packets and thus gather informations

- 342 about the system such as ip addresses, function codes of requests and to
 343 modify the operation of the devices.

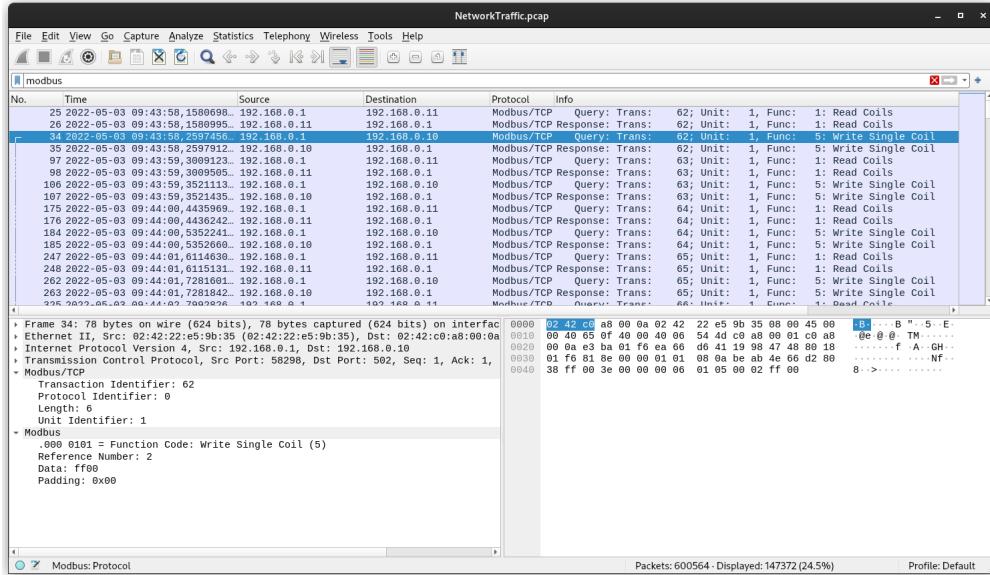


Figure 2.8: Example of packet sniffing on the Modbus protocol

344 To make the Modbus protocol more secure, an encapsulated version
 345 was developed within the *Transport Security Layer* (TLS) cryptographic
 346 protocol, also using mutual authentication. This version of the Modbus
 347 protocol is called **Secure Modbus** or **Modbus TLS**. In addition to this,
 348 Secure Modbus also includes X.509-type certificates to define permissions
 349 and authorisations [17].

350 2.2.6.2 EtherNet/IP

351 EtherNet/IP (where IP stands for *Industrial Protocol*) is an open industrial
 352 protocol that allows the *Common Industrial Protocol* (CIP) to run on a
 353 typical Ethernet network [18]. It is supported by ODVA [19].

354 EtherNet/IP uses the major Ethernet standards, such as IEEE 802.3 and
 355 the TCP/IP suite, and implements the CIP protocol stack at the upper layers
 356 of the OSI stack (see Figure 2.9). It is furthermore compatible with the
 357 main Internet standard protocols, such as SNMP, HTTP, FTP and DHCP,

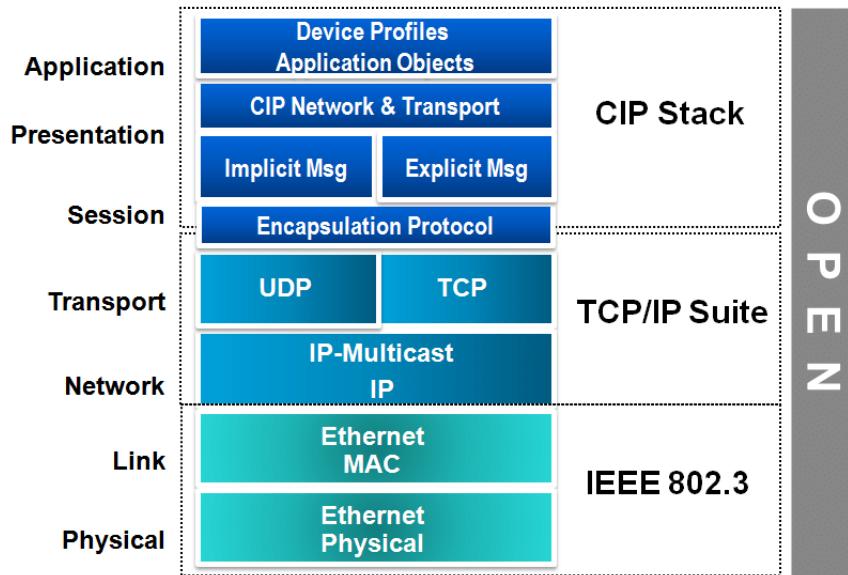


Figure 2.9: OSI model for EtherNet/IP stack

³⁵⁸ and other industrial protocols for data access and exchange such as *Open*
³⁵⁹ *Platform Communication* (OPC).

³⁶⁰ **Physical and Data Link layer** The use of the IEEE 802.3 standard allows
³⁶¹ EtherNet/IP to flexibly adopt different network topologies (star, linear,
³⁶² ring, etc.) over different connections (copper, fibre optic, wireless, etc.), as
³⁶³ well as the possibility to choose the speed of network devices.
³⁶⁴ IEEE 802.3 in addition defines at Data Link layer the *Carrier Sense Multiple*
³⁶⁵ *Access - Collision Detection* (CSMA/CD) protocol, which controls access to
³⁶⁶ the communication channel and prevents collisions.

³⁶⁷ **Transport layer** At the transport level, EtherNet/IP encapsulates mes-
³⁶⁸ sages from the CIP stack into an Ethernet message, so that messages can
³⁶⁹ be transmitted from one node to another on the network using the TCP/IP
³⁷⁰ protocol. EtherNet/IP uses two forms of messaging, as defined by CIP
³⁷¹ standard [18][20]:

- ³⁷² • **unconnected messaging:** used during the connection establishment
³⁷³ phase and for infrequent, low priority, explicit messages. Uncon-

374 nected messaging uses TCP/IP to transmit messages across the net-
375 work asking for connection resource each time from the *Unconnected*
376 *Message Manager* (UCMM).

- 377 • **connected messaging:** used for frequent message transactions or for
378 real-time I/O data transfers. Connection resources are reserved and
379 configured using communications services available via the UCMM.

380 EtherNet/IP has two types of message connection [18]:

- 381 – **explicit messaging:** *point-to-point* connections to facilitate *request-*
382 *response* transactions between two nodes. These connections use
383 TCP/IP service on port 44818 to transmit messages over Ether-
384 net.
- 385 – **implicit messaging:** this kind of connection moves application-
386 specific **real-time I/O data** at regular intervals. It uses multicast
387 *producer-consumer* model in contrast to the traditional *source-*
388 *destination* model and UDP/IP service (which has lower proto-
389 col overhead and smaller packet size than TCP/IP) on port 2222
390 to transfer data over Ethernet.

391 **Session, Presentation and Application layer** At the upper layers, Ether-
392 Net/IP implements the CIP protocol stack. We will discuss this protocol
393 more in detail in Section 2.2.6.3.

394 2.2.6.3 Common Industrial Protocol (CIP)

395 The *Common Industrial Protocol* (CIP) is an open industrial automation
396 protocol supported by ODVA. It is a **media independent** (or *transport in-*
397 *dependent*) protocol using a *producer-consumer* communication model and
398 providing a **unified architecture** throughout the manufacturing enterprise
399 [21][22].

400 CIP has been adapted in different types of network:

- 401 • **EtherNet/IP**, adaptation to *Transmission Control Protocol* (TCP) tech-
402 nologies
- 403 • **ControlNet**, adaptation to *Concurrent Time Domain Multiple Access*
404 (CTDMA) technologies
- 405 • **DeviceNet**, adaptation to *Controller Area Network* (CAN) technolo-
406 gies
- 407 • **CompoNet**, adaptation to *Time Division Multiple Access* (TDMA) tech-
408 nologies

409 **CIP objects** CIP is a *strictly object oriented* protocol at the upper layers:
410 each object of CIP has **attributes** (data), **services** (commands), **connec-
411 tions**, and **behaviors** (relationship between values and services of attributes)
412 which are defined in the **CIP object library**. The object library supports
413 many common automation devices and functions, such as analog and dig-
414 ital I/O, valves, motion systems, sensors, and actuators. So if the same
415 object is implemented in two or more devices, it will behave the same way
416 in each device [23].

417 **Security** [24] In EtherNet/IP implementation, security issues are the same
418 as in traditional Ethernet, such as network traffic sniffing and spoofing.
419 The use of the UDP protocol also exposes CIP to transmission route ma-
420 nipulation attacks using the *Internet Group Management Protocol* (IGMP)
421 and malicious traffic injection.

422 Regardless of the implementation used, it is recommended that certain
423 basic measures be implemented on the CIP network to ensure a high level
424 of security, such as *integrity, authentication and authorization*.

State of the Art

425 IN COVENTIONAL IT SYSTEMS, the objective of an attacker is to comprehend
426 the behavior of a program using diverse techniques in order to launch attacks that alter its execution flow, functionalities, or bypass limitations imposed by software licensing. These attack techniques involve an initial examination of the program, consisting of *static analysis* (i.e., analyzing the software without running it) and *dynamic analysis* (i.e., analyzing the program while it is running).

432 The outcome of these two investigative techniques is the *reverse engineering* of the software, which serves the purpose of identifying vulnerabilities or bugs and subsequently strategizing an attack.

435 In the context of OT systems, the notion of *reverse engineering* is not limited
436 to its conventional definition, but also includes the concept of **process
437 comprehension**. This term, introduced by Green et al. [25], refers to gaining a comprehensive understanding of the underlying physical process.

439 There is limited literature available concerning the gathering and analysis
440 of information related to the comprehension and operation of an Industrial
441 Control System (ICS). In Section 3.1, we will provide a brief overview of the existing literature on this topic, and in the subsequent sections, we
442 will specifically focus on one of the presented papers.

3.1 Literature on Process Comprehension

Keliris and Maniatikos The first approach presented in this section is by Keliris and Maniatikos [26]: they present a methodology for automating the reverse engineering of ICS binaries based on a *modular framework* (called ICSREF) that can reverse binaries compiled with CODESYS [27], one of the most popular and widely used PLC compilers, irrespective of the language used.

Yuan et al. Yuan et al. [28] propose a *data-driven* approach to discovering cyber-physical systems process behavior from data directly: to achieve this goal, they have implemented a framework whose purpose is to identify physical systems and transition logic inference, and to seek to understand the mechanisms underlying these processes, making furthermore predictions concerning their state trajectories based on the discovered models.

Feng et al. Feng et al. [29] developed a framework that can generate system *invariant rules* based on machine learning and data mining techniques from ICS operational data log. These invariants are then selected by systems engineers to derive IDS systems from them.

The experiment results on two different testbeds, the *Water Distribution system* (WaDi) and the *Secure Water Treatment system* (SWaT), both located at the iTrust - Center for Research in Cyber Security at the University of Singapore of Technology and Design [30], show that under the same false positive rate invariant-based IDSs have a higher efficiency in detecting anomalies than IDS systems based on a residual error-based model.

Pal et al. Pal et al. [31] work is somewhat related to Feng et al.'s: this paper describes a data-driven approach to identifying invariants automatically using *association rules mining* [32] with the aim of generating invariants sometimes hidden from the design layout. The study has the same objective of Feng et al.'s and uses too the iTrust SwaT System as testbed.

475 Currently this technique is limited to only pair wise sensors and
476 actuators: for more accurate invariants generation, the technique
477 adopted must be capable of deriving valid constraints across multiple
478 sensors and actuators.

479 **Winnicki et al.** Winnicki et al. [33] instead propose a different approach
480 to process comprehension based on the *attacker's perspective* and not
481 limited to mere *Denial of Service* (DoS): their approach is to discover
482 the dynamic behavior of the system, in a semi-automated and process-
483 aware way, through *probing*, that is, slightly perturbing the cyber
484 physical system and observing how it reacts to changes and how
485 it returns to its original state. The difficulty and challenge for the
486 attacker is to perturb the system in such a way as to achieve an ob-
487 servable change, but at the same time avoid this change being seen
488 as a system anomaly by the IDSs.

489 **Green et al.** Green et al. [25] also adopt an approach based on the at-
490 tacker's perspective: this approach consists of two practical exam-
491 ples in a *Man in the Middle* (MitM) scenario to obtain, correlate, and
492 understand all the types of information an attacker might need to
493 plan an attack to alter the process while avoiding detection.

494 The paper shows *step-by-step* how to perform a **ICS reconnaissance**, a
495 phase specifically designed to gather extensive intelligence on mul-
496 tiple fronts, including human factors, network and protocol infor-
497 mation, details about the manufacturing process, industrial applica-
498 tions, and potential vulnerabilities. The primary goal is to accumu-
499 late a wealth of information to enhance understanding and aware-
500 ness in these areas [34]).

501 Reconnaissance phase is fundamental to process comprehension and
502 thus to the execution of MitM attacks.

503 **Ceccato et al.** Ceccato et al. [9] propose a methodology based on a *black*
504 *box dynamic analysis* of an ICS using a reverse engineering tool to
505 derive from the scans performed on the memory registers of the ex-

26 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

506 posed PLCs and the Modbus protocol network scans an approximate
507 model of the physical process. This model is obtained by inferring
508 statistical properties, business process and system invariants from
509 data logs.

510 The proposed methodology was tested on a non-trivial case study,
511 using a virtualized testbed inspired by an industrial water treatment
512 plant.

513 In the next section we will examine this latest work in more detail,
514 which will be the basis for my work and thus the subsequent chap-
515 ters of this thesis.

516 3.2 Ceccato et al.'s black-box dynamic analysis 517 for water-tank systems

518 As previously mentioned, the paper introduces a methodology that re-
519 lies on black box dynamic analysis of an Industrial Control System (ICS)
520 and more particularly of its OT network. This methodology involves iden-
521 tifying potential Programmable Logic Controllers (PLCs) within the net-
522 work and scanning the memory registers of these identified controllers.
523 The purpose of this process is to obtain an approximate model of the con-
524 trolled physical process.

525 The primary goal of this black box analysis is to establish a correlation
526 between the different memory registers of the targeted PLCs and funda-
527 mental concepts of an OT network such as sensor values (i.e., measure-
528 ments), actuator commands, setpoints (i.e., range of values of a physical
529 variable), network communications, among others.

530 To accomplish this, the various types of memory registers are analyzed,
531 and attempts are made to determine the nature of the data they might
532 contain.

533 The second goal is to establish a relationship between the dynamic evolu-
534 tion of these fundamental concepts.

535 To accomplish this, Ceccato et al. have developed a prototype tool [35]
536 that facilitates the reverse engineering of the physical system. This tool
537 goes through four distinct phases:

- 538 1. **scanning of the system and data pre-processing:** this phase involves
539 gathering data to generate data logs for the registers of PLCs and for
540 Modbus network communications.
- 541 2. **graphs and statistical analysis:** The collected data is utilized to pro-
542 vide insights into the memory registers associated with the Modbus
543 protocol by leveraging graphs and statistical data. This analysis ap-
544 proach offers valuable information about the characteristics and pat-
545 terns of the memory registers.
- 546 3. **invariants inference and analysis:** generates system invariants, which
547 are used to identify specific patterns and regularities within the sys-
548 tem. Additionally, this phase provides users with the capability to
549 view invariants related to a particular sensor or actuator.
- 550 4. **business process mining and analysis:** Using event logs, this phase
551 involves reconstructing the business process that depicts how a pro-
552 cess is executed. This step enables a thorough understanding of the
553 sequence of events that occur in the system and how they are in-
554 terrelated, ultimately leading to a comprehensive overview of the
555 business process.

556 Figure 3.1 presents a schematic representation of the stages and the
557 workflow associated with this work, specifying tools and technologies
558 used. In the subsequent sections of this chapter, we will provide a detailed
559 exploration of each of these phases, offering a comprehensive understand-
560 ing of the entire process.

561 3.2.1 Testbed

562 Before delving into the description of the methodology's different phases,
563 let's first examine the testbed utilized to evaluate this approach. The testbed

28 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

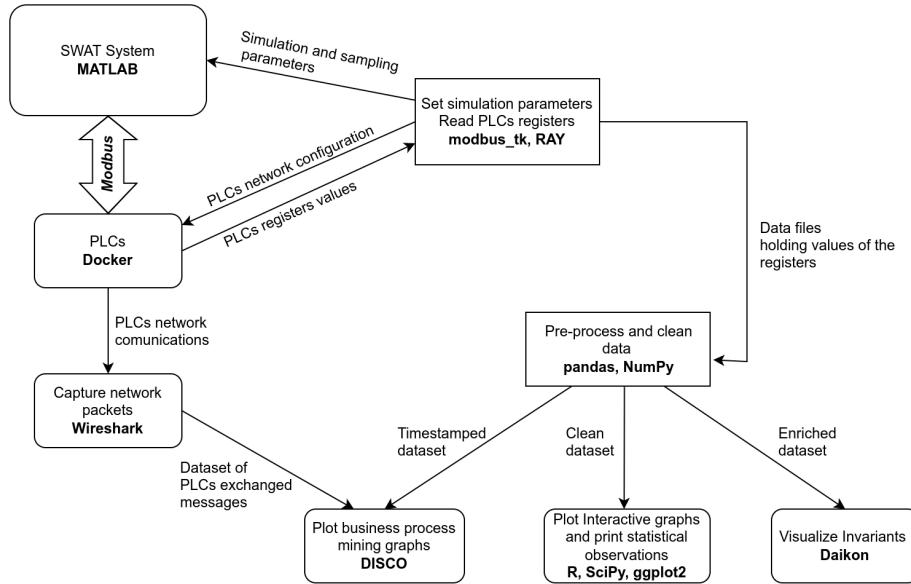


Figure 3.1: Workflow of Ceccato et al.'s stages and operations with used tools

564 employed for testing purposes is a (very) simplified rendition of the iTrust
 565 SWaT system [36], as implemented by Lanotte et al. [37]. Figure 3.2 pro-
 566 vides a graphical representation of the testbed. This simplified version
 567 comprises three stages, each governed by a dedicated PLC.

568 **Stage 1** During the initial stage, a **tank** referred to as T-201 with a capacity
 569 of 80 gallons is filled with raw water using the P-101 pump. Connected to the T-201 tank, the MV-301 motorized valve flushes out the
 570 accumulated water from the tank, directing it to the next stage. Initially,
 571 the water flows from the T-201 tank to the *filtration unit* (which
 572 is not specifically identified by any sensor), and subsequently to a
 573 **second tank** denoted as T-202, with a capacity of 20 gallons.
 574

575 **Stage 2** At the second stage, the water stored in tank T-202 flows into the
 576 *reverse osmosis unit* (RO), which serves as both a valve and a continu-
 577 ous water extractor. The purpose of the RO unit is to reduce organic
 578 impurities present in the water. Subsequently, the water flows from
 579 the *RO unit* to the third and final stage of the system.

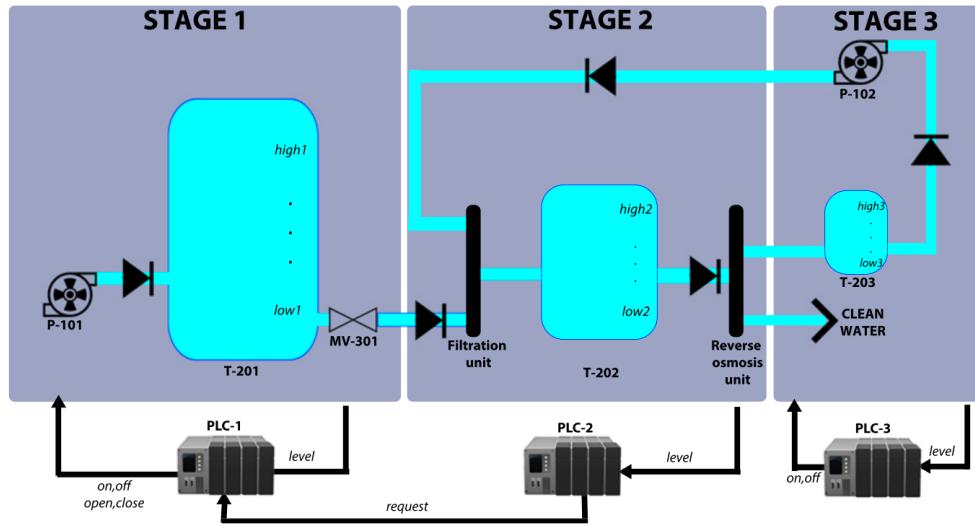


Figure 3.2: The simplified SWaT system used for running Ceccato et al. methodology

580 **Stage 3** At the third stage, the water coming from the *RO unit* undergoes
 581 division based on whether it meets the required standards. If the
 582 water is deemed clean and meets the standards, it is directed into
 583 the distribution system. However, if the water fails to meet the stan-
 584 dards, it is redirected to a *backwash tank* identified as T-203, which
 585 has a capacity of one gallon. The water stored in this tank is then
 586 pumped back to the stage 2 *filtration unit* using pump P-102.

587 As previously mentioned, each stage of the system is handled via a
 588 dedicated PLC, namely PLC1, PLC2, and PLC3, which are responsible for
 589 controlling their respective stages. Let's briefly explore the behavior of
 590 each PLC:

591 **PLC1** PLC1 monitors the level of tank T-201 and distinguishes three dif-
 592 ferent cases based on the level readings:

- 593 1. when the level of tank T-201 reaches the defined *low setpoint*
 594 *low1* (which is hardcoded in a specific memory register), PLC1
 595 **opens pump P-101 and closes valve MV-301**. This configura-
 596 tion allows the tank to be filled with water;

30 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

- 597 2. if the level of T-201 reaches the *high setpoint high1* (which is also
598 hardcoded in a specific memory register), then the pump **P-101**
599 **is closed**;
- 600 3. in cases where the level of T-201 is between the *low setpoint low1*
601 and the *high setpoint high1*, PLC1 waits for a request from PLC2
602 to open or close the valve MV-301. If a request to open the valve
603 MV-301 is received, water will flow from T-201 to T-202. How-
604 ever, if no request is received, the valve remains closed. In both
605 situations, the pump P-101 remains closed.

606 **PLC2** PLC2 monitors the level of tank T-202 and adjusts its behavior based
607 on the water level. There are three cases to consider:

- 608 1. when the water level in tank T-202 reaches *low setpoint low2* (also
609 hardcoded in the memory registers), PLC2 sends a request to
610 PLC1 through a Modbus channel to **open valve MV-301**. This
611 request is made in order to allow the water to flow from tank T-
612 201 to tank T-202. The transmission channel between the PLCs
613 is established by copying a boolean value from a memory reg-
614 ister of PLC2 to a corresponding register of PLC1.
- 615 2. when the water level in tank T-202 reaches the *high setpoint high2*
616 value (also hardcoded in the memory registers), PLC2 sends a
617 **close request to PLC1 for valve MV-301**. This request prompts
618 PLC1 to close the valve, stopping the flow of water from tank
619 T-201 to tank T-202.
- 620 3. In cases where the water level in tank T-202 is between the low
621 and high setpoints, the valve MV-301 remains in its current state
622 (open or closed) while the tank is either filling or emptying.

623 **PLC3** PLC3 monitors the level of the T-203 backwash tank and adjusts its
624 behavior accordingly. There are two cases to consider:

- 625 1. If the water level in the backwash tank reaches the *low setpoint*

626 *low3, PLC3 sets pump P103 to off.* This allows the backwash
627 tank to be filled.

- 628 2. If the water level in the backwash tank reaches the *high setpoint*
629 *high3, PLC3 opens pump P103.* This action triggers the pump-
630 ing of the entire content of the backwash tank back to the filter
631 unit of T-202.

632 **3.2.2 Scanning of the System and Data Pre-processing**

633 **Scanning tool** The Ceccato et al. scanning tool extends and generalizes
634 a project I did [38] for the "Network Security" and "Cyber Security for IoT"
635 courses taught by Professors Massimo Merro and Mariano Ceccato, re-
636 spectively, in the 2020/21 academic year. The original project involved,
637 in its first part, the recognition within a network of potential PLCs lis-
638 tening on the standard Modbus TCP port 502 using the Nmap module
639 for Python, obtaining the corresponding IP addresses: then a (sequential)
640 scan of a given range of the memory registers of the found PLCs was per-
641 formed to collect the register data. The data thus collected were saved to
642 a file in *JavaScript Object Notation* (JSON) format for later use in the second
643 part of my project.

644 The scanning tool by Ceccato et. al works in a similar way, but extends
645 what originally did by trying to discover other ports on which the Mod-
646 bus protocol might be listening (since in many realities Modbus runs on
647 different ports than the standard one, according to the concept of *security*
648 *by obscurity*) and, most importantly, by **parallelizing and distributing the**
649 **scan** of PLC memory registers through the Ray module [39], specifying
650 moreover the desired granularity of the capture. An example of raw data
651 capture can be seen at Listing 3.1:

```
652     "127.0.0.1/8502/2022-05-03 12_10_00.591": {  
653         "DiscreteInputRegisters": {"%IX0.0": "0"},  
654         "InputRegisters": {"%IW0": "53"},  
655         "HoldingOutputRegisters": {"%QW0": "0"},  
656         "MemoryRegisters": {"%MW0": "40", "%MW1": "80"},
```

32 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

```
657     "Coils": {"%QX0.0": "0"}}
```

Listing 3.1: Example of registers capture

658 The captured data includes PLC's IP address, Modbus port and timestamp
659 (first line), type and name of registers with their values read from the scan
660 (subsequent lines).

661 The tool furthermore offers the possibility, in parallel to the memory
662 registers scan, of **sniffing network traffic** related to the Modbus protocol
663 using the *Man in the Middle* (MitM) technique on the supervisory control
664 network using a Python wrapper for tshark/Wireshark [40] [41]. An ex-
665 ample of raw data obtained with this sniffing can be seen in Listing 3.2:

```
666     Time,Source,Destination,Protocol,Length,Function Code,  
667     ↳ Destination Port,Source Port,Data,Frame length on the  
668     ↳ wire,Bit Value,Request Frame,Reference Number,Info  
669 2022-05-03 11:43:58.158,IP_PLC1,IP_PLC2,Modbus/TCP,76,Read  
670     ↳ Coils,46106,502,,76,TRUE,25,,,"Response: Trans: 62;  
671     ↳ Unit: 1, Func: 1: Read Coils"
```

Listing 3.2: Example of raw network capture

672 **Data Pre-processing** The data collected by scanning the memory regis-
673 ters of the PLCs are then reprocessed by a Python script and converted
674 in order to create a distinct raw dataset in *Comma Separated Value* for-
675 mat (CSV) for each PLC, containing the memory register values associ-
676 ated with the corresponding controller registers. These datasets are repro-
677 cessed again through the Python modules for **pandas** [42] and **NumPy** [43]
678 by another script to first perform a **data cleanup**, removing all unused reg-
679 isters, **merged** into a single dataset, and finally **enriched** with additional
680 data, such as the **previous value** of all registers and the the **measurement**
681 **slope**, that is, the trend of the water level in the system tanks along the
682 system cycles.¹. See 3.2.7 for more detail.

683

¹Not all additional data are calculated and entered automatically by the tool: some are manually inserted.

684 This process leads to the creation of two copies of the full dataset: one en-
 685 riched with the additional data, but not timestamped, which will be used
 686 for the invariant analysis; the other unenriched, but timestamped, which
 687 will be used for business process mining.

688 **3.2.3 Graphs and Statistical Analysis**

689 The paper mentions the presence of a *mild graph analysis*, performed
 690 using the framework **R** [44] for statistical analysis at the time of data gath-
 691 ering to find any uncovered patterns, trends and identify measurements
 692 and/or actuator commands through the analysis of registers holding mu-
 693 table values.

694 There is actually no trace of this within the tool: *graph analysis* and *sta-*
 695 *tistical analysis* of the data contained in the PLC memory registers are in-
 696 stead performed using the **matplotlib libraries** and statistical algorithms
 697 made available by the **SciPy libraries** [45], through two separate Python
 698 scripts (see Figure 3.3).

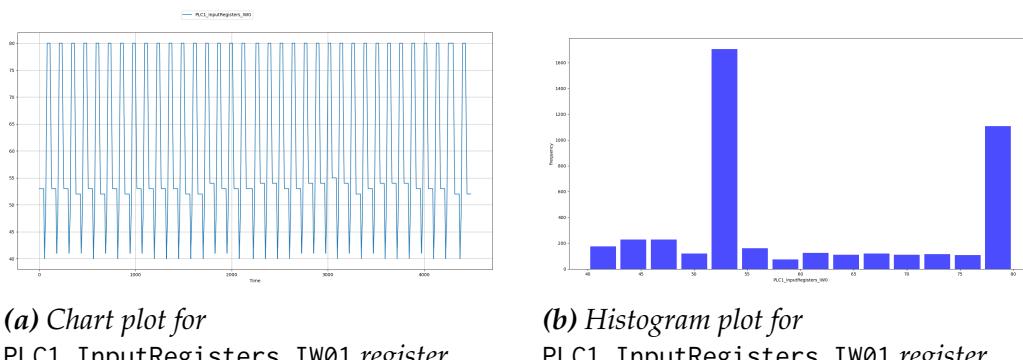


Figure 3.3: Output graphs from graph analysis

699 The first script plots the charts, one at the time, of certain registers en-
 700 tered by the user from the command line, plots in which one can see the
 701 trend of the data and get a first basic idea of what that particular regis-
 702 ter contains (a measurement, an actuation, a hardcoded setpoint, ...) and
 703 possibly the trend; the second script, instead, shows a **histogram and sta-**

34 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

704 **tistical informations** about the register entered as command-line input.
705 These informations include:

- 706 • the mean, median, standard deviation, maximum value and mini-
707 mum value
- 708 • two tests for the statistical distribution: *Chi-squared* test for unifor-
709 mity and *Shapiro-Wilk* test for normality, as shown in Listing 3.3:

```
710     Chi-squared test for uniformity
711     Distance      pvalue      Uniform?
712     12488.340    0.00000000    NO
713
714     Shapiro-Wilk test for normality
715     Test statistic   pvalue      Normal?
716     0.844        0.00000000    NO
717
718     Stats of PLC1_InputRegisters_IW0
719     Sample mean = 60.8881; Stddev = 13.0164; max = 80; min =
720     ↪ 40 for 4488 values
```

Listing 3.3: Statistical data for PLC1_InputRegisters_IW0 register

721 3.2.4 Invariants Inference and Analysis

722 For invariant analysis Ceccato et al. rely on **Daikon** [46], a framework
723 to **dynamically detect likely invariants** within a program. An *invariant*
724 is a property that holds at one or more points in a program, properties
725 that are not normally made explicit in the code, but within assert state-
726 ments, documentation and formal specifications: invariants are useful in
727 understanding the behavior of a program (in our case, of the cyber physi-
728 cal system).

729 Daikon uses *machine learning* techniques applied to arbitrary data with
730 the possibility of setting custom conditions for analysis by using a spe-
731 cific file [47] with a *.spinfo* extension (see Listing 3.4). The framework is
732 designed to find the invariants of a program, with various supported pro-
733 gramming languages, starting from the direct execution of the program

734 itself or passing as input the execution run (typically a file in CSV format):
 735 the authors of the paper tried to apply it by analogy also to the execution
 736 runs of a cyber physical system, to extract the invariants of this system.

```
737 PPT_NAME aprogram.point:::POINT
738 VAR1 > VAR2
739 VAR1 == VAR3 && VAR1 != VAR4
```

Listing 3.4: Generic example of a .spinfo file for customizing rules in Daikon

740 Therefore, Daikon is fed with the no-timestamp enriched dataset ob-
 741 tained in the pre-processing phase²: a simple bash script launches Daikon
 742 (optionally specifying the desired condition for analysis in the *.spinfo* file),
 743 which output is simply redirected to a text file containing the general in-
 744 variants of the system (i.e., valid regardless of any custom condition spec-
 745 ified), those generated based on the custom condition in the *.spinfo* file,
 746 and those generated based on the negation of the condition (see Listing
 747 3.5 below).

```
748 =====
749 aprogram.point:::POINT
750 PLC2_MemoryRegisters_MW1 == PLC3_MemoryRegisters_MW1
751 PLC1_MemoryRegisters_MW0 == 40.0
752 PLC1_MemoryRegisters_MW1 == 80.0
753 PLC1_Coils_QX00 one of { 0.0, 1.0 }
754 [...]
755 =====
756 aprogram.point:::POINT; condition="PLC1_InputRegisters_IW0
757     ↪ > 60"
758 PLC1_InputRegisters_IW0 > PLC1_MemoryRegisters_MW0
759 PLC1_InputRegisters_IW0 > PLC1_Min_safety
760 PLC1_MemoryRegisters_MW0 < prev_PLC1_InputRegisters_IW0
761 [...]
762 =====
763 aprogram.point:::POINT; condition="not(
764     ↪ PLC1_InputRegisters_IW0 > 60)"
765 PLC1_InputRegisters_IW0 < PLC1_MemoryRegisters_MW1
```

²In the paper, the timestamped dataset is erroneously mentioned as input: probably a simple typo during the writing of the paper

36 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

```
766     PLC1_InputRegisters_IW0 < PLC1_Max_safety  
767     PLC1_MemoryRegisters_MW1 > prev_PLC1_InputRegisters_IW0  
768     [...]
```

Listing 3.5: The three sections of Daikon analysis outcomes

769 When the analysis is finished, the user is asked to enter the name of a reg-
770 istry to view its related invariants.

771
772 Some examples of invariants derived from the enriched dataset may be:

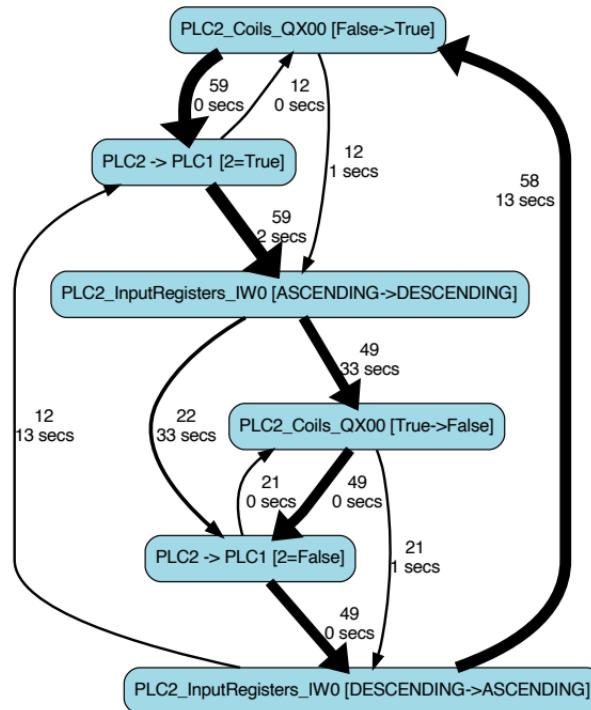
- 773 • measurements bounded by some setpoint;
- 774 • actuators state changes occurred in the proximity of setpoints or,
775 vice versa, proximity of setpoints upon the occurrence of an actuator
776 state change;
- 777 • state invariants of some actuators correspond to a specific trend in
778 the evolution of the measurements (ascending, descending, or sta-
779 ble) or, vice versa, the measurements trend corresponds to a specific
780 state invariant of some actuators.

781 3.2.5 Business Process Mining and Analysis

782 *Process mining* is the analysis of operational processes based on the
783 event log [48]: the aim of this analysis is to **extract useful informations**
784 from the event data to **reconstruct and understand the behavior** of the
785 business process and how it was actually performed.

786
787 In the considered system, process mining begins by analyzing the event
788 logs derived from scanning the memory registers of the PLCs and moni-
789 toring the network communications associated with the Modbus protocol,
790 as detailed in Subsection 3.2.2. These event logs serve as the *execution trace*
791 of the system. A Java program is utilized to extract and consolidate infor-
792 mation from these event logs, resulting in a CSV format file that captures
793 the relevant data.

This file is fed to **Disco** [49], a commercial process mining tool, which generates an *activity diagram* similar to UML Activity Diagram and whose nodes represent the activities while the edges represent the relations between these activities. In Figure 3.4 we can see an example of this diagram referred to PLC2 of the testbed: nodes represent the trend of register associated with measurement, actuator state changes, and communications between PLCs involving these state changes, while edges represent transitions with their associated time duration and frequency.



The *business process* obtained in this way provides an **overview of the system** and makes it possible to **make conjectures** about its behavior, particularly between changes in actuator state and measurement trends (i.e., a given change in state of some actuators corresponds to a specific measurement trend and vice versa), and with the possibility of **establishing causality** between Modbus communications and state changes within the physical system.

809 **3.2.6 Application**

810 In this section we will see how the black box analysis presented above
811 in its various phases is applied in practice, using the testbed described in
812 Subsection 3.2.1. The methodology supports a *top-down approach*: that
813 is, we start with an overview of the industrial process and then gradually
814 refine our understanding of the process by descending to a higher and
815 higher level of detail based on the results of the previous analyses and
816 focusing on the most interesting parts of the system for further in-depth
817 analysis.

818 **Data Collection and Pre-processing** According to what is described in
819 the paper, the data gathering process lasted six hours, with a granular-
820 ity of one data point per second (a full system cycle takes approximately
821 30 minutes). Each datapoint consists of 168 attributes (55 registers plus
822 a special register concerning the tank slope of each PLC) after the en-
823 richment. In addition, IP addresses are automatically replaced by an ab-
824 stract name identified by the prefix PLC followed by a progressive integer
825 (PLC1, PLC2, PLC3), in order to make reading easier.

826 **Graphs and Statistical Analysis** Graphs and Statistical Analysis revealed
827 three properties regarding the contents of the registers:

828 **Property 1:** PLC1_MemoryRegisters_MW0, PLC1_MemoryRegisters_MW1,
829 PLC2_MemoryRegisters_MW0, PLC2_MemoryRegisters_MW1,
830 PLC3_MemoryRegisters_MW0 and PLC3_MemoryRegisters_MW1
831 registers contain constant integer values (40, 80, 10, 20, 0, 10 respec-
832 tively)³. The authors speculate that they may be (relative) hardcoded
833 **setpoints**.

³From my tests on the original tool and dataset, the PLC3_MemoryRegisters_MW0 register is deleted during the *pre-processing* phase, as it is recognized as an unused register because of the constant value "0" it takes on. This leads me to assume that the properties are derived from a human read of the dataset prior to the *pre-processing* phase.

834 **Property 2:** PLC1_Coils_QX01, PLC1_Coils_QX02, PLC2_Coils_QX01,
 835 PLC2_Coils_QX02, PLC3_Coils_QX01 and PLC3_Coils_QX03 contain mu-
 836 table binary (Boolean) values. The authors speculate that these reg-
 837 isters can be associated with the **actuators** of the system.

838 **Property 3:** PLC1_InputRegisters_IW0, PLC2_InputRegisters_IW0 and
 839 PLC3_InputRegisters_IW0 registers contain mutable values.

840 Property 3 suggests that those registers might contain **values related to**
 841 **measurements**: it is therefore necessary to investigate further to see if the
 842 conjecture (referred to as *Conjecture 1* in the paper) is correct.

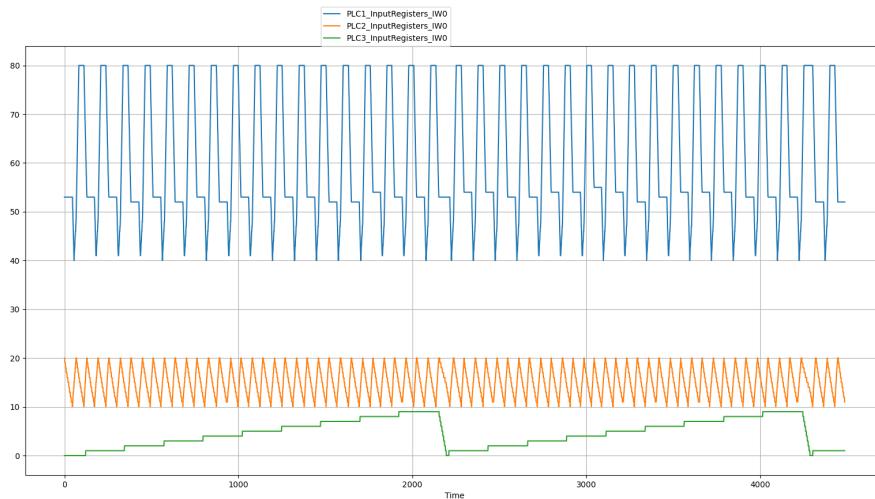


Figure 3.5: Execution traces of *InputRegisters_IW0* on the three PLCs

843 The graph analysis of the *InputRegisters_IW0* registers of the three
 844 PLCs (summarized in Figure 3.5 with a single plot) not only seems to con-
 845 firm the conjecture, but also allows the measurements to be correlated with
 846 the contents of the *MemoryRegisters_MW0* and *MemoryRegisters_MW1* regis-
 847 ters to the measurements, which may well represent the **relative setpoints**
 848 **of the measurements**. Hence, we have *Conjecture 2* described in the paper
 849 referring to the relative setpoints:

850

851 **Conjecture 2:**

852 - the relative setpoints for *PLC1_InputRegisters_IW0* are 40 and 80;

40 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

- 853 - the relative setpoints for PLC2_InputRegisters_IW0 are 10 and 20;
854 - the relative setpoints for PLC3_InputRegisters_IW0 0 and 9.

855 Further confirmation of this conjecture may come from statistical anal-
856 ysis. Indeed, in the example in Listing 3.1, some statistical data are given
857 for the register PLC1_InputRegisters_IW0, including the maximum value
858 and the minimum value: these values are, in fact, 80 and 40 respectively.

859 **Business Process Mining and Analysis** With Business Process Mining,
860 the authors aim to **visualize and highlight relevant system behaviors** by
861 relating PLC states and Modbus commands.

862 Through analysis of the activity diagrams shown in Figure 3.6, drawn
863 through Disco, they derive the following properties and conjectures:

864 **Property 4:** PLC2 sends messages to PLC1 (see Figure 3.6b) which are
865 recorded to PLC1_Coils_QX02.

866 **Conjecture 3:** PLC2_Coils_QX00 determines the trend in tank T-202 (Figure
867 3.6b).

868 When this register is set to *True*, the input register PLC2_InputRegisters_IW0
869 related to the tank controlled by PLC2 starts an **ascending trend**; vice
870 versa, when the coil register is set to *False*, the input register starts a
871 **descending trend**.

872 **Conjecture 4:** If PLC1_Coils_QX00 change his value to True, trend in tank
873 T-201, related to PLC1_InputRegisters_IW0 and controlled by PLC1,
874 become **ascending** (see Figure 3.6a)

875 **Conjecture 5:** PLC3_Coils_QX00 starts a **decreasing trend** in tank T-203, re-
876 lated to PLC3_InputRegisters_IW0 and controlled by PLC3, whereas
877 PLC3_Coils_QX02 starts an **increasing trend** on the tank (see Figure
878 3.6c)

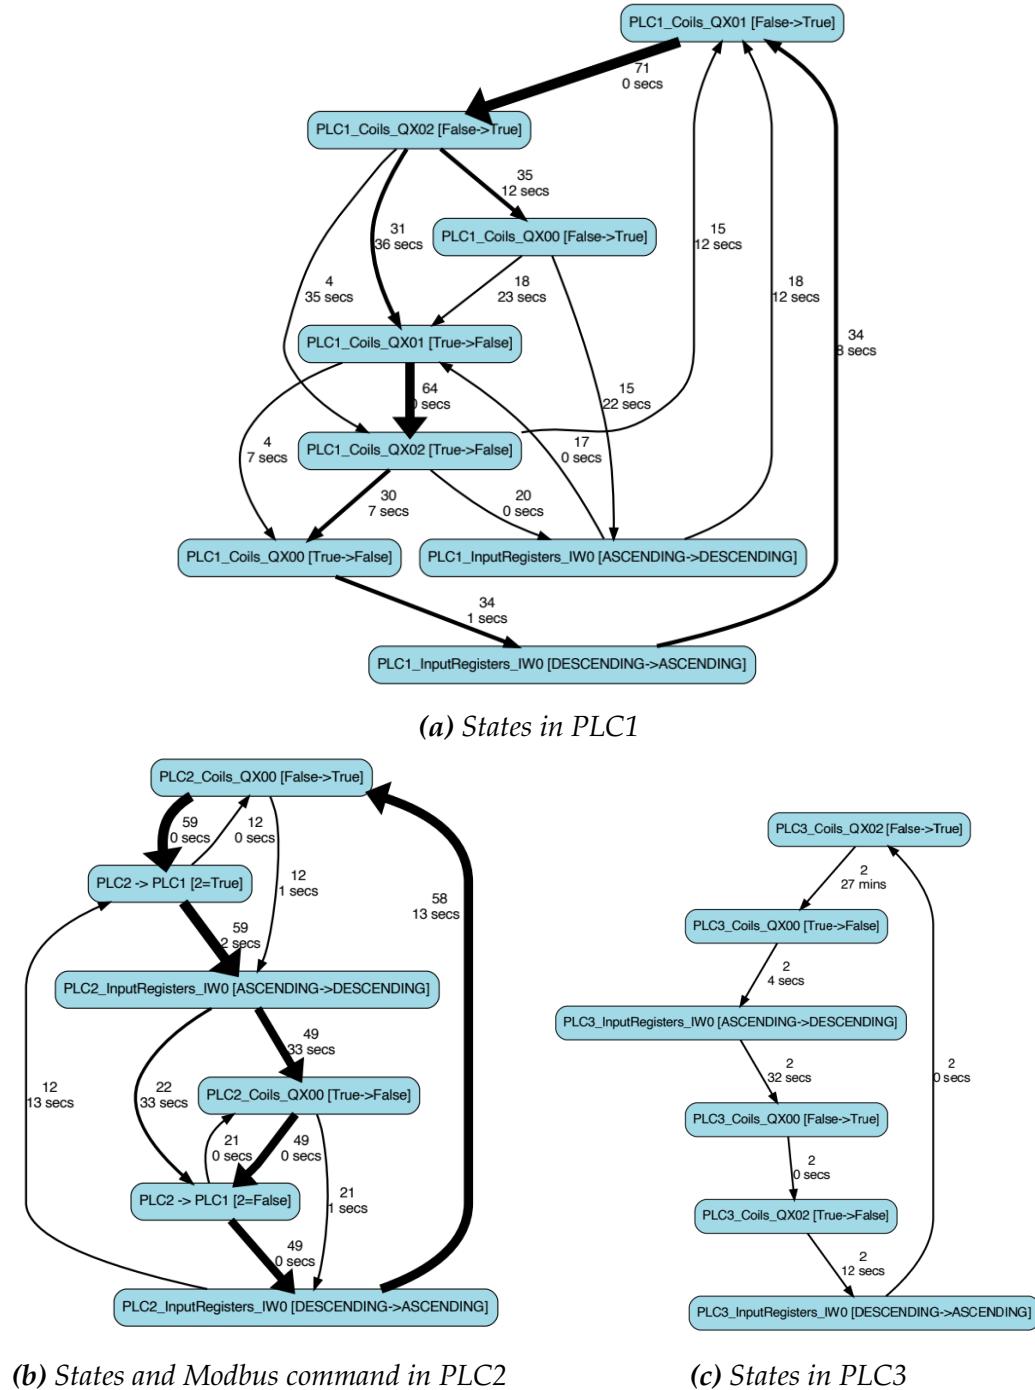


Figure 3.6: Business process with states and Modbus commands for the three PLCs

42 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

879 **Invariant Inference and Analysis** The last phase of the analysis of the
880 example industrial system is invariant analysis, performed through Daikon
881 framework. At this stage, an attempt will be made to confirm what has
882 been seen previously and to derive new properties of the system based on
883 the results of the Daikon analysis.

884 To get gradually more and more accurate results, the authors presum-
885 ably performed more than one analysis with Daikon, including certain
886 rules within the *splitter information file* (see Section 3.2.4 and Listing 3.4)
887 based on specific conditions placed on the measurements, for example, the
888 level of water contained in a tank. Given moreover the massive amount
889 of invariants generated by Daikon's output, it is not easy to identify and
890 correlate those that are actually useful for analysis: this must be done man-
891 ually.

892 However, it was possible to have confirmation of the conjectures made
893 in the previous stages of the analysis: starting with the setpoints, analyz-
894 ing the output of the invariants returned by Daikon⁴ reveals that

895
896 PLC1_InputRegisters_IW0 >= PLC1_MemoryRegisters_MW0 == 40.0
897 PLC1_InputRegisters_IW0 <= PLC1_MemoryRegisters_MW1 == 80.0
898 PLC2_InputRegisters_IW0 >= PLC2_MemoryRegisters_MW0 == 10.0
899 PLC2_InputRegisters_IW0 <= PLC2_MemoryRegisters_MW1 == 20.0
900 PLC3_InputRegisters_IW0 >= PLC3_MemoryRegisters_MW0 == 0.0
901 PLC3_InputRegisters_IW0 <= PLC3_MemoryRegisters_MW1 == 9.0
902
903 i.e., that the MemoryRegisters_MW0 and MemoryRegisters_MW1 registers of
904 each PLC contain the **absolute minimum and maximum setpoints**, re-
905 spectively (*Property 5*).

906 There is also a confirmation regarding *Property 4*: from the computed
907 invariants it can be seen that

⁴The invariants shown here are a manual summary and derivation of those actually returned in output by Daikon. I will discuss this more in Section 3.2.7

908
909 PLC1_Coils_QX01 == PLC1_Coils_QX02 == PLC2_Coils_QX00
910
911 and from this derive that there is a **communication channel between PLC2**
912 **and PLC1**, where the value of PLC2_Coils_QX00 is copied to PLC1_Coils_QX01
913 and PLC1_Coils_QX02 (*Property 6*).

914 Regarding the **relationships between actuator state changes and mea-**
915 **surement trends**, invariant analysis yields the results summarized in the
916 following rules:

917 **Property 7:** Tank T-202 level *increases* iff PLC1_Coils_QX01 == True. Oth-
918 erwise, if PLC1_Coils_QX01 == False will be *non-increasing*.

919 This is because if the coil is *True* the condition

920 PLC2_InputRegisters_IW0 == PLC2_MemoryRegisters_MW0 == 20.0 && PLC2_slope > 0
921 is verified. On the opposite hand, if the coil is *False*, the condition
922 PLC2_InputRegisters_IW0 == PLC2_MemoryRegisters_MW0 == 20.0 && PLC2_slope <= 0 is verified. The
923 *slope* is increasing if > 0, decreasing if < 0, stable otherwise.

924 **Property 8:** Tank T-201 level *increases* iff PLC1_Coils_QX00 == True. On the
925 other hand, if PLC1_Coils_QX00 == False and if PLC1_Coils_QX01 ==
926 True the level will be *non-decreasing*.

927 **Property 9:** Tank T-203 level *decreases* iff PLC3_Coils_QX00 == True. It will
928 be *non-decreasing* if PLC1_Coils_QX00 == False.

929 The last two properties concern the **relationship between actuator state**
930 **changes and the setpoints**: it is intended to check what happens to the
931 actuators when the water level reaches one of these setpoints. From the
932 analysis of the relevant invariants, the following properties are derived:

933 **Property 10:** Tank T-201 reaches the upper absolute setpoint when
934 PLC1_Coils_QX00 changes its state from *True* to *False*. If the coil changes
935 from *False* to *True*, the tank reaches its absolute lower setpoint.

44 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

936 **Property 11:** Tank T-203 reaches the upper absolute setpoint when
937 PLC3_Coils_QX00 changes its state from *True* to *False*. If the coil changes
938 from *False* to *True*, the tank reaches its absolute lower setpoint.

939 3.2.7 Limitations

940 The methodology proposed by Ceccato et al. is certainly valid and
941 offers a good starting point for approaching the reverse engineering of
942 an industrial control system from the attacker's perspective, while also
943 providing a tool to perform this task.

944 The limitations of this approach, however, all lie in the tool mentioned
945 above and also in the testbed described in Section 3.2.1. In this section
946 I will explain which are the criticisms of each phase, while in Chapter 4
947 I will formulate proposals to improve and make this methodology more
948 efficient.

949 **General Criticism** There are several critical aspects associated with the
950 application of this approach: the primary one concerns the fact that the
951 proposed tool seems to be built specifically for the testbed used and that
952 it is not applicable to other contexts, even to the same type of industrial
953 control system (water treatment systems, in this case).

954 What severely limits the analysis performed with the tool implemented
955 by Ceccato et al. is the use of *ad hoc* solutions and *a posteriori* interventions
956 done manually on the datasets after the data gathering process: I will dis-
957 cuss this last aspect in more detail later.

958 Moreover, there is the presence of many *hardcoded* variables and condi-
959 tions within the scripts: this makes the system unconfigurable and unable
960 to properly perform the various stages of the analysis as errors can occur
961 due to incorrect data and mismatches with the system under analysis.

962 Having considered, furthermore, only the Modbus protocol for network
963 communications between the PLCs is another major limiting factor and

964 does not help the methodology to be adaptable to different systems com-
965 municating with different protocols (sometimes even multiple ones on the
966 same system).

967 Let us now look at the limitations and critical aspects of each phase.

968 **Testbed** The testbed environment used by Ceccato et. al is entirely simu-
969 lated, from the physical system to the control system. The PLCs were built
970 with **OpenPLC** [50] in a Docker environment [51], while the physics part
971 was built through **Simulink** [52].

972 OpenPLC is an open source cross-platform software that simulates the
973 hardware and software functionality of a physical PLC and also offers a
974 complete editor for PLC program development with support for all stan-
975 dard languages: *Ladder Logic* (LD), *Function Block Diagram* (FBD), *Instruc-*
976 *tion List* (IL), *Structured Text* (ST), and *Sequential Function Chart* (SFC).

977 It is for sure an excellent choice for creating a zero-cost industrial or home
978 automation and *Internet of Things* (IoT) system that is easy to manage via a
979 dedicated, comprehensive and functional web interface. In spite of these
980 undoubted merits, however, there are (at the moment) **very few supported**
981 **protocols**: the main one and also referred to in the official documentation
982 is **Modbus**, while the other protocol is DNP3.

983 ***First limitation*** The biggest problem with the testbed, however, is not
984 with the controller part, but with the **physical part**: first of all, it
985 must be said that although this is something purely demonstrative
986 even though it is fully functional, the implemented Simulink model
987 is really **oversimplified** compared to the iTrust SWaT system, which
988 itself is a scaled-down version of a real water treatment plant. In
989 fact, in the entire system there are only three actuators, two of which
990 are connected to the same tank and controlled by the same PLC, and
991 sensors related only to the water level in the system's tanks: in a
992 real system there are many more *field devices*, which can monitor and
993 control other aspects of the system beyond the mere contents of the

46 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

994 tanks. Consider, for example, measuring and controlling the chemi-
995 cals in the water, the pressure of the liquid in the filter unit, or more
996 simply the amount of water flow at a given point or time.

997 All these must be considered and represent a number of additional
998 variables that makes analysis and consequently reverse engineering
999 of the system more difficult.

1000 ***Second limitation*** The second critical aspect concerns the **simulation of**
1001 **the physics of the liquid** inside the tanks: Simulink does not con-
1002 sider the fact that inside a tank that is filling (emptying) the liquid
1003 in it undergoes **fluctuations** which cause the level sensor not to see
1004 the water level constantly increasing (decreasing) or at most being
1005 stable at each point of detection. Figure 3.7 exemplifies more clearly
1006 with an example the concept just expressed: these oscillations cause
1007 a **perturbation** in the data.

1008 This issue leads to the difficulty, on a real physical system, of **cor-**
1009 **rectly calculating the trend of a measurement** by using the slope
1010 attribute: if this was obtained with a too low granularity, the trend
1011 will be oscillating between increasing and decreasing even when in
1012 reality this would be in general increasing (decreasing) or stable; on
1013 the other hand, if the slope was obtained with a too high granularity
1014 there is a loss of information and the trend may be "flattened" with
1015 respect to reality.

1016 In the present case, the slope in the Simulink model was calculated
1017 statically with a (very) low granularity, 5 and 6 seconds according
1018 to the Properties 7 and 9 described in the original paper: an aver-
1019 agey careful reader will have already guessed that this granularity
1020 is inapplicable to the real system in Figure 3.7b. As we will later see,
1021 we need to **operate on the data perturbations** to be able to obtain a
1022 suitable granularity and a correct calculation of the slope and conse-
1023 quently of the measurement trend.

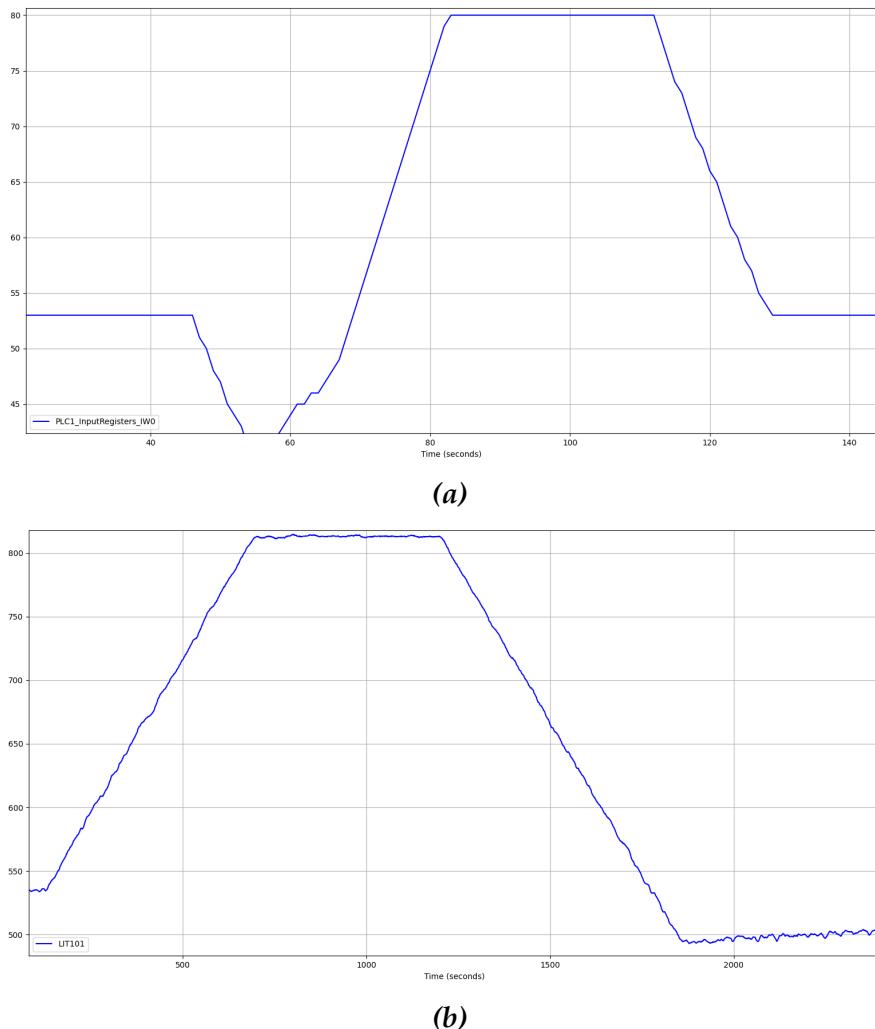


Figure 3.7: Water physics compared: simulated physics in the Simulink model (a) and physics in a real system (iTrust SWaT) (b). Fluctuations in the tank level in (b), almost completely absent in (a), can be appreciated.

1024 **Pre-processing** In the pre-processing phase, the authors make use of a
 1025 Python script to merge all the datasets of the individual PLCs into a single
 1026 dataset, remove the (supposedly) unused registers, and finally enrich the
 1027 obtained dataset with additional attributes. These attributes, as seen in
 1028 3.2.2, are:

- 1029 • the **previous value** of all registers;

48 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

- 1030 • some **additional relative setpoints** named PLC x _Max_safety and
1031 PLC x _Min_safety (where x is the PLC number), which represent a
1032 kind of alert on reaching the maximum and minimum water levels
1033 of the tanks;
- 1034 • the **measurement slope**.

1035 ***First limitation*** Merging the datasets of all individual PLCs into a single
1036 dataset representing the entire system can be a sound practice if the
1037 system to be analyzed is (very) small as is the testbed analyzed here,
1038 consisting of a few PLCs and especially a few registers. If, however,
1039 the complexity of the system increases, this type of merging can be-
1040 come counterproductive and make it difficult to analyze and under-
1041 stand the data obtained in subsequent steps.

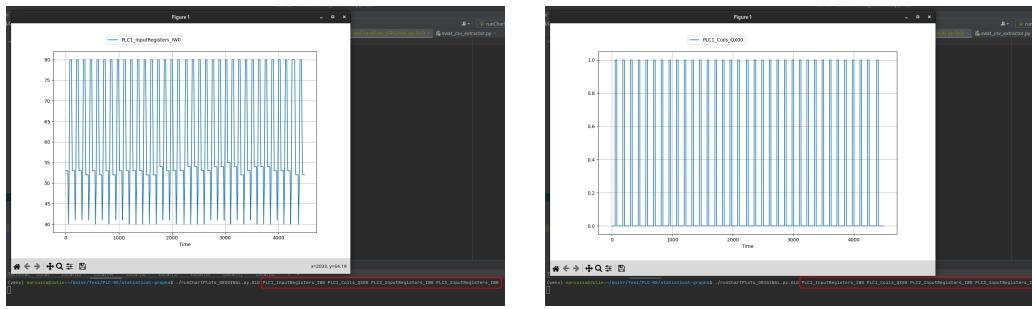
1042 In short, there is no possibility to analyze only a subsystem and thus
1043 make the analysis faster and more understandable. Moreover, a data
1044 gathering can take up to days, and the analyst/attacker may need to
1045 make an analysis of the system isolating precise time ranges, ignor-
1046 ing everything that happens before and/or after: all of this, with the
1047 tool we have seen, cannot be done.

1048 ***Second limitation*** Regarding the additional attributes, looking at the code
1049 of the script that performs the enrichment, we observed that **some at-**
1050 **tributes were manually inserted** after the merging phase: we are re-
1051 ferring in particular to the attributes PLC x _Max_safety and PLC x _Min_safety,
1052 whose references were moreover hardcoded into the script, and the
1053 *slope* whose calculation method we mentioned in the previous para-
1054 graph about the testbed limitations.

1055 In the end, only the attribute *prev* related to the value at the previous
1056 point of the detection is inserted automatically for all registers, more-
1057 over without the possibility to choose whether this attribute should
1058 be extended to all registers or only to a part.

1059 **Graphs and Statistical Analysis** Describing the behavior of graphical
1060 analysis in Section 3.2.3 we had already mentioned that only one register

1061 plot at a time was shown and not, for example, a single window containing
 1062 the charts of all registers entered by the user as input from the com-
 1063 mand line, such as in Figure 3.5. Figure 3.8 shows the actual behavior
 1064 of graphical analysis: note that although we have specified four registers
 1065 (highlighted in red in the figures) as command-line parameters, only one
 1066 at a time is shown and it is necessary to close the current chart in order to
 1067 display the next one.



(a) Chart for PLC1_InputRegisters_IW0

(b) Chart for PLC1_Coils_QX00

Figure 3.8: Behavior of the Graph Analysis on the Ceccato et al.'s tool

1068 ***First limitation*** While displaying charts for individual registers still pro-
 1069 vides useful information about the system such as the distinction
 1070 between actuators and measurements and the general trend of the
 1071 latter, single display does not allow one to catch, or at least makes it
 1072 difficult, the relationship that exists between actuators and measure-
 1073 ments, where it exists, because a view of the system as a whole is
 1074 missing.

1075 In this way, the risk is to make conjectures about the behavior of the
 1076 system that may prove to be at least imprecise, if not inaccurate.

1077 ***Second limitation*** On the other hand, regarding the statistical analysis,
 1078 two observations need to be made: the first is that for the given sys-
 1079 tem, I personally was unable to appreciate the usefulness of the gen-
 1080 erated histogram in Figure 3.3b, as it does not provide any particular
 1081 new information that has not already been obtained from the graph-
 1082 ical analysis (except maybe something marginal); the second obser-

50 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

1083 vation pertains to the presentation of statistical information obtained
1084 from the histogram plot. In certain cases, the histogram plot itself
1085 can overshadow the displayed statistical information. These statis-
1086 tics are actually shown on the terminal from which the script is exe-
1087 cuted. However, to an inattentive or unfamiliar user, these statistics
1088 may be mistaken for debugging output or warnings, as they coin-
1089 cide with the display of the histogram plot window, which takes the
1090 focus (see Figure 3.9).

1091 In general, however, little statistical information is provided.

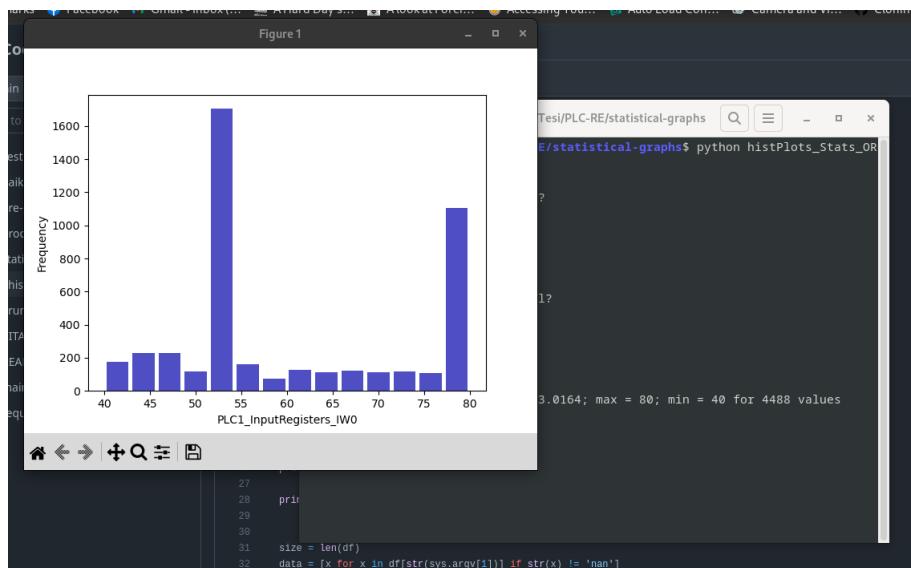


Figure 3.9: Histogram plot overshadowing statistical information shown on the terminal window in the background

1092 In general, however, little statistical information is provided.

1093 **Business Process Mining and Analysis** Concerning the data mining,
1094 this is a purely ad *hoc* solution, designed to work under special conditions:
1095 first, the timestamped dataset of the physical process and the one obtained
1096 after the packet sniffing operation of Modbus traffic on the network need
1097 to be synchronized and have the same granularity, in this case one event
1098 per second.

1099 It is relatively easy, therefore, to find correspondences between Modbus

1100 commands sent over the network and events occurring on the physical
1101 system, such as state changes in actuators, due in part to the fact that the
1102 number of communications over the network is really small (see Section
1103 3.2.1).

1104 ***First limitation*** In a real system, network communications are much more
1105 numerous and involve many more devices even in the same second:
1106 finding the exact correspondence with what is happening in the cy-
1107 ber physical system becomes much more difficult.

1108 Since this is, as mentioned, an *ad hoc* solution, only the Modbus pro-
1109 tocol is being considered: as widely used as this industrial protocol
1110 is, other protocols that are widely used [53] such as EtherNet/IP (see
1111 Section 2.2.6.2) or Profinet should be considered in order to extend
1112 the analysis to other industrial systems that use a different commu-
1113 nication network.

1114 ***Second limitation*** The other limiting aspect of the business process min-
1115 ing phase is the **process mining software** used to generate the ac-
1116 tivity diagram. As mentioned in Section 3.2.5, the process mining
1117 software used by Ceccato et al. is **Disco**: this is commercial soft-
1118 ware, with an academic license lasting only 30 days (although free of
1119 charge), released for Windows and MacOS operating systems only,
1120 which makes its use under Linux systems impossible except by us-
1121 ing emulation environments such as Wine.

1122 For what is my personal vision and training as a computer scientist,
1123 it would have been preferable to use a *cross-platform, freely licensed*
1124 *open source* software alternative to Disco: one such software could
1125 have been **ProM Tools** [54], a framework for process mining very
1126 similar to Disco in functionality, but fitting the criteria just described,
1127 or use Python libraries such as **PM4PY** [55], which offer ready-to-use
1128 algorithms suitable for various process mining needs.

52 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

1129 **Invariants Inference and Analysis** The limitation in this case is principally
1130 Daikon: this software is designed to compute the invariants of a software from its live execution or from a file containing its execution flow, not
1131 to find the invariants of a cyber physical system. Since there are currently
1132 no better consolidated alternatives for inferring invariants, however, an
1133 attempt was still made to use Daikon as best as possible.
1134

```
daikon_results_cond.txt
~/UniVr/Tesi/PLC-RE/daikon/Daikon_Invariants

Daikon version 5.8.14, released October 6, 2022; http://plse.cs.washington.edu/daikon.
Reading splitter info files
(read 1 spinfo file, 1 splitter)
Reading declaration files .aprogram.point:::POINT: 1 of 1 splitters successful

(read 1 decls file)
Processing trace data; reading 1 dtrace file:

Warning: No non-obvious non-suppressed exclusive invariants found in
aprogram.point:::POINT
Warning: No non-obvious non-suppressed exclusive invariants found in
aprogram.point:::POINT
=====
aprogram.point:::POINT
PLC2_MemoryRegisters_MW1 == PLC3_MemoryRegisters_MW1
PLC1_MemoryRegisters_MW0 == 40.0
PLC1_MemoryRegisters_MW1 == 80.0
PLC1_Coils_QX00 one of { 0.0, 1.0 }
PLC1_Coils_QX01 one of { 0.0, 1.0 }
PLC1_Coils_QX02 one of { 0.0, 1.0 }
PLC2_MemoryRegisters_MW1 == 10.0
PLC2_MemoryRegisters_MW2 == 20.0
PLC2_Coils_QX00 one of { 0.0, 1.0 }
PLC3_InputRegisters_IW0 >= 0.0
PLC3_Coils_QX00 one of { 0.0, 1.0 }
PLC3_Coils_QX02 one of { 0.0, 1.0 }
prev_PLC1_Coils_QX00 one of { 0.0, 1.0 }
prev_PLC1_Coils_QX01 one of { 0.0, 1.0 }
prev_PLC2_Coils_QX00 one of { 0.0, 1.0 }
prev_PLC3_InputRegisters_IW0 >= 0.0
prev_PLC3_Coils_QX00 one of { 0.0, 1.0 }
prev_PLC3_Coils_QX02 one of { 0.0, 1.0 }
PLC1_Max_safety == 77.0
```

Figure 3.10: Example of Daikon's output

1135 **First limitation** The biggest problem with Daikon applied to the computation of invariants of an industrial system is the difficult reading
1136 of the resulting output: the software in fact returns a very long list
1137 of invariants, one invariant per line, many of no use and without
1138 correlating invariants that may have common features or deriving
1139

1140 additional information from them. The process of screening and rec-
1141ognizing the significant invariants, as well as the correlation between
1142them, must be done by a human: certainly not an easy task given the
1143volume of invariants one could theoretically be faced with (hundreds
1144and hundreds of invariants). An example of Daikon's output can be
1145seen in Figure 3.10.

1146 **Second limitation** The bash script used in this phase of the analysis does
1147not help at all in deriving significant invariants more easily: it merely
1148launches Daikon and saves its output to a text file by simply redirect-
1149ing the stdout to file. No data reprocessing is done during this step.
1150 In addition, if a condition is to be specified to Daikon before perform-
1151ing the analysis, it is necessary each time to edit the .spinfo file by
1152manually entering the desired rule, an inconvenient operation when
1153multiple analyses are to be performed with different conditions each
1154time.

54 3.2 Ceccato et al.'s black-box dynamic analysis for water-tank systems

A Framework to Improve Ceccato et al.'s Work.

1155 IN CHAPTER 3, I presented the state of the art of *process comprehension* of
1156 an Industrial Control System (ICS) focusing later on the methodology
1157 proposed by Ceccato et al. [9][Section 3.2], explaining what it consists of,
1158 its practical application on a testbed, and most importantly highlighting
1159 its limitations and critical issues (see Section 3.2.7).

1160 In this chapter I will present **my proposals to improve the methodology**
1161 presented in the previous chapter, overcoming (or at least trying to do
1162 so) the criticalities mentioned above by almost completely rewriting the
1163 original framework, enhancing its functionalities and inserting new ones
1164 where possible, while keeping its general structure and approach: the sys-
1165 tem analysis will in fact consist of the same four steps as in the original
1166 methodology (Data Pre-processing, Graph and Statistical Analysis, Busi-
1167 ness Process Mining and Invariants Inference), but each of them will be
1168 deeply revised in order to provide a richer, clearer and more complete
1169 process comprehension of the industrial system to be analyzed and its be-
1170 havior.

1171 As it may have already been noted, my proposals do not involve im-
1172 proving the data gathering phase: this is due simply to the fact that the
1173 novel framework will not be tested on the same case study used by Cec-

1174 cato et, al. (Section 3.2.1), but on a different case study, the iTrust SWaT
1175 system [36], of which (some) datasets containing the execution trace of
1176 the physical system and the network traffic scan are already provided by
1177 iTrust itself. For more details about this case study, see Chapter 5.

1178 4.1 The novel Framework

1179 The implementation of the novel framework for ICSs analysis starts
1180 from several assumptions:

- 1181 1. it must be implemented in a **single programming language**
- 1182 2. it must be **independent of the system** to be analyzed
- 1183 3. It must provide greater **flexibility and ease of use** for the user at
1184 every stage

1185 In the following, these three points will be discussed in more detail.

1186 **Single Programming Language** The original tool was implemented us-
1187 ing various programming languages in each of the different phases:
1188 from Python up to Java, passing through Bash scripting.

1189 In my opinion, this heterogeneity makes it more difficult and less
1190 intuitive for the user to operate on the tool: moreover, the use of
1191 multiple technologies makes it more difficult to maintain the code
1192 and add new features, particularly if only a single person is manag-
1193 ing the code (he/she might be proficient in one language, but little
1194 of the others).

1195 For these reasons, I decided to use a single programming language,
1196 to ensure homogeneity to the framework and ease of use and main-
1197 tenance of the code for anyone who wants to manage it in the future:
1198 I chose to use Python, because of its simplicity and easy readability
1199 combined with its versatility and powerfulness: moreover, Python
1200 can count on a massive number of available libraries and packages
1201 that meet all kinds of needs.

1202 **System Independence** One of the biggest limitations of Ceccato et al.'s
1203 tool that I highlighted in Section 3.2.7 is the fact that it is **highly de-**
1204 **pendent on the testbed used:** that is, it is *not* possible to configure
1205 any of the tool's parameters to analyze different industrial systems.
1206 To overcome this issue and make my framework independent of the
1207 system to be analyzed, also eliminating all references to hardcoded
1208 variables and values present in the previous tool, I decided to use a
1209 **general configuration file**, named *config.ini*, in which the user can, at
1210 will, customize all the parameters necessary to perform the analysis
1211 of the targeted system.

1212 **Flexibility and Ease on Use** The lack of flexibility and ease of use in a tool
1213 can be a significant disadvantage, limiting its effectiveness and mak-
1214 ing it challenging for the user to get the desired outcomes. The orig-
1215 inal tool suffered from these limitations, with users having to run
1216 scripts from the command line, with little to no options or param-
1217 eters available to customize the analysis. As a result, the tool was
1218 not user-friendly and lacked the flexibility to adapt to specific user
1219 needs.

1220 To settle these issues, I enhanced the command-line interface in the
1221 novel framework by adding new options and parameters. These
1222 new features provide the user with greater flexibility, enabling to
1223 specify parameters and options that allow for more in-depth anal-
1224 ysis and focused results analyzing data more effectively and effi-
1225 ciently. With these enhancements, the framework has become more
1226 user-friendly, reducing the learning curve and making it more acces-
1227 sible to a wider range of users.

1228 This, in turn, makes the framework more valuable and useful, in-
1229 creasing its adoption and effectiveness across a range of industries
1230 and applications.

1231 Moreover, with new options and parameters users no longer have to
1232 rely solely on the command line interface, which can be challenging
1233 and intimidating for those with limited technical expertise. Instead,

1234 users can now access a range of customizable options and parameters,
 1235 making the tool more intuitive and user-friendly.

1236 Overall, the enhancements made to the framework represent a significant
 1237 step forward in making it more effective, efficient, and user-friendly.

1238 4.1.1 Framework Structure

1239 The structure of the novel framework mostly follows the structure of
 1240 the original tool: it is divided into four main directories each representing
 1241 the different phases of the analysis (data pre-processing, graphs and sta-
 1242 tistical analysis, process mining, and invariant analysis), and containing
 1243 the relevant Python scripts that perform the analysis, as well as subdirec-
 1244 tories and any input/output files necessary for the proper behavior of the
 1245 framework.

```

1246 .
1247    --- config.ini
1248    --- daikon
1249      --- Daikon_Invariants
1250      --- daikonAnalysis.py
1251      --- runDaikon.py
1252    --- network-analysis
1253      --- data
1254      --- export_pcap_data.py
1255      --- swat_csv_extractor.py
1256    --- pre-processing
1257      --- mergeDatasets.py
1258      --- system_info.py
1259    --- process-mining
1260      --- data
1261      --- process_mining.py
1262    --- statistical-graphs
1263      --- histPlots_Stats.py
1264      --- runChartSubPlots.py

```

Listing 4.1: Novel Framework structure and Python scripts

1265 Ahead of these directories there is the most important part, that allows
1266 the framework to be independent of the industrial control system being
1267 analyzed: the *config.ini* file. Here the user can configure general parame-
1268 ters and options, such as paths to read from or write files to, or related to
1269 individual analysis phases.

1270 The file is divided into sections, each covering a different aspect of the con-
1271 figuration: each section contains user-customizable parameters that will
1272 then be called within the Python scripts that constitute the framework.

1273 Sections of *config.ini* are:

- 1274 • [PATHS]: defines general paths such as the project root directory and
1275 some source directories for datasets
- 1276 • [PREPROC]: contains some parameters needed for the pre-processing
1277 phase
- 1278 • [DAIKON]: defines parameters needed for invariant analysis with
1279 Daikon
- 1280 • [DATASET]: defines settings and parameters used during the dataset
1281 enrichment stage and possibly in further phases
- 1282 • [MINING]: contains parameters used during the process mining phase
- 1283 • [NETWORK]: Contains specific settings for extracting the data ob-
1284 tained from the packet sniffing phase on the ICS network and con-
1285 verting it to CSV format. It also defines the network protocols that
1286 are to be analyzed

1287 4.1.2 Python Libraries and External Tools

1288 Since the framework has been entirely developed in Python, I have
1289 tried to make use of external tools as little as possible, with the idea of in-
1290 tegrating all the various functionalities within the framework and making
1291 it independent of further software: the only external tool remaining from
1292 the old Ceccato et al. tool is Daikon, precisely because there is currently

1293 no better alternative or Python packages that performs the same function-
1294 alities.

1295 Instead, large use of Python libraries is made for handling functionality
1296 and input data: the fundamental libraries upon which the framework is
1297 based are:

- 1298 • **Pandas**, also used in the previous tool for dataset management, but
1299 whose use here has been deepened and extended
- 1300 • **NumPy**, often used together with Pandas to perform some opera-
1301 tions to support it
- 1302 • **MatPlotLib**, for managing and plotting graphical analysis
- 1303 • other scientific libraries such as **SciPy**, **StatsModel** [56] and **Net-
1304 workX** [57], for mathematical, statistical and analysis operations on
1305 the data
- 1306 • **GraphViz**, for the creation of activity diagrams in the process mining
1307 phase

1308 Having now seen the structure of the framework, in the next sections we
1309 will go into more detail describing my proposals and what I have done to
1310 improve the various stages of the analysis.

1311 4.2 Analysis Phases

1312 4.2.1 Phase 1: Data Pre-processing

1313 *Data Pre-processing phase* is probably the most delicate and significant
1314 one: depending on how large the industrial system to be analyzed is, the
1315 data collected, and how it is enriched using the additional attributes, the
1316 subsequent system analysis will provide more or less accurate outcomes.

1317 The previous tool has many limitations, especially at this stage: it is not
1318 possible to isolate a subsystem (either on a temporal basis or on the num-
1319 ber of PLCs to be analyzed - the system is considered in its whole), and
1320 many of the additional attributes were actually added manually: more-
1321 over, for those automatically entered, there is no way to specify which
1322 register type to associate the additional attribute with.

1323 All this, combined with the fact that in the tool code many references
1324 to attributes and registers are hardcoded, makes the analysis of the sys-
1325 tem much more difficult and the obtained results less accurate in terms of
1326 quantity and quality.

1327 In the novel framework these problems have been overcome by intro-
1328 ducing the possibility, starting from the datasets of individual PLCs ob-
1329 tained from data gathering process, to select a subsystem from the com-
1330 mand line both on a temporal basis and of the PLCs to be considered;
1331 I have also redesigned the whole process of enrichment of the resulting
1332 dataset, eliminating the manual entry of additional attributes and giving
1333 the user the possibility to be able to decide which type of additional at-
1334 tribute to associate with a given register. In addition to this, at the end
1335 of the pre-processing operation, it is possible to perform a brief prelimi-
1336 nary analysis of the obtained dataset in order to estimate which registers
1337 are connected to actuators, which to measurements, and which represent
1338 hardcoded relative setpoints or constants: this operation also makes it pos-
1339 sible to be able to refine the enrichment step by setting the relevant param-
1340 eters in the *config.ini* file

1341
1342 In the next sections we will look in more detail at what has been accom-
1343 plished.

1344 **4.2.1.1 Subsystem Selection**

1345 In the previous tool, the datasets in CSV format referring to each single
1346 PLC are placed in a fixed directory (hardcoded in the script) from which
1347 the dedicated script later perform merge and enrichment of them all, re-
1348 sulting in a single dataset representing the entire process trace of the in-
1349 dustrial system as an output. As mentioned, the script makes no provision
1350 to choose the individual PLCs to be analyzed, nor to decide on a temporal
1351 range over which to perform the analysis: in fact, it may happen that dur-
1352 ing the period of scanning and data gathering there is a so-called *transient*,
1353 i.e., a general state in which the industrial system is still initializing before
1354 actually reaching full operation; or, more simply, there is the need to ana-
1355 lyze the process of only a specific part of the industrial system in a certain
1356 period of interest: whatever the motivation, the lack of elasticity and op-
1357 tions to provide to the user makes the analysis much more complex than
1358 it might be, affecting even the later phases, as the number of variables to
1359 be analyzed becomes enormously higher.
1360 In addition, it is not possible to specify an output CSV file where to save
1361 the resulting dataset: at each dataset creation and enrichment operation,
1362 therefore, the resultant file will be overwritten. This is very awkward
1363 when making comparisons between two different execution traces, for ex-
1364 ample, unless the files are renamed manually.

1365 Let's see how all these issues were solved in the novel framework I
1366 developed: first of all, in the general *config.ini* file there are some general
1367 default settings about paths, and among them the one concerning the di-
1368 rectory where to place the datasets of the individual PLCs to be processed.
1369 In addition to this option, there are other ones that define further aspects
1370 related to the operations performed in this phase. Listing 4.2 shows the
1371 settings in question:

```
1372 [PATHS]
1373 root_dir = /home/marcuzzo/UniVr/Tesi
1374 project_dir = %(root_dir)s/PLC-RE
1375 input_dataset_directory = %(root_dir)s/datasets_SWaT/2015
```

```

1376     net_csv_path = %(root_dir)s/datasets_SWaT/2015/Network_CSV
1377
1378 [DEFAULTS]
1379 dataset_file = PLC_SWaT_Dataset.csv # Default output
1380     ↪ dataset
1381 granularity = 10 # slope granularity
1382 number_of_rows = 20000 # Seconds to consider
1383 skip_rows = 100000 # Skip seconds from beginning

```

Listing 4.2: Paths and parameters for the Pre-processing phase in config.ini file

1384 Concurrently, the same options can be specified by the user via the
1385 command line of the new Python script (named `mergeDatasets.py` and
1386 contained in the directory `pre-processing` of the project) and will override
1387 the default ones found in `config.ini`. These options are:

- 1388 • **-s or --skiprows:** seconds to jump from the beginning of the file. This
1389 option is useful in case the system has an initial transient or to start
1390 the analysis from a certain point in the dataset
- 1391 • **-n or --nrows:** reference temporal period in seconds (rows) for the
1392 analysis from the beginning of the dataset or from the point specified
1393 in the `-s` option or in the corresponding setting in `config.ini`.
1394 This option makes a **selection** on the data of the dataset.
- 1395 • **-p or --plcs:** PLCs to be merged and enriched. The user can specify
1396 the desired PLCs by indicating the CSV file names of the associated
1397 datasets with no limitations on number.
1398 This option makes a **projection** on the data of the dataset.
- 1399 • **-d or --directory:** performs the merge and enrichment of all CSV files
1400 contained in the directory specified by user, overriding the default
1401 setting in `config.ini`. It is in fact the old functionality of the previous
1402 tool, maintained here to give the user more flexibility and conve-
1403 nience in case he wants to perform the analysis on the whole system.
1404 This is also the default behavior in case the `-p` option is not specified.

- 1405 • **-o or --output:** specifies the name of the file in which the obtained
 1406 dataset will be saved. It must necessarily be a file in CSV format.
- 1407 • **-g or --granularity:** specifies a granularity that will be used to calculate
 1408 the measurement slope during the dataset enrichment phase.
 1409 We will discuss this later in Section 4.2.1.2.

1410 **4.2.1.2 Dataset Enrichment**

1411 After a step in which a function is applied to each PLC-related dataset
 1412 that eliminates its registers that have not been used within the system (this
 1413 is especially true if the Modbus register scan has been performed, in which
 1414 ranges of registers are scanned: it is assumed that unused registers have
 1415 constant value zero), the **dataset enrichment operation** is performed.

1416 This operation differs from the previous version not only in the fact that
 1417 it is performed on each individual dataset and not on the resulting dataset,
 1418 but also in the additional attributes: not only are they greater in number,
 1419 but they are automatically calculated and inserted by the `mergeDatasets.py`
 1420 script into the dataset and, most importantly, it is possible to decide through
 1421 the parameters in the `config.ini` configuration file under the [DATASET] sec-
 1422 tion to which registers these attributes should be assigned.

1423 In Listing 4.3 we can see the list of additional attributes and how they
 1424 should be associated with the registers of the dataset:

```
1425      [DATASET]
1426      timestamp_col = Timestamp
1427      max_prefix = max_
1428      min_prefix = min_
1429      max_min_cols_list = lit|ait|dpit
1430      prev_cols_prefix = prev_
1431      prev_cols_list = mv[0-9]{3}|p[0-9]{3}
1432      trend_cols_prefix = trend_
1433      trend_cols_list = lit
1434      trend_period = 150
1435      slope_cols_prefix = slope_
```

1436 slope_cols_list = lit

Listing 4.3: config.ini parameters for dataset enriching

1437 Following is a brief explanation of the parameters just seen:

1438 **timestap_col** indicates the name of the column that contains the data
1439 timestamps. This parameter is used not only in this phase, but is
1440 also referred to in the Process Mining phase. In the previous work,
1441 this parameter was hardcoded and not configurable (and thus caus-
1442 ing errors if the system being analyzed changed)

1443 **max_prefix, min_prefix, max_min_cols_list** refer to any relative maximum
1444 or minimum values (*relative setpoints*) of one or more measures and
1445 that can be found and inserted as new columns within the dataset.
1446 The first two parameters indicate the prefix to be used in the column
1447 names affected by this additional attribute, while the third specifies
1448 of which type of registers we want to know the maximum and/or
1449 minimum value reached (several options can be specified using the
1450 logical operator | - or).

1451 If, for example, we want to know the maximum value of the regis-
1452 ters associated with the tanks, indicated in the iTrust SWaT system
1453 by the prefix LIT, we only need to specify the necessary parameter in
1454 the *config.ini* file, so `max_min_cols_list = lit`.

1455 The result will be to have in the dataset thus enriched a new column
1456 named `max_P1_LIT101`.

1457 **prev_cols_prefix, prev_cols_list** refer to the values at the previous step
1458 of the registers specified in `prev_cols_list`. It is possible to specify
1459 registers using *regex*, as in the example shown. It may be useful in
1460 some cases to have this value available to check, for example, when
1461 a change of state of a single given actuator occurs. The behavior of
1462 these parameters is the same as described in the point above.

1463 **slope_cols_prefix, slope_cols_list** are related to the calculation of the
1464 slope of a specific register (usually a measure), that is, its trend. The

1465 slope can be ascending (if its value is greater than zero), descending
 1466 (if less than zero) or stable (if approximately equal to zero). I will dis-
 1467 cuss the slope calculation in more detail in the next paragraph, as it
 1468 is related to the attributes `trend_cols_prefix`, `trend_cols_list` and
 1469 `trend_period`

1470 Initially, the parameters for registers to be associated with each addi-
 1471 tional attribute may be left blank, assuming that we do not know the sys-
 1472 tem at all and therefore do not know which registers may be actuators,
 1473 which measures, and which further. This information can be obtained
 1474 from the **brief analysis** following the datasets merging operation: this
 1475 analysis, performed at the user's choice, indicates which may be likely
 1476 sensors, which actuators, and further information of various kinds: these
 1477 indications allow the user to be able to set the desired values in *config.ini*
 1478 file and hence refine the enrichment process by re-launching the
 1479 `mergeDatasets.py` script again.

1480 **Slope Calculation** The *slope* is an attribute that indicates the trend of
 1481 the measurement we are considering and is useful, in our context, during
 1482 the inference and invariant analysis phase in order to derive information
 1483 about this trend given specific conditions: this trend can be, in general,
 1484 increasing ($slope > 0$), decreasing ($slope < 0$) or stable ($slope = 0$).
 1485 Normally, the slope is calculated through a simple mathematical formula:
 1486 given an interval a, b relative to the measurement l , the slope is given by
 1487 the difference of these two values divided by the amount of time t that the
 1488 measurement takes to reach b from a :

$$slope = \frac{l(b) - l(a)}{t(b) - t(a)}$$

1489 In the novel framework as in the old tool, this time interval (also called
 1490 **granularity**) can be either long or short, depending on the accuracy de-
 1491 sired on the slope: the lower the granularity, the more the slope will re-
 1492 flect the actual measurement trend; the higher the granularity, on the other
 1493 hand, the more the slope data will be flattened. Each time interval into

1494 which the measure is divided corresponds to a slope, the set of which in-
1495 serted as an additional attribute in the dataset will later be used to define
1496 the trend of the measure itself in specific situations.

1497 Calculating the slope directly from the raw measurement data may be
1498 an acceptable solution for those systems whose measurements are not sig-
1499 nificantly affected by **perturbations** (such as the oscillations of the liquid
1500 inside a tank during the filling and emptying phases, leading to fluctu-
1501 ating level readings): in this case, granularity can be kept low and thus
1502 obtain a very accurate overall trend calculation close to the actual mea-
1503 surement trend. This is the case, for example, with the tanks of the testbed
1504 used by Ceccato et al.

1505 However, in the case where these perturbations significantly afflict the de-
1506 tections on the measurement, the slope calculation on the individual time
1507 intervals of the measurement may lead to an erroneous result in trend def-
1508 inition, regardless of the granularity used.

1509 Figure 4.1 demonstrates this assertion: the measurement, in blue, refers
1510 to the LIT101 tank of the iTrust SWaT system; in red, the slope calculation
1511 related to the measurement with three different granularities: 30 (Figure
1512 4.1a), 60 (Figure 4.1b) and 120 seconds (Figure 4.1c). It can be seen that in
1513 addition to the flattening of slope values as the granularity increases, in
1514 the time interval between seconds 1800 and 4200 the level of LIT101 has
1515 a generally increasing trend, but the slope values vary from positive to
1516 negative: the result is that in the invariant analysis the general increasing
1517 trend will not be detected thus losing the information.

1518 The possibility of having (strongly) perturbed data was not considered
1519 in the previous tool, so I was faced with this issue to solve.

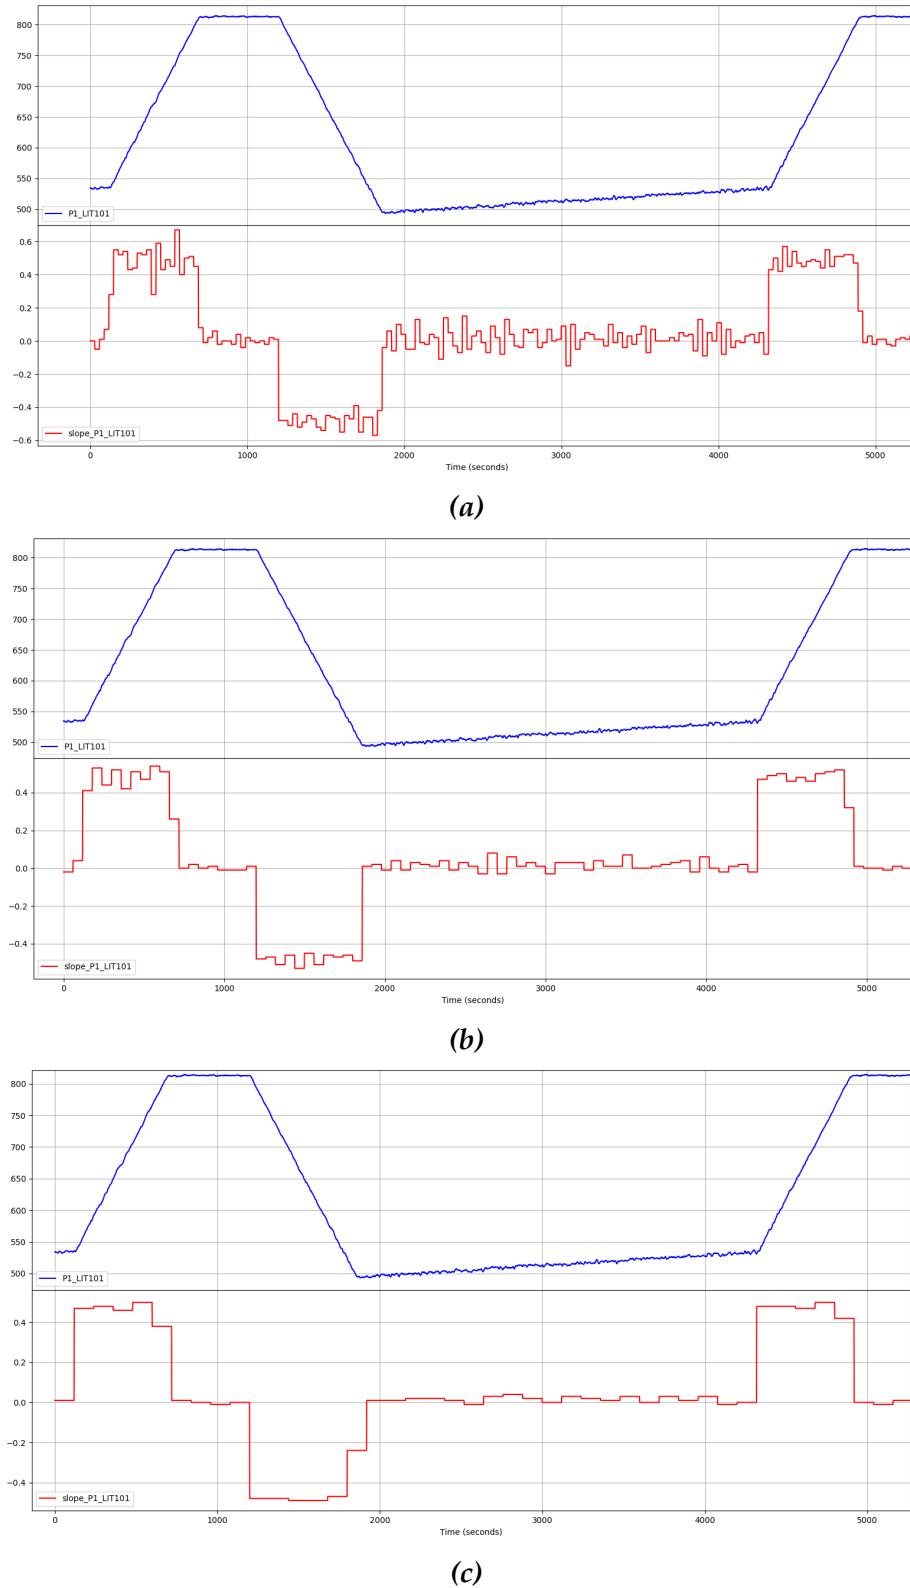


Figure 4.1: Slope comparison with granularity 30 (a), 60 (b) and 120 seconds (c)

1520 The solution to this problem is trying to remove as much "noise" as
 1521 possible from the data, in order to get a more linear trend in the curve rep-
 1522 resenting the measurement and consequently be able to calculate slopes
 1523 more accurately.

1524 There are several ways to smooth out the noise, but two of them seemed
 1525 to me to be the most suitable and that I considered and evaluated: using
 1526 **polynomial regression**, thus creating a filter on the noise, or a **seasonal**
 1527 **decomposition**, and more specifically the part concerning the **trending**.

1528 With regard to polynomial regression, I evaluated the *Savitzky-Golay* algo-
 1529 rithm [58], and with regard to seasonal decomposition, I evaluated *Seasonal-*
 1530 *Trend decomposition using LOESS (STL)* [59].

1531 For reasons of the length of this paper I cannot describe these two solutions
 1532 in detail (for this, I refer to the bibliographical notes): Figure 4.2, however,
 1533 shows a quick graphical comparison of them compared with the original
 1534 data. The solution chosen is the STL decomposition, which is more effec-
 1535 tive in attenuating noise than the Savitzky-Golay filter although at the cost
 1536 of more delay (still present in all algorithms of this kind) in some parts.

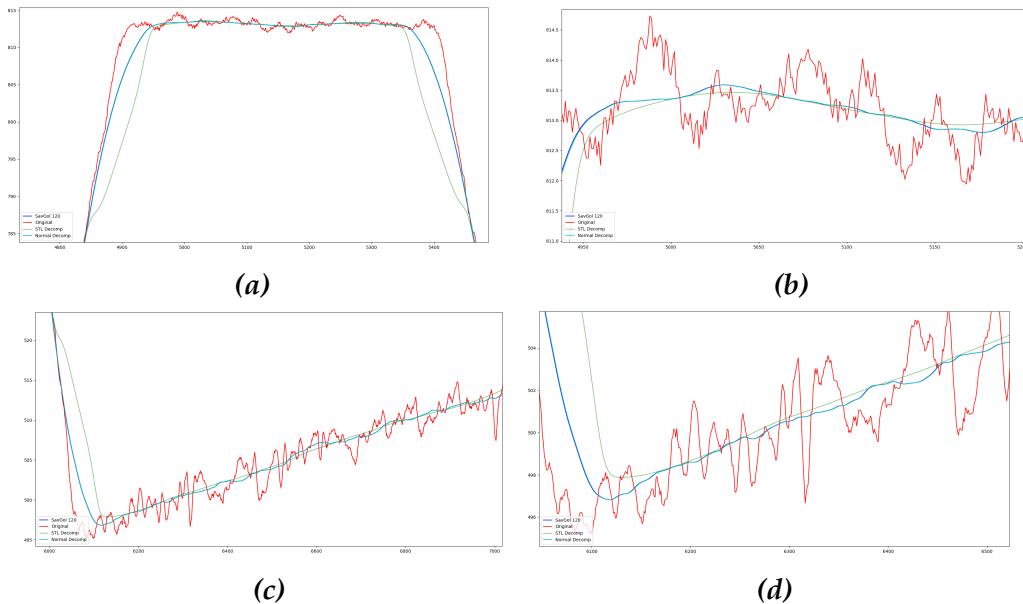


Figure 4.2: Savitzky-Golay filter (blue line) and STL decomposition (green) comparison

1537 The application of the STL decomposition results in a significant im-
 1538 provement in slope calculation even when using low granularity: in Fig-
 1539 ure 4.3 it can be seen that, with the same granularity used for the example
 1540 in Figure 4.1a, the slope values, although with unavoidable fluctuations,
 1541 remain within the same trend, corresponding to the trend of the data curve
 1542 net of the introduced lag caused by the periodicity set for the decomposi-
 1543 tion.

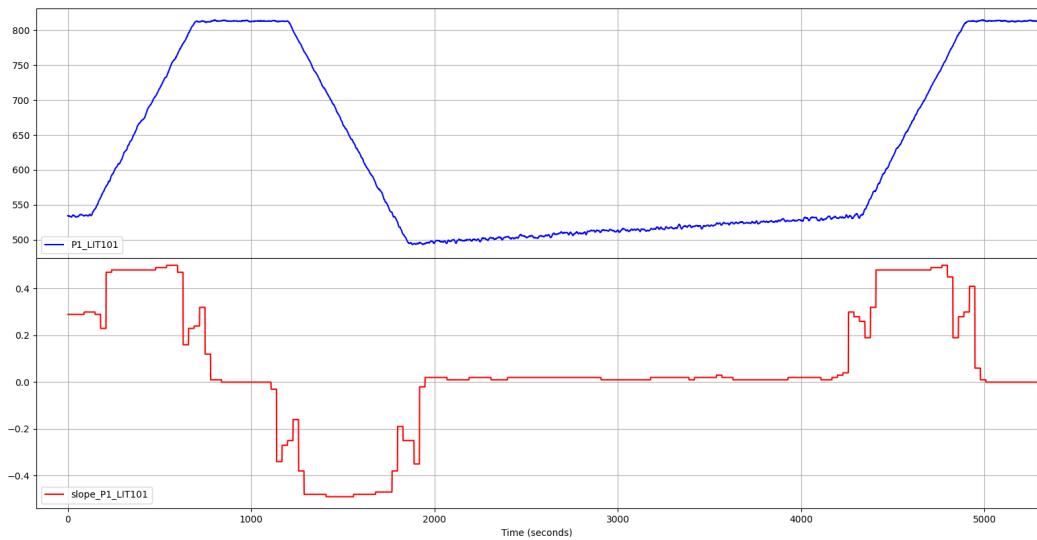


Figure 4.3: Slope after the application of the STL decomposition

1544 This periodicity, which indicates the sampling time window for de-
 1545 composition and thus the level of noise smoothing, can be set in the *con-*
 1546 *fig.ini* configuration file in the *trend_period* directive.

1547 The slope calculation will then be performed on the data from the addi-
 1548 tional measurement trend attributes specified in the *trend_cols_list* di-
 1549 rective in the configuration file, and no longer on the original unfiltered
 1550 data.

1551 Finally, to enable Daikon to correctly interpret the slope data, the deci-
 1552 mal values corresponding to each calculated slope are converted into three
 1553 numerical values -1, 0, and 1, which correspond to the decreasing (if the
 1554 slope is less than zero), stable (if it is equal to zero), and increasing (if it is

1555 greater than zero) trends, respectively. In Figure 4.4, the new slope can be
 1556 seen, along with the curve obtained from the STL decomposition:

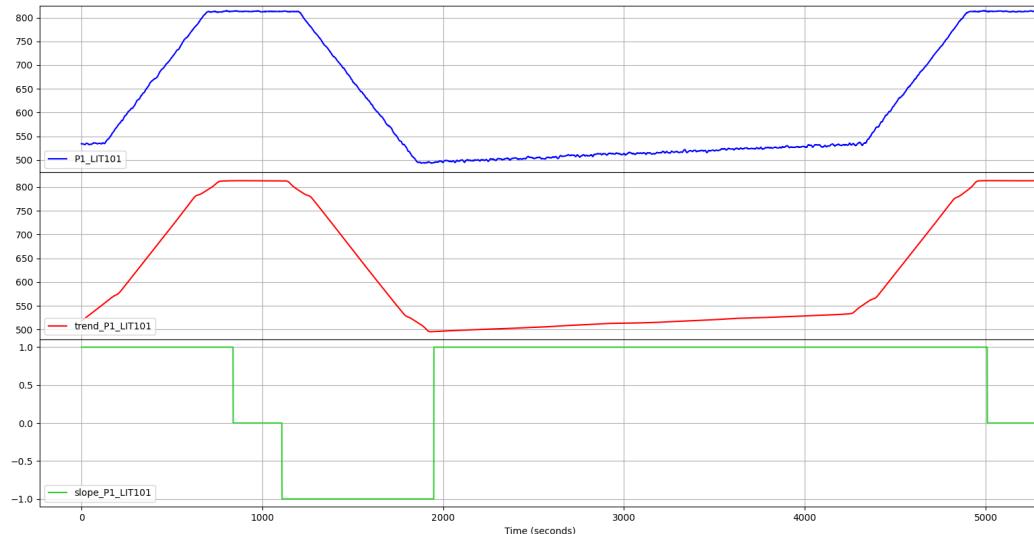


Figure 4.4: The new slope representation (green line) and the smoothed measurement data obtainind with the STL decompositon (red)

1557 **4.2.1.3 Datasets Merging**

1558 In this step the individual datasets are merged, obtaining two new dis-
 1559 tinct datasets: one without the enrichment and that will be use in the pro-
 1560 cess mining phase, and the other, containing the additional attributes but
 1561 with the timestamp column removed, intended for inference and invariant
 1562 analysis.

1563 The first of the two datasets obtained will be saved in CSV format by
 1564 default in the `$(project-dir)/process-mining/data` directory while the
 1565 second one in the `$(project-dir)/daikon/Daikon_Invariants` directory
 1566 (both paths are however configurable via `config.ini`).

1567 The dataset name is specified in `config.ini` or via the `-o` command-line op-
 1568 tion: in the case of the dataset intended for process mining, the script will
 1569 automatically add a `_TS` suffix to the filename, indicating the fact that the
 1570 dataset includes the timestamp.

1571 The opportunity for the user to specify a different filename for the output
 1572 each time allows the user to save the execution trace of the selected sub-
 1573 system without overwriting the previous ones and thus to use all of them
 1574 in the subsequent phases of the analysis.

1575 **4.2.1.4 Brief Analysis of the Obtained Subsystem**

1576 At the end of the datasets merging and saving step, the user is asked
 1577 whether to perform a brief optional analysis of the final resulting dataset
 1578 to extract preliminary data, with the purpose of obtaining basic informa-
 1579 tion about the (sub)system and possibly refining the enrichment.
 1580 If the user responds affirmatively to the request, `mergeDatasets.py` launches
 1581 within it a further Python script located in the same
 1582 `$(project-dir)/pre-processing` directory, called `system_info.py`. This
 1583 script, using a combination of Daikon and Pandas analysis, performs a
 1584 quick analysis of the dataset contents trying to **recognize**, however roughly,
 1585 **the register types**, with possible maximum and minimum values and hard-
 1586 coded setpoints. In addition, using the additional attribute `prev_`, it is ca-
 1587 pable of deriving measurement values in correspondence with state changes
 1588 of individual actuators.
 1589 Listing 4.4 shows an example of this brief analysis elated to PLC1 of the
 1590 iTrust SWaT system (for brevity, only one measurement is reported in the
 1591 analysis of actuator state changes):

```
1592     Do you want to perform a brief analysis of the dataset? [y
1593     ↵ /n]: y
1594
1595     Actuators:
1596     P1_MV101 [0.0, 1.0, 2.0]
1597     P1_P101 [1.0, 2.0]
1598
1599     Sensors:
1600     P1_FIT101 {'max_lvl': 2.7, 'min_lvl': 0.0}
1601     P1_LIT101 {'max_lvl': 815.1, 'min_lvl': 489.6}
1602
1603     Hardcoded setpoints or spare actuators:
```

```
1604 P1_P102 [1.0]

1605

1606 Actuator state changes:
1607     P1_LIT101   P1_MV101   prev_P1_MV101
1608     669       800.7170      0           2
1609     1850      499.0203      0           1
1610     4876      800.5992      0           2
1611     6052      498.9026      0           1
1612     9071      800.7170      0           2
1613     10260     499.1381      0           1
1614     13268     801.3058      0           2
1615     14435     498.4315      0           1
1616     17423     801.4628      0           2
1617     18603     498.1567      0           1

1618
1619     P1_LIT101   P1_MV101   prev_P1_MV101
1620     677       805.0741      1           0
1621     4885      805.7414      1           0
1622     9079      805.7806      1           0
1623     13276     805.1133      1           0
1624     17432      804.4068      1           0

1625
1626     P1_LIT101   P1_MV101   prev_P1_MV101
1627     1858      495.4483      2           0
1628     6060      497.9998      2           0
1629     10269     495.9586      2           0
1630     14443     495.8016      2           0
1631     18611     494.5847      2           0

1632
1633     P1_LIT101   P1_P101   prev_P1_P101
1634     118       536.0356      1           2
1635     4322      533.3272      1           2
1636     8537      542.1591      1           2
1637     12721     534.8581      1           2
1638     16883     540.5890      1           2

1639
1640     P1_LIT101   P1_P101   prev_P1_P101
1641     1190      813.0031      2           1
1642     5395      813.0031      2           1
```

```

1643      9597     811.8256      2      1
1644      13776    812.7283      2      1
1645      17938    813.3171      2      1
1646
1647      Actuator state durations:
1648      P1_MV101 == 0.0
1649      9   9   10   9   9   10   9   9   10   9
1650
1651      P1_MV101 == 1.0
1652      1174   1168   1182   1160   1172
1653
1654      P1_MV101 == 2.0
1655      669   3019   3012   3000   2981
1656
1657      P1_P101 == 1.0
1658      1073   1074   1061   1056   1056
1659
1660      P1_P101 == 2.0
1661      118   3133   3143   3125   3108

```

Listing 4.4: Example of preliminar system analysis

1662 From these results we can see that:

- 1663 • the probable actuators are P1_MV101, which assumes three states iden-
1664 tified by the values 0, 1 and 2, and P1_P101, which instead assumes
1665 two states identified by the values 1 and 2
- 1666 • there are two probable measures: P1_FIT101 whose values range
1667 from 2.7 to 0.0, and P1_LIT101 whose values range from 815.1 to
1668 489.6. One conjecture could already be made about the topology of
1669 the system: P1_LIT101 represents a tank
- 1670 • apparently there are no related *hardcoded setpoints*, but a probable
1671 spare actuator, P1_P102, whose value is always 1. From this data,
1672 another conjecture can be made: the value 1 is the close state for that
1673 particular type of actuators, while 2 represents the open state
- 1674 • from the analysis of state changes, in summary, we derive some *rela-*
1675 *tive setpoints*: for example, we know that P1_P101 changes state from

1676 value 1 to value 2 when the level of P1_LIT101 is about 813, while it
1677 changes from value 2 to 1 when the level of P1_LIT101 is about 535.
1678 We can deduce that P1_P101 is responsible for emptying the tank

- 1679 • as the last information we have duration of actuator states for each
1680 cycle of the system: this information can be useful for making as-
1681 sumptions and conjectures about the behavior of an actuator in a
1682 specific state or, by observing the duration values of each cycle, high-
1683 lighting anomalies in the system.

1684 In the specific case, the very short duration of P1_MV101 in state 0 is
1685 observable: since we are unable, at this preliminary stage, to make
1686 assumptions about this, we will try to understand the behavior of
1687 the system in this actuator state in the next stages of the analysis

1688 The information obtained here can be used, as mentioned above, to
1689 refine the enrichment of the dataset by setting directives in the [DATASET]
1690 section of the *config.ini* file, should this be empty or only partially set, or to
1691 make the first conjectures about the system, as we have just seen.

1692 The *system_info.py* file can also run in standalone mode if needed:
1693 it takes as command-line arguments the dataset to be analyzed, a list of
1694 actuators, and a list of sensors. For analysis related to state changes, the
1695 dataset must mandatorily be of the enriched type.

1696 4.2.2 Phase 2: Graphs and Statistical Analysis

1697 The new *graph analysis* arises from the need to give the user an overview
1698 of the (sub)system obtained in the previous pre-processing phase, identi-
1699 fying more easily the typology of the registers and grasping more effec-
1700 tively the relationships and the dynamics that may exist between the reg-
1701 isters controlled by one or more PLCs, confirming the initial conjectures if
1702 the brief analysis described in the previous section has been performed, or
1703 making new ones thanks to the visual graph support.

1704 In the previous tool, as already pointed out in Section 3.2.7, it is pos-
 1705 sible to view the chart of one only register at a time: this certainly makes
 1706 it possible to identify, or at least to hypothesize, the type of that register,
 1707 but it makes it very complicated to be able to relate it to the other compo-
 1708 nents of the system and thus to derive conjectures about the behavior of
 1709 the latter, conjectures to be possibly confirmed in the later phases. Hence
 1710 the need to create a tool that was better than the previous one and that
 1711 provided more information in an easier way.

1712 The first thing I thought was that an approach similar to Figure 3.5,
 1713 with all the graphs contained within a single plot, might be the most suit-
 1714 able one: unfortunately, I quickly realized that this solution presented sev-
 1715 eral issues, which are highlighted in figure 4.5.

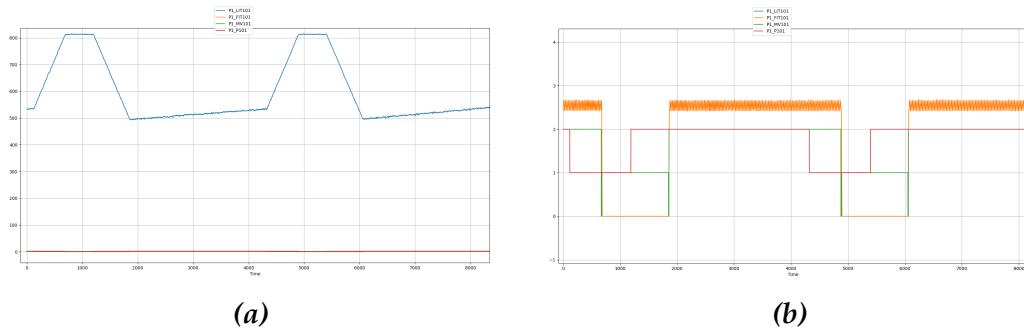
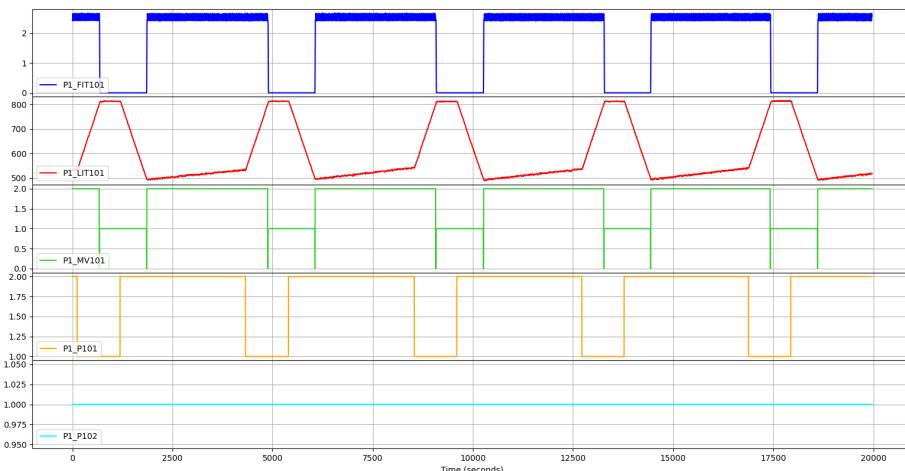


Figure 4.5: Plotting registers on the same y-axis

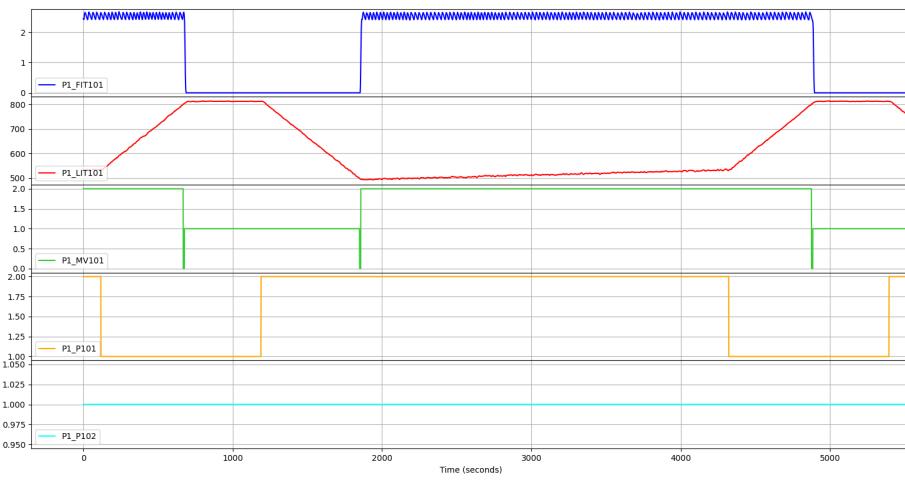
1716 In figure 4.5a it can be seen that the biggest issue is the use of the same
 1717 y-axis, related to the values of the individual registers, for all charts: if the
 1718 range between register values is wide, the risk of some charts resembling a
 1719 single flat line, or at any rate being in fact indistinguishable and very diffi-
 1720 cult to read immediately; furthermore, as shown in Figure 4.5b, if registers
 1721 have similar values to each other, the respective graphs may be confusing
 1722 to each other, making them more difficult to understand.

1723 The solution to this issue is simple and effective: the use of **subplots**.
 1724 Basically, each register corresponds to a subplot of the graph that shares
 1725 the time axis (the x-axis) with the other subplots, but keeps the y-axis of

the values of each register independent. This maintains the readability and comprehensibility of the charts, while simultaneously being able to immediately grasp the relationships between them. In addition, by sharing the time axis, it is possible to zoom in on a particular area of one of the charts and automatically the other ones will be zoomed in as well, thus not losing any information and no connection between registers. Figure 4.6 illustrates more clearly what has just been explained: the charts refer to the PLC1 registers of the iTrust SWaT system.



(a) Example of plotting charts of a PLC registers using subplots



(b) Zooming on a particular zone of the charts

Figure 4.6: Example of the new graph analysis

1734 To demonstrate the behavior and effectiveness of the new graph anal-
1735 ysis process, we observe in particular, Figure 4.6b: we can already have
1736 some validation on the conjectures made in the brief analysis from the
1737 previous step:

- 1738 • the chart of P1_FIT101 seems to confirm that this register can be as-
1739 sociated with a **measurement**; moreover, we can see that the trend
1740 of this register is closely related to the periods in which P1_LIT101
1741 has an increasing trend and to the evolution of P1_MV101. Its values,
1742 between about 0 and 2.5, are too narrow to say that this register rep-
1743 presents the tank, and the general trend also leads in that direction:
1744 we can therefore assume that it is a **pressure or flow sensor**
- 1745 • P1_LIT101, because of the above, would therefore seem to **represent**
1746 **a tank**, as already assumed in the brief analysis
- 1747 • P1_MV101 assumes value 2 at the beginning of the ascending trend of
1748 P1_LIT101, while it has value 1 in correspondence with the *plateau*
1749 that is seen when the measurement is about 800 and when the trend
1750 of it is descending. Thus, we can have a confirmation of the ini-
1751 tial conjecture: 1 represents the OFF state of the sensor, 2 the ON
1752 state, and P1_MV101 is the **actuator responsible for the rising level**
1753 of P1_LIT101
- 1754 • in the brief analysis we observed the short duration of P1_MV101 in
1755 state 0, but we were not able to speculate on the reason for this.
1756 Graph analysis shows that P1_MV101 switches and stays in the 0
1757 state acting as a kind of "transient" between states 1 and 2. It can
1758 be assumed that that is the period of time it actually takes for the
1759 actuator to change state
- 1760 • P1_P101 assumes value 2 at the beginning of the descending trend of
1761 P1_LIT101, after the *plateau*, and returns to value 1 at the change of
1762 slope of the (likely) tank increase: taken by itself this fact is not very
1763 clear, so it needs to be compared with P1_LIT101 and P1_MV101 to un-
1764 derstand its behavior: it can be seen that when P1_101 and P1_MV101

1765 are in state 1 the water level remains stable, whereas, when P1_P101
1766 changes to state 2 the level of the measurement drops rapidly; when
1767 P1_MV101 then also returns to state 2, the level of the measurement
1768 slowly starts rising again and then has a sudden rise at the change
1769 of state of P1_P101 from 2 to 1. We can therefore infer that P1_P101 is
1770 the **actuator responsible for emptying** the (presumed) tank: state 1
1771 represents the OFF state, while state 2 represents the ON state

- 1772 • P1_P102 seems to play no role in the system, since its state remains at
1773 1 all the time. It seems improbable that it could be a relative setpoint,
1774 so it can be assumed, according to the brief analysis, that it may be a
1775 **spare actuator**

1776 As you have probably already noticed, the vast majority of the graphs
1777 shown in the previous sections and chapters were generated precisely us-
1778 ing the new graph analysis script.

1779 The script that performs the new graph analysis is `runChartsSubPlots.py`,
1780 located in the `$(project-dir)/statistical-graphs` directory, which is writ-
1781 ten in Python and makes use, like the previous one, of the `matplotlib` li-
1782 braries to generate the graph plots.

1783 The script accepts the following command-line parameters:

- 1784 • **-f or --filename:** specifies the dataset, in CSV format, from which
1785 to read the data. The dataset must be located within the directory
1786 containing the enriched datasets for the invariant analysis phase. If
1787 subsequent parameters are not specified, the script will show all reg-
1788 isters in the dataset, except for additional attributes
- 1789 • **-r or --registers:** specifies one or more specific registers to be dis-
1790 played
- 1791 • **-a or --addregisters:** adds one or more registers to the default visual-
1792 ization. This option is useful in case an additional attribute such as
1793 slope is to be analyzed

- 1794 • **-e or --excluderegisters:** excludes one or more specific registers from
1795 the default visualization. This option is useful to avoid displaying
1796 hardcoded setpoints or spare registers

1797 This script, like the previous ones, is designed to provide the maximum
1798 flexibility and ease of use for the user, combined with greater power and
1799 effectiveness in deriving useful information about the analyzed system.

1800 **Statistical Analysis** A quick mention of the statistical analysis part: after
1801 careful evaluation, I decided **not to include this aspect** of the previous
1802 tool within the framework, because I saw no real practical use for it. The
1803 actual statistical information can be easily integrated into the brief analysis
1804 immediately following the pre-processing phase, while the histogram has
1805 relative utility and, in my opinion, is outdated by the new graph analysis
1806 I implemented.

1807 However, the Python `histPlots_Stats.py` script remains in the directory,
1808 essentially unchanged from the one developed by Ceccato et al., in case
1809 future use is needed.

1810 4.2.3 Phase 3: Invariant Inference and Analysis

1811 The phase of invariant inference and analysis has undergone a redesign
1812 and improvement to offer the user a more comprehensive and effortless
1813 approach to identify invariants. This has been achieved through the ap-
1814 plication of new criteria to analyze and reorganize the Daikon analysis
1815 results. The outcome of this is a more compact presentation of informa-
1816 tion that highlights the possible relationships among invariants.

1817 The new design not only enables the identification of undiscovered aspects
1818 of the system behavior but also confirms the hypotheses made during the
1819 earlier stages of analysis. This step is now semi-automated, unlike before,
1820 and allows for the analysis of invariants on individual actuator states and
1821 their combinations.

4.2.3.1 Revised Daikon Output

To streamline the process of identifying invariants quickly and efficiently, it is necessary to revise the output generated by standard Daikon analysis. The goal is to create a more compact and readable format for the output.

The current Daikon results consist of three sections:

1. the first section containing generic invariants, i.e., valid regardless of whether a condition is specified for the analysis
2. the second section containing invariants obtained by specifying a condition for the analysis in the *.spinfo* file
3. a third section containing the invariants that are obtained from the negations of the condition possibly specified in the *.spinfo* file

In each section only a single invariant per row is shown, without relating it in any way to the others: this makes it difficult to identify significant invariants and any invariant chains that might provide much more information about the behavior of the system than the single invariant.

A brief example of the structure and format of this output related to PLC1 of the iTrust SWaT system is shown in Listing 4.5, where a condition was specified on the measurement P1_LIT101 and on actuator P1_MV101:

```
aprogram.point:::POINT
P1_P102 == prev_P1_P102
P1_FIT101 >= 0.0
P1_MV101 one of { 0.0, 1.0, 2.0 }
P1_P101 one of { 1.0, 2.0 }
P1_P102 == 1.0
max_P1_LIT101 == 816.0
min_P1_LIT101 == 489.0
slope_P1_LIT101 one of { -1.0, 0.0, 1.0 }
[...]
P1_LIT101 > P1_MV101
P1_LIT101 > P1_P101
```

```

1854      P1_LIT101 > P1_P102
1855      P1_LIT101 < max_P1_LIT101
1856      P1_LIT101 > min_P1_LIT101
1857      [...]
1858      P1_MV101 < min_P1_LIT101
1859      P1_MV101 < trend_P1_LIT101
1860      P1_P101 >= P1_P102
1861      P1_P101 < max_P1_LIT101
1862      [...]
1863      =====
1864     aprogram.point:::POINT;condition="P1_MV101 == 2.0 &&
1865      ↪ P1_LIT101 < max_P1_LIT101 - 16 && P1_LIT101 >
1866      ↪ min_P1_LIT101 + 15"
1867      P1_MV101 == prev_P1_MV101
1868      P1_P102 == slope_P1_LIT101
1869      P1_MV101 == 2.0
1870      P1_FIT101 > P1_MV101
1871      P1_FIT101 > P1_P101
1872      P1_FIT101 > P1_P102
1873      P1_FIT101 > prev_P1_P101
1874      P1_MV101 >= P1_P101
1875      P1_MV101 >= prev_P1_P101
1876      P1_P101 <= prev_P1_P101
1877      =====
1878     aprogram.point:::POINT;condition="not(P1_MV101 == 2.0 &&
1879      ↪ P1_LIT101 < max_P1_LIT101 - 16 && P1_LIT101 >
1880      ↪ min_P1_LIT101 + 15)"
1881      P1_P101 >= prev_P1_P101
1882      Exiting Daikon.

```

Listing 4.5: Standard Daikon output for PLC1 of the iTrust SWaT system

1883 In the framework I am presenting, this output is reduced to two sec-
 1884 tions, the general one and the elated to the user-specified condition, elim-
 1885 inating the one related to the negated condition because it is not useful in
 1886 this context and a source of potential misinterpretation; moreover, the re-
 1887 lationships between invariants will be highlighted instead by using **transi-**
 1888 **tive closures:** transitive closure of a relation R is another relation, typically
 1889 denoted R^+ that adds to R all those elements that, while not necessarily

¹⁸⁹⁰ related directly to each other, can be reached by a *chain* of elements related
¹⁸⁹¹ to each other. In other words, the transitive closure of R is the smallest (in
¹⁸⁹² set theory sense) transitive relation R such that $R \subset R^+$ [60].

To implement the transitive closure of the invariants generated by Daikon, I first divided the invariants in each section by type according to their equivalence relation, not considering all those invariants that refer to additional attributes except for the slope, and for each of them I generated a **graph** using the NetworkX libraries, connecting with an arc the registers (represented by nodes) that have a common endpoint in the invariant through the transitive property. To reconstruct the individual invariant chains, then, I used a simple *Depth-first Search* (DFS) on each of the graphs, thus obtaining the desired outcome. Figure 4.7 shows some examples of these graphs:

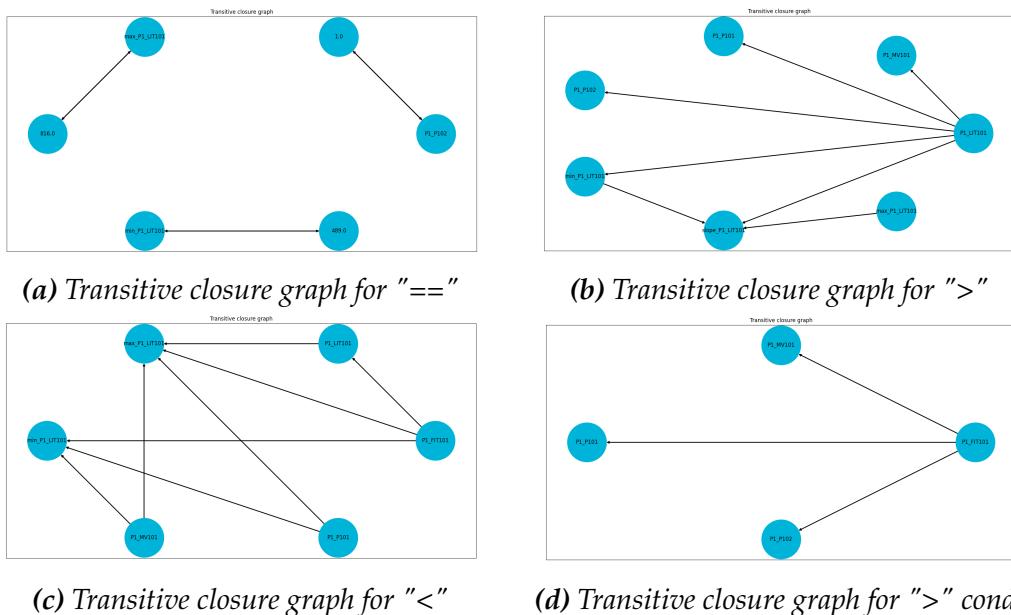


Figure 4.7: Example of transitive closure graphs for invariants in PLC1 of iTrust SWaT system

At the end of this process, still applied to PLC1 of the iTrust SWaT system and with the same analysis condition, we get the following complete output:

```

1906   1 =====
1907   2 Generic
1908   3 =====
1909   4 P1_MV101 one of { 0.0, 1.0, 2.0 }
1910   5 P1_P101 one of { 1.0, 2.0 }
1911   6 slope_P1_LIT101 one of { -1.0, 0.0, 1.0 }
1912   7 P1_FIT101 != P1_P101, P1_P102
1913   8 P1_P102 == 1.0
1914   9 max_P1_LIT101 == 816.0
1915  10 min_P1_LIT101 == 489.0
1916  11 P1_LIT101 > P1_MV101
1917  12 P1_LIT101 > P1_P101
1918  13 P1_LIT101 > P1_P102
1919  14 P1_LIT101 > min_P1_LIT101 > slope_P1_LIT101
1920  15 P1_FIT101 >= 0.0
1921  16 P1_P101 >= P1_P102 >= slope_P1_LIT101
1922  17 =====
1923  18 =====
1924  19 P1_MV101 == 2.0 && P1_LIT101 < max_P1_LIT101 - 16 &&
1925  20 ↪ P1_LIT101 > min_P1_LIT101 + 15
1926  21 =====
1927  22 slope_P1_LIT101 == P1_P102
1928  23 P1_MV101 == 2.0
1929  24 P1_FIT101 > P1_MV101
1930  25 P1_FIT101 > P1_P101
1931  26 P1_FIT101 > P1_P102
1932  27 P1_MV101 >= P1_P101

```

Listing 4.6: Revised Daikon output with transitive closures for PLC1 of the iTrust SWaT system

Transitive closures can be appreciated in lines 7, 14 and 16 of Listing 4.6. In general, the output has been reduced in the number of effective rows (19 versus the 61 in the output of Listing 4.5, making it certainly better to read and identify significant invariants) and the invariant chains make it more immediate to grasp the relationships between registers.

1938 **4.2.3.2 Types of Analysis**

1939 Compared to Ceccato et al.'s solution, in which individual analyses
1940 are performed manually, my proposal is to implement two types of semi-
1941 automated analysis: the first performs an analysis of all states for each
1942 individual register, while the second performs the analysis on the current
1943 system configuration based on the actual states of the actuators.
1944 Both analyses refer to a specific measurement, selected by the user.

1945 These two types of analysis will be handled by the Python script
1946 `daikonAnalysis.py`, contained in the default directory `$(project_dir)/daikon`.
1947 The script accepts the following command-line arguments:

- 1948 • **-f or --filename**: specifies the enriched dataset, in CSV format, from
1949 which to read the data. The dataset must be located within the di-
1950 rectory containing the enriched datasets specified in the `config.ini` file
1951 (by default `$(project_dir)/daikon/Daikon_Invariants`).
- 1952 • **-s or --simpleanalysis**: performs the analysis on the states of indi-
1953 vidual actuators
- 1954 • **-c or --customanalysis**: performs the analysis on combinations of ac-
1955 tual states of the actuators
- 1956 • **-u or --uppermargin**: defines a percentage margin on the maximum
1957 value of the measurement
- 1958 • **-l or --lowermargin**: defines a percentage margin on the minimum
1959 value of the measurement

1960 One or both types of analysis can be selected. The last two parameters set
1961 a condition on the value of the measurement that is meant to bypass the
1962 transient periods between the actuator state change and the actual trend
1963 change at the maximum and minimum values: this expedient is especially
1964 useful for the first type of analysis, allowing for generally more accurate
1965 data on measurement trends.

1966 **Analysis on single actuator states** Analysis on the states of individual
 1967 actuators is the simplest: after the user is prompted to input the measure-
 1968 ment, chosen from a list of likely available measurements, the script rec-
 1969ognizes the likely actuators and the relative states of each, using the same
 1970 mixed Daikon/Pandas technique adopted in the brief analysis during the
 1971 pre-processing phase.
 1972 For each actuator and each state it assumes, a single Daikon analysis is
 1973 performed, eventually placing the condition on the maximum and mini-
 1974 mum level of the measurement.
 1975 The result of these analyses are saved in the form of text files in a directory
 1976 having the name corresponding to the analyzed actuator and contained in
 1977 the default parent directory `$(project_dir)/daikon/Daikon_Invariants/results`:
 1978 each file generated by the analysis is identified by the name of the actuator,
 1979 the state and the condition, if any, on the measurement (see Figure 4.8).

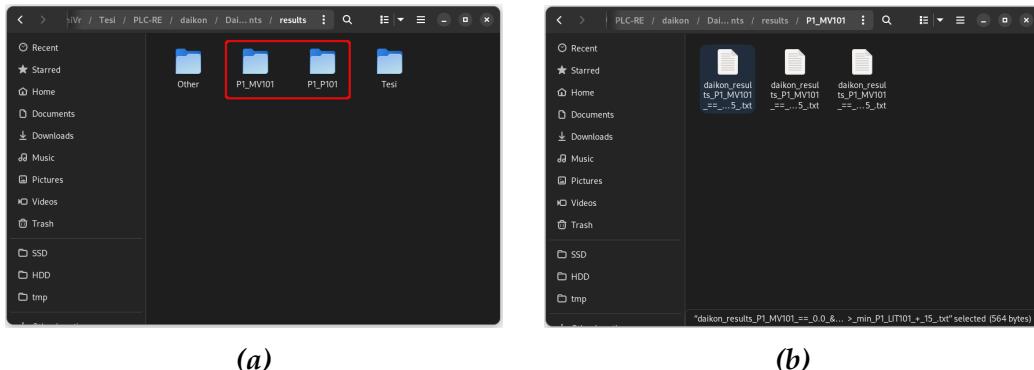


Figure 4.8: Directory (a) and outcome files (b) for the single actuator states analysis

1980 An example of the outcome of this analysis is Listing 4.6, where the
 1981 actuator analyzed is P1_MV101 in state 2: from the generic invariants we
 1982 can observe that P1_P101 is a likely actuator (line 5), assuming binary val-
 1983 ues; from line 8, however, we note that P1_P102 is permanently equal to 1,
 1984 so it can be confirmed that it is a spare actuator rather than a setpoint;
 1985 furthermore, lines 9 and 10 show the maximum and minimum values
 1986 reached by P1_LIT101, which we have assumed to be the register con-
 1987 nected to the tank. The most interesting information, however, comes

from the condition-generated invariants: at line 21 we notice that the slope of P1_LIT101 is 1 (thus increasing) when P1_MV101 is set to 2: this confirms what we assumed in the previous steps, namely, that P1_MV101 represents the ON state of the actuator and is responsible for filling the tank P1_LIT101; moreover, at line 23 we see that P1_FIT101 is greater than P1_MV101: this means that when the actuator assumes the value 2 the sensor P1_FIT101 detects data (in contrast, if P1_MV101 is 1 P1_FIT101 is 0, detecting nothing), confirming the original hypothesis made that this may be a pressure or flow sensor.

Analysis of the Current System Configuration The analysis of the current system configuration based on the actual states of the actuators is more complex, but at the same time offers more interesting outcomes as it provides better evidence of actuator behavior in relation to the selected measurement.

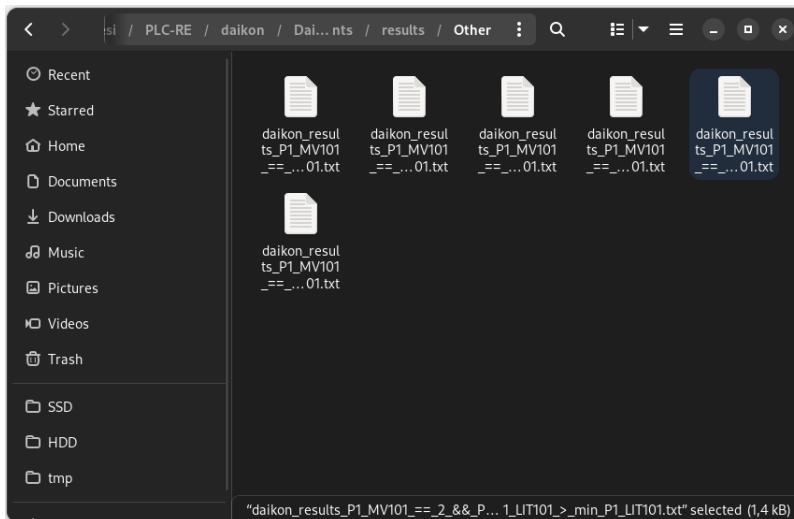


Figure 4.9: Daikon outcome files for system configuration analysis. Each file represents a single system state

For this analysis, the script automatically recognizes the actual configurations of the actuators (e.g., P1_MV101 == 2, P1_P101 == 1) that represent a system state, and for each of these configurations Daikon analysis is automatically run, after prompting the user for the measurement and

2006 eventually the actuators whose configurations are to be studied (If the user
 2007 does not select any actuators, all those previously detected by the Python
 2008 script will be considered): the analysis outcomes are saved, still in the text
 2009 format, within a specific directory under the same parent directory of the
 2010 previous type of analysis and, as before, their name is characterized by the
 2011 rule used for the analysis (see Figure 4.9).

2012

2013 An example of the obtained outcomes can be seen in Listing 4.7:

```

2014 1 =====
2015 2 Generic
2016 3 =====
2017 4 P1_MV101 <= P1_P101 ==> P1_FIT101 >= 0.0
2018 5 P1_MV101 <= P1_P101 ==> P1_MV101 one of { 0.0, 1.0, 2.0
2019 6   ↪ }
2020 7 P1_MV101 <= P1_P101 ==> P1_P101 one of { 1.0, 2.0 }
2021 8 P1_MV101 <= P1_P101 ==> slope_P1_LIT101 one of { -1.0,
2022 9   ↪ 0.0, 1.0 }
2023 10 P1_MV101 > P1_P101 ==> P1_FIT101 > P1_MV101
2024 11 P1_MV101 > P1_P101 ==> P1_FIT101 > P1_P101
2025 12 P1_MV101 > P1_P101 ==> P1_FIT101 > P1_P102
2026 13 P1_MV101 > P1_P101 ==> P1_FIT101 > slope_P1_LIT101
2027 14 P1_MV101 > P1_P101 ==> P1_MV101 == 2.0
2028 15 P1_MV101 > P1_P101 ==> P1_MV101 > P1_P102
2029 16 P1_MV101 > P1_P101 ==> P1_MV101 > slope_P1_LIT101
2030 17 P1_MV101 > P1_P101 ==> P1_P101 == 1.0
2031 18 P1_MV101 > P1_P101 ==> P1_P101 == P1_P102
2032 19 P1_MV101 > P1_P101 ==> P1_P101 == slope_P1_LIT101
2033 20 P1_MV101 > P1_P101 ==> slope_P1_LIT101 == 1.0
2034 21 [...]
2035 22 =====
2036 23 P1_MV101 == 2 && P1_P101 == 1 && P1_LIT101 < max_P1_LIT101
2037 24   ↪ && P1_LIT101 > min_P1_LIT101
2038 25 =====
2039 26 slope_P1_LIT101 == P1_P102 == P1_P101 == 1.0
2040 27 P1_MV101 == 2.0
2041 28 P1_FIT101 > P1_MV101
2042 29

```

2043 27 P1_FIT101 > P1_P101

Listing 4.7: Daikon outcomes for the system configuration P1_MV101 == 2, P1_P101 == 1 on P1_LIT101

2044 Compared to Listing 4.6, this time we can observe in the general invari-
2045 ants section the presence of implications, which were previously absent
2046 (the remaining generic invariants have been omitted for reasons of space).
2047 Such implications can provide very useful information, as in this case: for
2048 example, the invariant on line 18 tells us that if the value of P1_MV101 is
2049 greater than that of P1_P101 then the value of the slope is 1, that is, the
2050 tank level is increasing. Consequently, we have further confirmation that
2051 the state P1_MV101 == 2 is the ON state for that actuator and that it is the
2052 actuator responsible for filling the tank P1_LIT101. Comparing the other
2053 analysis outcomes, we will discover that P1_P101 is instead responsible
2054 for emptying the tank and that its ON and OFF states are 2 and 1, respec-
2055 tively.

2056 In addition, the invariant at line 8 also indicates that when the actuator
2057 P1_MV101 is in state 2 and P1_P101 is in state 1 then the value of P1_FIT101
2058 is greater than 2: consequently, it follows that the corresponding sensor is
2059 measuring something relative to the tank P1_LIT101.

2060 The above is confirmed by the invariants related to the analysis condition:
2061 at line 24 we have that the slope is indeed equal to 1 (thus increasing) and
2062 that P1_FIT101, at line 26, takes values greater than 2 when P1_MV101 is
2063 equal to 2.

2064 **Refining the Analysis** In some cases, it may happen that the outcomes
2065 provided by the semi-automated analyses are not satisfactory to the user
2066 (e.g., the value of the slope may not emerge clearly) or simply the user
2067 wants to investigate a particular aspect of the system in more detail by
2068 trying to discover additional invariants that did not emerge previously: in
2069 this case, it is possible to run a new and more specific invariant analysis
2070 using the Python script `runDaikon.py`, which allows for more punctual
2071 analyses of the system.

2072 The script, also contained in the default directory \$(project_dir)/daikon,
2073 accepts three command-line parameters:

2074 • **-f or --filename**: specifies the enriched dataset, in CSV format, from
2075 which to read the data. Even in this case, the dataset must be located
2076 within the directory containing the enriched datasets

2077 • **-c or --condition**: specifies the condition for the analysis, which will
2078 be automatically overwritten in the *.sinfo* file. It is possible to spec-
2079 ify more than one condition, but it is recommended to use the logical
2080 operator &&

2081 • **-r or --register**: specifies the directory where to save the text file with
2082 the outcomes

2083 Performing this single analysis will produce a single output file contain-
2084 ing the invariants discovered, file that will be saved in the directory spec-
2085 ified by the user via the **-r** command-line option (by default the file will be
2086 saved in the directory \$(project_dir)/daikon/Daikon_Invariants/results)
2087 to be examined later by the user.

2088 In conclusion, the two types of analysis presented, along with the pos-
2089 sibility of allowing for more refined analysis at a later date and Daikon's
2090 redefined output, make this stage much more complete, clear, and power-
2091 ful than before

2092 **4.2.4 Phase 4: Business Process Analysis**

Case study: the iTrust SWaT System

2093 HAVING introduced the innovative framework and highlighted its po-
2094 tential in the preceding chapter, we now turn our attention to the
2095 case study where we will apply this framework. As previously mentioned
2096 in Chapter 4 and demonstrated through various examples in the same
2097 chapter, our focus will be on the **iTrust SWaT system** [36], developed by
2098 the iTrust – Center for Research in Cyber Security of University of Singa-
2099 pore for Technology and Design [30]. The acronym SWaT represents *Secure*
2100 *Water Treatment*.

2101 The iTrust SWaT system is a testbed that replicates on a small scale
2102 a real water treatment plant arises to support research in the area of cy-
2103 ber security of industrial control systems and has been operational since
2104 March 2015: it is still being used by students at the University of Singapore
2105 for educational and training purposes and is available to organizations to
2106 train their operators on cyber physical incidents.

2107 5.1 Architecture

2108 In contrast to the virtualized testbed discussed in Section 3.2.1 by Cec-
2109 cato et al., the iTrust SWaT system is composed entirely of physical hard-
2110 ware components. It encompasses various elements, starting from field

2111 devices and extending to PLCs, HMI, SCADA workstations, and the SCADA
2112 server (also referred to as the *historian*). The historian is responsible for
2113 recording data from field devices for further analysis. In the upcoming
2114 sections, we will delve deeper into the architecture of the physical process
2115 and the communication network.

2116 **5.1.1 Physical Process**

2117 The physical process of the SWaT consists of six stages, denoted P1
2118 through P6. These stages are [61][62]:

- 2119 1. **taking in raw water:** feeds unfiltered water into the system
- 2120 2. **chemical dosing:** adds chemicals to water useful for initial pretreat-
2121 ment
- 2122 3. **Ultra Filtration (UF) system:** the water is filtered through a semi-
2123 permeable membrane (ultrafiltration membrane) using the liquid pres-
2124 sure, effectively capturing impurities and suspended solids, as well
2125 as removing bacteria, viruses, and other pathogens present in the
2126 water.
- 2127 4. **dechlorination:** removes residual chlorine from disinfected water
2128 using ultraviolet lamps
- 2129 5. **Reverse Osmosis (RO):** performs further filtration of the water
- 2130 6. **backwash process:** cleans the membranes in UF using the water pro-
2131 duced by RO

2132 Figure 5.1 shows a graphical representation of the architecture and the six
2133 stages of the SWaT system.

2134 The SWaT system incorporates an array of sensors that play a crucial
2135 role in monitoring the system's operations and ensuring their safe. These
2136 sensors are responsible for continuously collecting data and providing
2137 valuable insights into the functioning of the system. These sensors are:

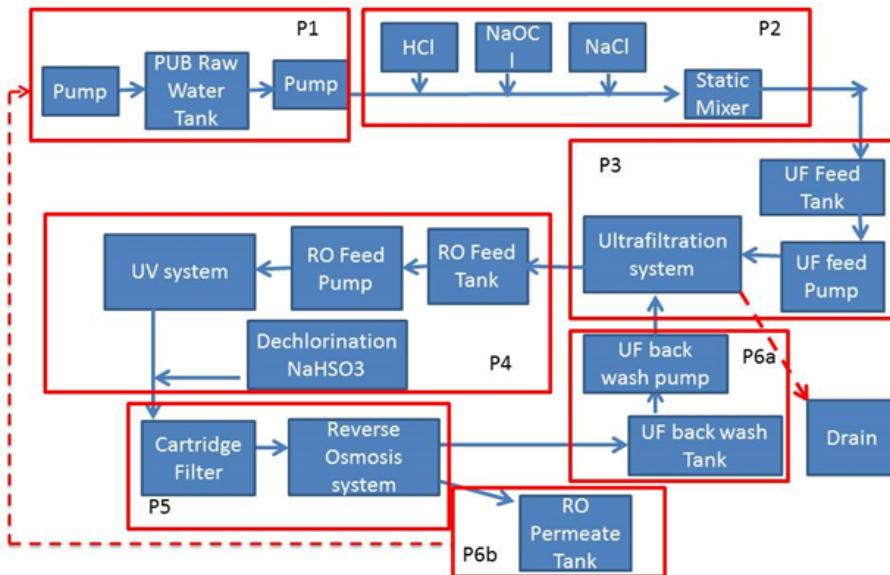


Figure 5.1: SWaT architecture

- Level Indication Transmitter (measured in mm)
- Flow Indication Transmitter (m³/hr)
- Analyser Indicator Transmitter
 - Conductivity ($\mu\text{S}/\text{cm}$)
 - pH
 - Oxidation Reduction Potential (mV)
- Differential Pressure Indicator Transmitter (kPa)
- Pressure Indicator Transmitter (kPa)

The sensors and actuators associated with each PLC are shown in Figure 5.2.

Sensors and actuators are mapped to tags by the communication protocol used (see 5.1.2): a tag can be addressed via string descriptor defined by the system designer (e.g. MV101, to indicate motorized valve number 1 at stage 1) or by referring directly to the analog/digital pins of the PLC I/O unit [62].

Raw Water	Pre-Treatment	Ultra-Filtration	De-Chlorination	Reverse Osmosis	RO Product
P-101 Stopped	P-201 Stopped	P-301 Stopped	P-401 Stopped	P-501 Stopped	P-601 Stopped
P-102 Stopped	P-202 Stopped	P-302 Stopped	P-402 Stopped	P-502 Stopped	P-602 Stopped
MV-101 Closed	P-203 Stopped	MV-301 Closed	P-403 Stopped	MV-501 Closed	P-603 Stopped
LIT-101 520 mm	P-204 Stopped	MV-302 Closed	P-404 Stopped	MV-502 Closed	LS-601 Normal
FIT-101 0.00 m³/h LL	P-205 Stopped	MV-303 Closed	UV-401 Stopped	MV-503 Closed	LS-602 HIGH
FIT-201 0.00 m³/h LL	P-206 Stopped	MV-304 Closed	LS-401 Normal	MV-504 Closed	LS-603 LOW
	P-207 Stopped	PSH-301 Normal	LIT-401 1008 mm H	PSL-501 Normal	FIT-601 0.00 m³/h LL
	P-208 Stopped	DPSH-301 Normal	FIT-401 0.00 m³/h LL	PSH-501 Normal	
	MV-201 Closed	LIT-301 1012 mm H	AIT-401 0.17 ppm	AUT-501 6.89	
	LS-201 Normal	FIT-301 0.00 m³/h	AIT-402 275.70 mV	AUT-502 204.20 mV	
	LS-202 Normal	DPT-301 0.95 kPa LL		AUT-503 264.23 µS/cm H	
	LS-203 Normal			AUT-504 14.27 µS/cm H	
	AIT-201 142.18 µS/cm L			FIT-501 0.00 m³/h L	
	AIT-202 7.20 H			FIT-502 0.00 m³/h HH	
	AIT-203 293.59 mV L			FIT-503 0.00 m³/h HH	
				FIT-504 0.00 m³/h LL	
				PIT-501 2.64 kPa LL	
				PIT-502 0.00 kPa H	
				PIT-503 0.00 kPa	

Figure 5.2: Sensors and actuators associated with each PLC

2153 5.1.2 Control and Communication Network

2154 The SWaT system's network architecture follows the principles of lay-
 2155 ering and zoning, which enable segmentation and control of traffic within
 2156 the network.

2157

2158 Five layers are present starting from the highest to the lowest:

- 2159 • Layer 3.5 – Demilitarized Zone (DMZ)
- 2160 • Layer 3 – Operation Management (Historian)
- 2161 • Layer 2 – Supervisory Control (Touch Panel, Engineering Worksta-
2162 tion, HMI Control Clients)
- 2163 • Layer 1 – Plant Control Network (PLCs) (Star Network)
- 2164 • Layer 0 – Process (Actuator/Sensors and Input/output modules)
2165 (Ring Network)

2166 PLCs at Layer 1 communicate with their respective sensors and actuators
 2167 at Layer 0 through a conventional ring network topology based on Ether-
 2168 net/IP, to ensure that the system can tolerate the loss of a single link

without any adverse impact on data or control functionality.
 PLCs between the different process stages at Layer 1 communicate with each other through a star network topology using the CIP protocol on EtherNet/IP, previously discussed in Section 2.2.6.3.

Regarding zoning, the SWaT system is divided into three zones, each containing one or more layers. These zones are, in descending order of security level:

- **Plant Control Network, or Control System:** includes layers from 0 to 2
- **DMZ:** includes Layer 3.5
- **Plant Network:** includes Layer 3

Figure 5.3 provides a clearer visualization of the zoning and layer division within the network architecture of the SWaT system. This diagram highlights the distinct zones and their corresponding layers, offering a comprehensive overview of the system's network structure.

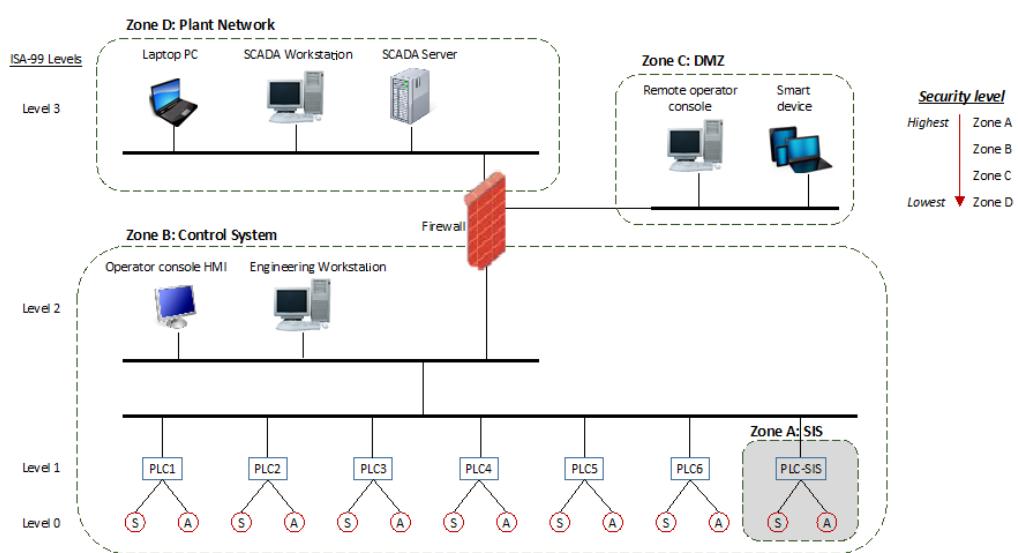


Figure 5.3: SWaT network architecture

2184 A specific IP address is associated with each device: in Table 5.1 we
2185 report the addresses for the PLCs, historian, and Touch Panel in the *Plant*
2186 *Control Network* (PCN) zone:

IP Address	Device
192.168.1.10	PLC1
192.168.1.20	PLC2
192.168.1.30	PLC3
192.168.1.40	PLC4
192.168.1.50	PLC5
192.168.1.60	PLC6
192.168.1.100	Touch Panel (PCN)
192.168.1.200	Historian Server

Table 5.1: main IP addresses of the six PLCs and SCADA in SWaT

2187 5.2 Datasets

2188 To facilitate the study and testing of technologies related to cyber security in Industrial Control Systems and critical infrastructure, iTrust offers 2189 researchers worldwide the opportunity to access a range of datasets [63]. 2190 These datasets consist of a collection of data obtained from the SWaT system, 2191 encompassing information on both the physical processes and network 2192 communications. The data is organized into different years, and 2193 researchers can request access to these datasets for their analysis and 2194 experimentation purposes. 2195

2196 **Physical Process Datasets** The datasets containing information about 2197 the physical processes are provided in CSV format files. These files 2198 encompass data collected during different time intervals, which can vary 2199 from a few hours to entire days. The granularity of the data is typically 2200 at a one-second interval, although there may be some exceptions. The collected data primarily consists of timestamped sensor measurements and 2201

2202 actuator status values for each PLC, describing the physical properties of
2203 the testbed in operational mode.

2204 **Network Communications Datasets** Network communications are typi-
2205 cally available in the form of *Packet Capture* format (PCAP) files. These files
2206 contain captures of communication network traffic, allowing researchers
2207 to analyze and examine the network interactions. In some instances, CSV
2208 files are provided instead of PCAP files, featuring different characteristics
2209 for the collected data.

2210 **5.2.1 Our Case Study: the 2015 Dataset**

2211 The dataset selected as a case study to apply the framework discussed
2212 in the previous chapter is specifically the dataset from the year 2015 [64].
2213 The main reason for this choice is the unique characteristics found in the
2214 physical process dataset that are not present in datasets from subsequent
2215 years.

2216 **Physical Process Data** The data collection process lasted 11 consecutive
2217 days, 24 hours per day. During the first 7 days, the system operated nor-
2218 mally without any recorded attacks. However, attacks were observed dur-
2219 ing the remaining 4 days. The collected data reflects the impact of these
2220 attacks, leading to the creation of two separate CSV files: one containing
2221 the recorded data of SWaT during the system's regular operations, and
2222 the other containing data recorded during the days of the attacks. To en-
2223 sure accurate information about the system, the dataset pertaining to the
2224 normal operations, which spans seven days, was chosen for analysis.

2225 Data collection occurs at a frequency of one data point per second,
2226 with the assumption that significant attacks cannot occur within a shorter
2227 time frame. Additionally, the firmware of the PLCs remains unchanged
2228 throughout the data collection period.

2229 At the beginning of data gathering the tanks are empty and the system
2230 must be initialized in order to then reach full operation: it typically takes

2231 around five hours for all tanks to be fully filled and for the system to sta-
 2232 bilize and reach the appropriate operational state (see Figure 5.4).

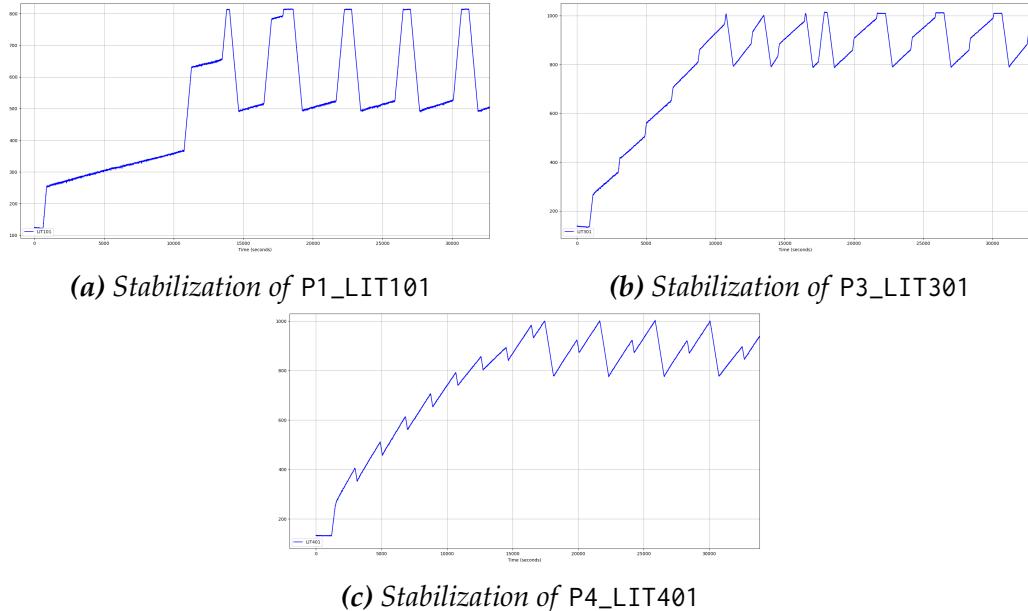


Figure 5.4: SWaT stabilization

2233 In total, the dataset consists of thousand and thousand of samples re-
 2234 lated to 51 attributes.

2235 **Network Traffic** The network traffic was collected using an appliance
 2236 from a well-known network hardware manufacturer and was made avail-
 2237 able only in CSV format and not PCAP format. Table 5.2 shows some of
 2238 the main data captured:

Category	Description
Date	Date of Log
Time	Time of Log
Origin	IP of server
Source IP	IP address of source
Destination IP	IP address of destination
Protocol	Network protocol

Application Name	Name of application
Modbus Function Code	Function Code
Modbus Function Description	Description of function
Modbus Transaction ID	Transaction ID
SCADA Tag	Sensor or actuator ID
Modbus Value	Value transmitted
Service / Destination Port	Port number of destination IP
Source Port	Port number of source IP

Table 5.2: SWaT network traffic data

2239 Unfortunately, the data provided are partial as it only includes read-
 2240 ings from sensors and does not include information on actuator status.

2241

2242 Do not mislead by the word *Modbus* in the captured data fields: this is
 2243 most likely a feature of the sniffing appliance, which may have encapsu-
 2244 lated EtherNet/IP-CIP protocol packets in Modbus frames, as described
 2245 by T. A. Snide in [65].

2246 To provide evidence that the protocol being used is CIP over EtherNet/IP,
 2247 we can examine the "Source / Destination port" field in the captured data.

2248 The presence of TCP port 44818 as the destination port aligns with the
 2249 TCP port used for transporting **explicit messages** in the EtherNet/IP pro-
 2250 tocol, as explained in Section 2.2.6.2. Additionally, the "Application name",
 2251 "Modbus Function Code" and "Modbus Function Description" fields fur-
 2252 ther support this conclusion. The "Application name" explicitly states
 2253 "CIP_read_tag_service," while the "Modbus Function Code" field contains
 2254 a decimal value of 76, which does not correspond to any known Modbus
 2255 Function Calls. When converted to hexadecimal, this value is 4C, match-
 2256 ing the CIP protocol read tag request code as indicated in the "Modbus
 2257 Function Description" field.

2258 In summary, the datasets for subsequent years were not selected due
 2259 to certain limitations. In the case of the year 2017, the physical system

2260 data had a wide granularity, resulting in significant information loss. Ad-
2261 ditionally, in some cases, the datasets were considered "spurious" as there
2262 was no clear differentiation between data obtained during normal SWaT
2263 operations and data associated with attacks.

2264 Regarding network traffic data, there were instances where the data was
2265 either missing or fragmented into large PCAP files weighing several gi-
2266 gabytes (more than 6 GB each!), despite containing only a few minutes'
2267 worth of data. This made it impractical to manage such files given the
2268 available resources.

2269 Despite their large quantity, the CSV files associated with network traffic
2270 for the year 2015 are comparatively smaller in size, just slightly over 100
2271 MB each. These files cover a more extensive time period, which facilitates
2272 easier management and analysis. Prioritizing the physical process dataset
2273 as crucial, I have selected the year 2015 as a case study, even though the
2274 network data may be incomplete.

Our framework at work: reverse engineering of the SWaT system

6.1 Pre-processing

6.2 Graph Analysis

6.2.1 Conjectures About the System

6.3 Invariants Analysis

6.3.1 Actuators Detection

6.3.2 Daikon Analysis and Results Comparing

6.4 Extra information on the Physics

6.5 Business Process Analysis

Chapter

7

Conclusions

7.1 Discussions

7.2 Guidelines

7.3 Future work

List of Figures

2.1	ICS architecture schema	6
2.2	PLC architecture	9
2.3	PLC communication schema	10
2.4	Comparison between ST language and Ladder Logic	11
2.5	Example of HMI for a water treatment plant	14
2.6	Modbus Request/Response schema	16
2.7	Modbus RTU frame and Modbus TCP frame	17
2.8	Example of packet sniffing on the Modbus protocol	19
2.9	OSI model for EtherNet/IP stack	20
3.1	Workflow of Ceccato et al.'s stages and operations with used tools	28
3.2	The simplified SWaT system used for running Ceccato et al. methodology	29
3.3	Output graphs from graph analysis	33
3.4	An example of Disco generated activity diagram for PLC2 .	37
3.5	Execution traces of InputRegisters_IW0 on the three PLCs .	39
3.6	Business process with states and Modbus commands for the three PLCs	41
3.7	Water physics compared: simulated physics in the Simulink model (a) and physics in a real system (iTrust SWaT) (b). Fluctuations in the tank level in (b), almost completely absent in (a), can be appreciated.	47

3.8	Behavior of the Graph Analysis on the Ceccato et al.'s tool	49
3.9	Histogram plot overshadowing statistical information shown on the terminal window in the background	50
3.10	Example of Daikon's output	52
4.1	Slope comparison with granularity 30 (a), 60 (b) and 120 seconds (c)	68
4.2	Savitzky-Golay filter (blue line) and STL decomposition (green) comparison	69
4.3	Slope after the application of the STL decomposition	70
4.4	The new slope representation (green line) and the smoothed measurement data obtaind with the STL decomposition (red)	71
4.5	Plotting registers on the same y-axis	76
4.6	Example of the new graph analysis	77
4.7	Example of transitive closure graphs for invariants in PLC1 of iTrust SWaT system	83
4.8	Directory (a) and outcome files (b) for the single actuator states analysis	86
4.9	Daikon outcome files for system configuration analysis. Each file represents a single system state	87
5.1	SWaT architecture	93
5.2	Sensors and actuators associated with each PLC	94
5.3	SWaT network architecture	95
5.4	SWaT stabilization	98

List of Tables

2.1	differences between Information Technology (IT) and Industrial Control Systems (ICSs)	4
2.2	Modbus Function Codes list	18
5.1	main IP addresses of the six PLCs and SCADA in SWaT	96
5.2	SWaT network traffic data	99

Listings

3.1	Example of registers capture	31
3.2	Example of raw network capture	32
3.3	Statistical data for PLC1_InputRegisters_IW0 register	34
3.4	Generic example of a .spinfo file for customizing rules in Daikon	35
3.5	The three sections of Daikon analysis outcomes	35
4.1	Novel Framework structure and Python scripts	58
4.2	Paths and parameters for the Pre-processing phase in <i>config.ini</i> file	62
4.3	config.ini parameters for dataset enriching	64
4.4	Example of preliminar system analysis	72
4.5	Standard Daikon output for PLC1 of the iTrust SWaT system	81
4.6	Revised Daikon output with transitive closures for PLC1 of the iTrust SWaT system	84
4.7	Daikon outcomes for the system configuration P1_MV101 == 2, P1_P101 == 1 on P1_LIT101	88

References

- [1] NIST. *ICS defintion*. URL: https://csrc.nist.gov/glossary/term/industrial_control_system (visited on 2023-04-06).
- [2] Wikipedia. *Stuxnet*. URL: <https://en.wikipedia.org/wiki/Stuxnet> (visited on 2023-05-18).
- [3] Wikipedia. *Industroyer*. URL: <https://en.wikipedia.org/wiki/Industroyer> (visited on 2023-05-18).
- [4] F. Robles and N. Perlroth. “‘Dangerous Stuff’: Hackers Tried to Poison Water Supply of Florida Town”. In: *The New York Times* (2021-02-08). URL: <https://www.nytimes.com/2021/02/08/us/oldsmar-florida-water-supply-hack.html> (visited on 2023-05-18).
- [5] A. Akbarzadeh. “Dependency based risk analysis in Cyber-Physical Systems”. PhD thesis. Norwegian University of Science, Technology - Faculty of Information Technology, and Electrical Engineering, 2023. ISBN: 978-82-326-6986-8 (electronic version).
- [6] G. M. Makrakis, C. Kolias, G. Kambourakis, C. Rieger, and J. Benjamin. “Industrial and Critical Infrastructure Security: Technical Analysis of Real-Life Security Incidents”. In: *IEEE Access* (2021-12-06). DOI: <https://dx.doi.org/10.1109/ACCESS.2021.3133348>.
- [7] NIST. *PLC defintion*. URL: https://csrc.nist.gov/glossary/term/programmable_logic_controller (visited on 2023-04-06).

- [8] W. Bolton. *Programmable Logic Controllers, 6th edition*. Newnes, 2015, pp. 7–9.
- [9] M. Ceccato, Y. Driouich, R. Lanotte, M. Lucchese, and M. Merro. “Towards Reverse Engineering of Industrial Physical Processes”. In: CPS4CIP@ESORICSAt 2022 (Copenhagen, Denmark, Sept. 30, 2022). 2022, 273–290. DOI: https://doi.org/10.1007/978-3-031-25460-4_15.
- [10] H. S. G. Pussewalage, P. S. Ranaweera, and V. Oleshchuk. “PLC security and critical infrastructure protection”. In: 2013 IEEE 8th International Conference on Industrial and Information Systems (Dec. 2013). 2013. DOI: <https://dx.doi.org/10.1109/ICIInfs.2013.6731959>.
- [11] Wikipedia. *Remote Terminal Unit*. URL: https://en.wikipedia.org/wiki/Remote_terminal_unit (visited on 2023-04-08).
- [12] What is SCADA? URL: <https://www.inductiveautomation.com/resources/article/what-is-scada> (visited on 2023-04-05).
- [13] NIST. *HMI defintion*. URL: https://csrc.nist.gov/glossary/term/human_machine_interface (visited on 2023-04-11).
- [14] Zscaler. *What Is the Purdue Model for ICS Security?* URL: <https://www.zscaler.com/resources/security-terms-glossary/what-is-purdue-model-ics-security> (visited on 2023-05-17).
- [15] Schneider Electrics. *What is Modbus and how does it work?* URL: <https://www.se.com/us/en/faqs/FA168406/> (visited on 2023-04-13).
- [16] O. Morando. *Hacking: Modbus - Cyber security OT/ICS*. URL: <https://blog.omarmorando.com/hacking/ot-ics-hacking/hacking-modbus> (visited on 2023-04-15).
- [17] Modbus.org. “MODBUS/TCP Security”. In: pp. 7–8. URL: https://modbus.org/docs/MB-TCP-Security-v21_2018-07-24.pdf (visited on 2023-04-16).

- [18] ODVA Inc. "EtherNet/IP - CIP on Ethernet Technology". In: URL: https://www.odva.org/wp-content/uploads/2021/05/PUB00138R7_Tech-Series-EtherNetIP.pdf (visited on 2023-04-18).
- [19] *Odva, Inc.* URL: <https://www.odva.org> (visited on 2023-04-21).
- [20] *Introduction to EtherNet/IP Technology*. URL: https://scadahacker.com/library/Documents/ICS_Protocols/Intro%20to%20EthernetIP%20Technology.pdf (visited on 2023-04-19).
- [21] ODVA Inc. *Common Industrial Protocol*. URL: <https://www.odva.org/technology-standards/key-technologies/common-industrial-protocol-cip/> (visited on 2022-12-26).
- [22] Wikipedia. *Common Industrial Protocol*. URL: https://en.wikipedia.org/wiki/Common_Industrial_Protocol (visited on 2023-04-21).
- [23] *What is the Common Industrial Protocol (CIP)?* URL: <https://www.motioncontroltips.com/what-is-the-common-industrial-protocol-cip/> (visited on 2023-04-21).
- [24] *CIP (Common Industrial Protocol): CIP messages, device types, implementation and security in CIP*. URL: <https://resources.infosecinstitute.com/topic/cip/> (visited on 2023-04-21).
- [25] B. Green, M. Krotofil, and A. Abbasi. "On the Significance of Process Comprehension for Conducting Targeted ICS Attacks". In: Workshop on Cyber-Physical Systems Security and PrivaCy (Nov. 2017). ACM, 2017, pp. 57–67. DOI: <https://doi.org/10.1145/3140241.3140254>.
- [26] A. Keliris and M. Maniatakos. "ICSREF: A Framework for Automated Reverse Engineering of Industrial Control Systems Binaries". In: 26th Annual Network and Distributed System Security Symposium (NDSS) (Feb. 2019). 2019. DOI: <https://doi.org/10.48550/arXiv.1812.03478>.
- [27] Wikipedia. *CODESYS*. URL: <https://en.wikipedia.org/wiki/CODESYS> (visited on 2023-05-18).

- [28] Y. Yuan, X. Tang, W. Zhou, W. Pan, X. Li, H. T. Zhang, H. Ding, and J. Goncalves. "Data driven discovery of cyber physical systems". In: *Nature Communications* 10 (2019-10-25). URL: <https://www.nature.com/articles/s41467-019-12490-1> (visited on 2023-04-20).
- [29] C. Feng, V.R. Palleti, A. Mathur, and D. Chana. "A Systematic Framework to Generate Invariants for Anomaly Detection in Industrial Control Systems". In: NDSS Symposium 2019 (San Diego, CA, USA, Feb. 24–27, 2019). 2019. DOI: <https://dx.doi.org/10.14722/ndss.2019.23265>.
- [30] Singapore University of Technology and Design. *iTrust - Center for Research in Cyber Security*. URL: <https://itrust.sutd.edu.sg/> (visited on 2023-05-15).
- [31] K. Pal, A. Adepu, and J. Goh. "Effectiveness of Association Rules Mining for Invariants Generation in Cyber-Physical Systems". In: 2017 IEEE 18th International Symposium on High Assurance Systems Engineering (HASE) (Jan. 2017). 2017. DOI: <https://doi.org/10.1109/HASE.2017.21>.
- [32] *Association rules mining*. URL: https://en.wikipedia.org/wiki/Association_rule_learning (visited on 2023-04-21).
- [33] K. Pal, A. Adepu, and J. Goh. "Cyber-Physical System Discovery: Reverse Engineering Physical Processes". In: 3rd ACM Workshop on Cyber-Physical System Security (Apr. 2017). 2017, pp. 3–14. DOI: <https://doi.org/10.1145/3055186.3055195>.
- [34] B. Perelman. "Reconnaissance in Industrial Networks: What You Don't See Can Hurt You". In: *SecurityWeek* (2017-04-04). URL: <https://www.securityweek.com/reconnaissance-industrial-networks-what-you-dont-see-can-hurt-you/> (visited on 2023-05-18).
- [35] Ceccato *et al.* *reverse engineering prototype tool*. URL: <https://github.com/SeUniVr/PLC-RE> (visited on 2023-04-23).

- [36] iTrust Singapore University of Technology and Design. *Secure Water Treatment*. URL: <https://itrust.sutd.edu.sg/testbeds/secure-water-treatment-swat/> (visited on 2023-05-15).
- [37] R. Lanotte, M. Merro, and A. Munteanu. "Industrial Control Systems Security via Runtime Enforcement". In: *ACM Transactions on Privacy and Security* 26.4 (2023-02), pp. 1–41. DOI: <https://doi.org/10.1145/3546579>.
- [38] M. Oliani. *AttackPLC - Project for Network Security teaching*. URL: <https://github.com/marcooliani/AttackPLC> (visited on 2023-04-23).
- [39] Ray - Productionizing and scaling Python ML workloads simply. URL: <https://www.ray.io/> (visited on 2023-04-25).
- [40] Tshark. URL: <https://tshark.dev/> (visited on 2023-04-25).
- [41] Wireshark. URL: <https://www.wireshark.org/> (visited on 2023-04-25).
- [42] pandas - Python Data Analysis library. URL: <https://pandas.pydata.org/> (visited on 2023-04-25).
- [43] NumPy - The fundamental package for scientific computing with Python. URL: <https://numpy.org/> (visited on 2023-04-25).
- [44] R project for statistical computing. URL: <https://www.r-project.org/> (visited on 2023-04-25).
- [45] SciPy - Fundamental algorithms for scientific computing with Python. URL: <https://scipy.org/> (visited on 2023-04-25).
- [46] The Daikon dynamic invariant detector. URL: <http://plse.cs.washington.edu/daikon/> (visited on 2023-04-25).
- [47] Enhancing Daikon output. URL: <https://plse.cs.washington.edu/daikon/download/doc/daikon/Enhancing-Daikon-output.html#Splitter-info-file-format> (visited on 2023-04-25).
- [48] Process mining. URL: https://en.wikipedia.org/wiki/Process_mining (visited on 2023-04-26).

- [49] *Fluxicon Disco*. URL: <https://fluxicon.com/disco/> (visited on 2023-04-26).
- [50] *OpenPLC - Open source PLC software*. URL: <https://openplcproject.com/> (visited on 2023-04-23).
- [51] *Docker*. URL: <https://www.docker.com/> (visited on 2023-04-26).
- [52] MathWorks. *Simulink*. URL: <https://it.mathworks.com/products/simulink.html> (visited on 2023-04-26).
- [53] T. Carlsson. "Industrial network market shares 2023". In: *HMS Networks* (2023-05-05). URL: <https://www.hms-networks.com/news-and-insights/news-from-hms/2023/05/05/industrial-network-market-shares-2023> (visited on 2023-05-18).
- [54] *ProM Tools - The Process Mining Framework*. URL: <https://promtools.org/> (visited on 2023-04-26).
- [55] Fraunhofer Institute for Applied Information Technology. *pm4py - Process Mining for Python*. URL: <https://promtools.org/> (visited on 2023-04-26).
- [56] statsmodels. *statsmodels - statistical models, hypothesis tests, and data exploration*. URL: <https://www.statsmodels.org/stable/index.html> (visited on 2023-05-05).
- [57] NetworkX. *NetworkX - Network Analysis in Python*. URL: <https://networkx.org/> (visited on 2023-05-05).
- [58] Wikipedia. *Savitzky–Golay filter*. URL: https://en.wikipedia.org/wiki/Savitzky%E2%80%93Golay_filter (visited on 2023-05-06).
- [59] Otext. *STL decomposition*. URL: <https://otexts.com/fpp2/stl.html> (visited on 2023-05-06).
- [60] Wikipedia. *Transitive closures*. URL: https://en.wikipedia.org/wiki/Transitive_closure (visited on 2023-05-11).

- [61] iTrust Singapore University of Technology and Design. *Secure Water Treatment (SWaT) Testbed*. URL: https://itrust.sutd.edu.sg/wp-content/uploads/2021/07/SWaT_technical_details-160720-v4.4.pdf (visited on 2023-05-15).
- [62] A. Mathur and N. O. Tippenhauer. "SWaT: a water treatment testbed for research and training on ICS security". In: 2016 International Workshop on Cyber-physical Systems for Smart Water Networks (CySWater) (Apr. 2016). 2016. DOI: <http://dx.doi.org/10.1109/CySWater.2016.7469060>.
- [63] iTrust Singapore University of Technology and Design. *Datasets*. URL: https://itrust.sutd.edu.sg/itrust-labs_datasets/ (visited on 2023-05-15).
- [64] S. Adepu, J. Goh, K.N. Junejo, and A. Mathur. "A Dataset to Support Research in the Design of Secure Water Treatment Systems". In: The 11th International Conference on Critical Information Infrastructures Security (France, Oct. 2016). 2016.
- [65] T. A. Snide. *CIP-Modbus Integration*. 2008. URL: https://www.modbus.org/docs/CIP%20Modbus%20Integration%20Hanover%20Fair_0408.pdf (visited on 2023-05-17).