

UNIVERSITY OF VERONA

Department of COMPUTER SCIENCE

Master's Degree in
COMPUTER SCIENCE AND ENGINEERING

Master Thesis

**Towards Process Comprehension of Industrial
Control Systems: a Framework for Analyzing
Industrial Systems**

Candidate:

Marco OLIANI
VR457249

Supervisor:

Prof. **Massimo MERRO**

Co-supervisor:

Prof. **Ruggero LANOTTE**
University of Insubria

Academic Year 2022/2023

*“If you spend more on coffee than on IT security, you
will be hacked. What’s more, you deserve to be hacked”*
(Richard Clarke)

Abstract

Bla bla bla

Contents

1	Introduction	1
1.1	Contribution	1
1.2	Outline	1
2	Background on Industrial Control Systems	3
2.1	Industrial Control Systems Architecture	5
2.2	Operational Technology Networks	7
2.2.1	Field I/O Devices Layer	7
2.2.2	Controller Network Layer	8
2.2.2.1	Programmable Logic Controllers	8
2.2.2.2	Remote Terminal Units	12
2.2.3	Area Control Layer	13
2.2.3.1	Supervisory Control And Data Acquisition	13
2.2.3.2	Human-Machine Interface	13
2.2.4	Operations/Control Layer	14
2.2.5	Demilitarized Zone	14
2.2.6	Industrial Protocols	15
2.2.6.1	Modbus	16
2.2.6.2	EtherNet/IP	19
2.2.6.3	Common Industrial Protocol (CIP)	21
3	State of the Art	23
3.1	Literature on Process Comprehension	24

3.2	Ceccato et al.'s methodology for analyzing water-tank systems	26
3.2.1	Testbed	28
3.2.2	Scanning of the System and Data Pre-processing . . .	30
3.2.3	Graphs and Statistical Analysis	32
3.2.4	Invariants Inference and Analysis	34
3.2.5	Businness Process Mining and Analysis	35
3.2.6	Application	36
3.2.7	Limitations	43
4	A Framework to Improve Ceccato et al.'s Work.	53
4.1	The novel Framework	54
4.1.1	Framework Structure	56
4.1.2	Python Libraries and External Tools	57
4.2	Analysis Phases	58
4.2.1	Phase 1: Data Pre-processing	58
4.2.1.1	Subsystem Selection	60
4.2.1.2	Dataset Enrichment	62
4.2.1.3	Datasets Merging	69
4.2.1.4	Brief Analysis of the Obtained Subsystem .	70
4.2.2	Phase 2: Graphs and Statistical Analysis	73
4.2.3	Phase 3: Invariant Inference and Analysis	78
4.2.3.1	Revised Daikon Output	79
4.2.3.2	Types of Analysis	83
4.2.4	Phase 4: Businness Process Analysis	88
5	Case study: the iTrust SWaT System	89
5.1	Architecture	89
5.1.1	Physical Process	90
5.1.2	Control and Communication Network	92
5.2	Datasets	94
5.2.1	Our Case Study: the 2015 Dataset	95

6	Our framework at work: reverse engineering of the SWaT system	99
6.1	Pre-processing	99
6.2	Graph Analysis	99
6.2.1	Conjectures About the System	99
6.3	Invariants Analysis	99
6.3.1	Actuators Detection	99
6.3.2	Daikon Analysis and Results Comparing	99
6.4	Extra information on the Physics	99
6.5	Business Process Analysis	99
7	Conclusions	101
7.1	Discussions	101
7.2	Guidelines	101
7.3	Future work	101
	List of Figures	103
	List of Tables	105
	Listings	107
	References	109

Introduction

LOREM ipsum dolor bla bla bla. Ma dove metto l'abstract? Prova di interlinea che direi posso anche andare bene, ma bisogna poi vedere il tutto come si incastra alla fine, in modo da ottenere un bel risultato alla vista.

1.1 Contribution

Lorem Ipsum is simply dummy text of the printing and typesetting industry. Lorem Ipsum has been the industry's standard dummy text ever since the 1500s, when an unknown printer took a galley of type and scrambled it to make a type specimen book. It has survived not only five centuries, but also the leap into electronic typesetting, remaining essentially unchanged. It was popularised in the 1960s with the release of Letraset sheets containing Lorem Ipsum passages, and more recently with desktop publishing software like Aldus PageMaker including versions of Lorem Ipsum.

1.2 Outline

The thesis is structured as follows:

Chapter 2: provides background on the topics covered in this thesis: Industrial Control Systems (ICSs), Supervisory Control And Data Acquisition (SCADA), Programmable Logic Controllers (PLCs) and other devices, industrial communication protocols.

Chapter 3:

Chapter 4:

Chapter 5:

Chapter 6:

Chapter 7:

Chapter 2

Background on Industrial Control Systems

1 INDUSTRIAL CONTROL SYSTEMS (ICSs) are information systems used to
2 control industrial processes such as manufacturing, product handling,
3 production, and distribution [1]. ICSs are often found in critical infrastruc-
4 ture facilities such as power plants, oil and gas refineries, and chemical
5 plant ICSs are different from traditional IT systems in several key ways.

6 Firstly, ICSs are designed to control physical processes, whereas IT sys-
7 tems are designed to process and store data. This means that ICSs have dif-
8 ferent requirements for availability, reliability, and performance. Secondly,
9 ICSs are typically deployed in environments that are harsh and have lim-
10 ited resources, such as extreme temperatures and limited power. Thirdly,
11 the protocols hardware and software used in ICSs are often proprietary.

12 ICSs are becoming increasingly connected to the internet and other net-
13 works, which has led to increased concerns about their security. Industrial
14 systems were not originally designed with security in mind, and many of
15 them have known vulnerabilities that could be exploited by attackers. Ad-
16 ditionally, the use of legacy systems and equipment can make it difficult
17 to implement security measures. As a result, ICSs are increasingly seen
18 as a potential target for cyber attacks, which could have serious conse-
19 quences for the safe and reliable operation of critical infrastructure: some
20 notorious examples of cyber attacks are (i) the **STUXnet** worm [2], which

21 purpose was to sabotage the nuclear centrifuges of the enrichment plant
 22 at the Natanz nuclear facility in Iran; (ii) **Industroyer** [3], also referred
 23 as *Crashoverride*, responsible for the attack on the Ukrainian power grid
 24 on December 17, 2016; (iii) the attack on February, 2021 to a water treat-
 25 ment plant in Oldsmar, Florida [4], where the level of sodium hydroxide
 26 was intentionally increased to a level approximately 100 times higher than
 27 normal.

28
 29 The increasing connectivity of ICSs and the associated security risks have
 30 led to a growing interest in the field of ICS security. Researchers and prac-
 31 titioners are working to develop new security technologies, standards, and
 32 best practices to protect ICSs from cyber attacks. This includes efforts to
 33 improve the security of ICS networks and devices, as well as the develop-
 34 ment of new monitoring and detection techniques to identify and respond
 35 to cyber attacks.

36
 37 Table 2.1 summarizes the differences between traditional IT and ICSs [5]:

	Traditional IT	ICSs
Focus	Data	Asset
Update Frequency	High	Low
Priority	Confidentiality Integrity Availability	Availability Integrity Confidentiality
Operating System	Standardized	Proprietary
Protocols	Standardized	Proprietary
Attacker Motivation	Monetization	Disruption

Table 2.1: differences between Information Tecnology (IT) and Industrial Control Systems (ICSs)

2.1 Industrial Control Systems Architecture

In the past, there has been a clear division between *Information Technology* (IT) and *Operational Technology* (OT), both at the technical and organizational levels. Each domain has maintained its own distinct technology stacks, protocols, and standards. However, with the emergence of Industry 4.0 and the rapid expansion of industrial automation, which heavily relies on IT tools for monitoring and controlling critical infrastructures, the boundary between IT and OT has started to blur. This trend has paved the way for greater integration between these two domains, thus improving productivity and process quality.

General ICS architecture consists in **six levels** each representing a functionality: this architecture comprising the OT and IT parts is represented in Figure 2.1 [6][5], according to the *Purdue Enterprise Reference Architecture* (PERA), or simply **Purdue Model**:

- Level 0 (**Processes, or Field I/O Devices**): contains **field devices**.
- Level 1 (**Intelligent Devices, or Controller Network**): includes **local or remote controllers** that sense, monitor and control the physical process, such as **PLCs** (2.2.2.1) and **RTUs** (2.2.2.2). Controllers interface directly to the field devices reading data from sensors and sending commands to actuators.
- Level 2 (**Control Systems, or Area Control**): contains computer systems used to supervising and monitoring the physical process: they provide a **Human-Machine Interface** (HMI, 2.2.3.2) and *Engineering Workstations* (EW) for operator control.
- Level 3 (**Manufacturing/Site Operations, or Operations/Control**): comprises systems used to manage the production workflow for plant-wide control: they collate informations from the previous levels and store them in Data Historian servers.

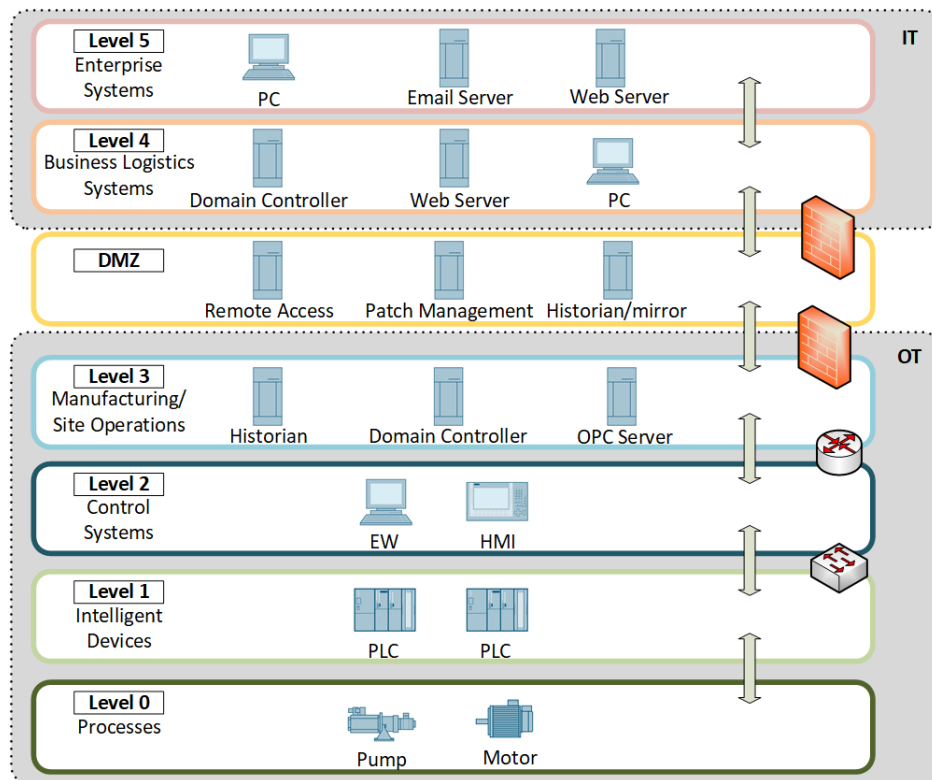


Figure 2.1: ICS architecture schema

- Industrial Demilitarized Zone (DMZ):** intermediate level that connects the *Operational Technology* (OT) part (levels 0-3) with the *Information Technology* (IT) part of the system (levels 4 and 5). Communication takes place indirectly through services such as *proxy servers* and *remote access servers*, which act as intermediaries between the two environments.
- Level 4 (Business Logistics Systems, or Business Planning/Logistics):** collect and aggregates data from the Manufacturing/Site Operations level overseeing the IT-related activities to generate **reporting** to the Enterprise System layer. At this layer we can find application and e-mail servers, and *Enterprise Resource Planning* (ERP) systems.
- Level 5 (Enterprise Systems):** represents the enterprise network, used for the business-to-business activities and for business-to-client pur-

pose services. At Enterprise Systems level are typical IT services such as mail servers, web servers and all the systems used to manage the ongoing process.

As previously discussed, the gap between IT and OT is steadily narrowing. Nowadays, it is increasingly common to encounter IT elements within the OT realm. For example, desktop PCs are now frequently found in OT environments, and industrial devices are interconnected using standard IT communication protocols like TCP and UDP.

2.2 Operational Technology Networks

Operational Technology primarily encompasses the **tangible aspects** of Industrial Control Systems and directly interfaces with the physical processes of the monitored systems. Its main purpose is to **manage and control the procedures** involved in creating and correcting physical value in various equipment.

This section will focus on the key aspects and components of Operational Technology network, with specific reference to the first four levels of the Purdue model previously seen.

2.2.1 Field I/O Devices Layer

This level concerns all aspects related to the physical environment and the physical elements that are part of it, which have the ability to actively influence the environment.

These physical elements are represented by **Field Devices**, i.e., **sensors** and **actuators** used to collect data from the process and control it: sensors are the elements responsible for reading specific values related to the physical environment (e.g., the level of a liquid), while actuators change its behavior and characteristics (e.g., opening or closing a valve to make the liquid flow).

Examples of field devices include temperature sensors, pressure sensors, valves and pumps.

2.2.2 Controller Network Layer

Controller Network layer includes devices that handle data from and to the *Field I/O Devices* layer. This kind of device is capable of gathering data from sensors, updating its internal state, and activating actuators (for example opening or close a pump that controls the level of a tank), making decisions based on a customized program, known as its control logic. Commonly found within this layer are *Programmable Logic Controllers* (PLCs) and *Remote Terminal Units* (RTUs): in the upcoming sections, we will examine these elements in detail.

2.2.2.1 Programmable Logic Controllers

A *Programmable Logic Controller* (PLC) is a **small and specialized industrial computer** having the capability of controlling complex industrial and manufacturing processes [7].

Compared to relay systems and personal computers, PLCs are optimized for control tasks and industrial environments: they are rugged and designed to withstand harsh conditions such as dust, vibrations, humidity and temperature: they have more reliability than personal computers, which are more prone to crash, and they are more compact and require less maintenance than a relay system.

Furthermore, I/O interfaces are already on the controller, so PLCs are easier to expand with additional I/O modules (if in a rack format) to manage more inputs and outputs, without reconfiguring hardware as in relay systems when a reconfiguration occurs.

PLCs are more *user-friendly*: they are not intended (only) for computer programmers, but designed for engineers with a limited knowledge in programming languages: control program can be entered with a simple and intuitive language based on logic and switching operations instead of a general-purpose programming language (*i.e.* C, C++, ...).

PLC Architecture The basic hardware architecture of a PLC consists of these elements [8]:

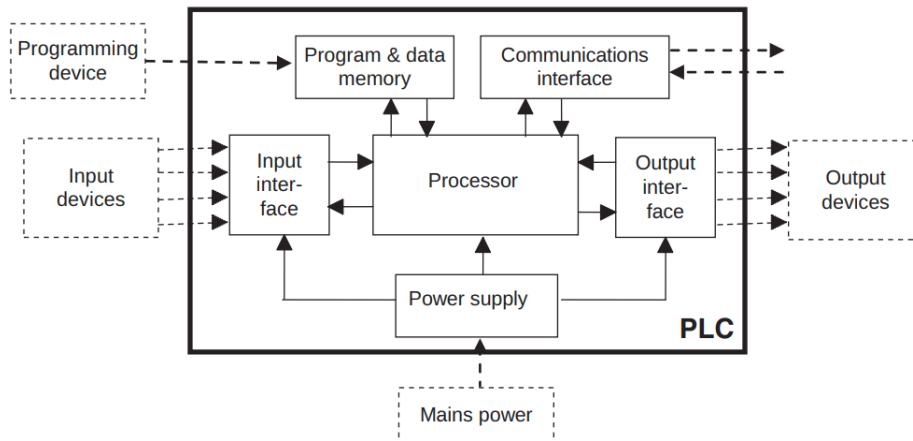


Figure 2.2: PLC architecture

- **Processor unit (CPU):** contains the microprocessor. This unit interpretes the input signals from I/O modules, executes the control program stored in the Memory Unit and sends the output signals to the I/O Modules. The processor unit also sends data to the Communication interface, for the communication with additional devices.
- **Power supply unit:** converts AC voltage to low DC voltage.
- **Programming device:** is used to store the required program into the memory unit.
- **Memory Unit:** consists in RAM memory and ROM memory. RAM memory is used for storing data from inputs, ROM memory for storing operating system, firmware and user program to be executed by the CPU.
- **I/O modules:** provide interface between sensors and final control elements (actuators).
- **Communications interface:** used to send and receive data on a network from/to other PLCs.

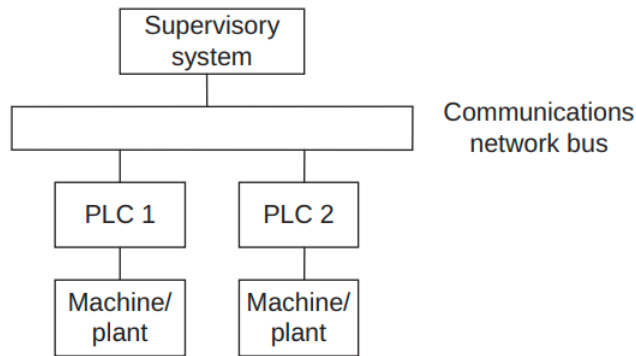


Figure 2.3: PLC communication schema

155 **PLC Programming** Two different programs are executed in a PLC: the
 156 **operating system** and the **user program**.

157 The operating system tasks include executing the user program, man-
 158 aging memory areas and the *process image table* (memory registers where
 159 inputs from sensors and outputs for actuators are stored).

160 The user program needs to be uploaded on the PLC via the program-
 161 ming device and runs on the process image table in *scan cycles*: each scan
 162 is made up of three phases [9]:

- 163 1. reading inputs from the process images table
- 164 2. execution of the control code and computing the physical process
 165 evolution
- 166 3. writing output to the process image table to have an effect on the
 167 physical process. At the end of the cycle, the process image table is
 168 refreshed by the CPU

169 Standard PLCs **programming languages** are basically of two types:
 170 **textuals** and **graphicals**. Textual languages include languages such as
 171 *Instruction List* (IL) and *Structured Text* (ST), while *Ladder Diagrams* (LD),
 172 *Function Block Diagram* (FBD) and *Sequential Function Chart* (SFC) belong
 173 to the graphical languages.

Graphical languages are more simple and immediate comparing to the textual ones and are preferred by programmers because of their features and simplicity, in particular the **Ladder Logic programming** (see Figure 2.4 for a comparison).

```

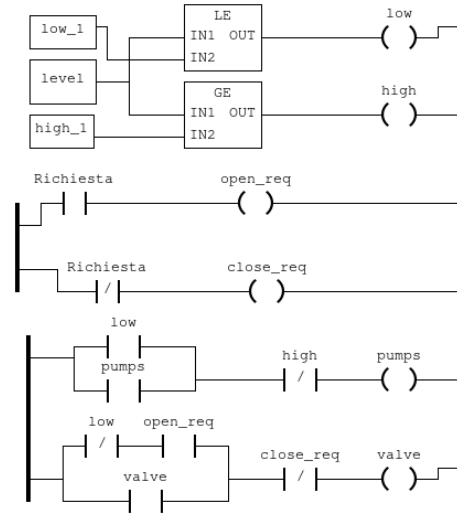
PROGRAM PLC1
VAR
  level AT %IW0 : INT;
  Richiesta AT %QX0.2 : BOOL;
  request AT %IW1 : INT;
  pumps AT %QX0.0 : BOOL;
  valve AT %QX0.1 : BOOL;
  low AT %MX0.0 : BOOL;
  high AT %MX0.1 : BOOL;
  open_req AT %MX0.3 : BOOL;
  close_req AT %MX0.4 : BOOL;
  low_l AT %MW0 : INT := 40;
  high_l AT %MW1 : INT := 80;
END_VAR
VAR
  LE3_OUT : BOOL;
  GE7_OUT : BOOL;
END_VAR

LE3_OUT := LE(level, low_l);
low := LE3_OUT;
GE7_OUT := GE(level, high_l);
high := GE7_OUT;
open_req := Richiesta;
close_req := NOT(Richiesta);
pumps := NOT(high) AND (low OR pumps);
valve := NOT(close_req) AND (open_req AND NOT(low) OR valve);
END_PROGRAM

CONFIGURATION Config0
RESOURCE Res0 ON PLC
TASK task0[INTERVAL := T#20ms,PRIORITY := 0];
PROGRAM instance0 WITH task0 : PLC1;
END_RESOURCE
END_CONFIGURATION

```

(a) Example of ST programming



(b) Example of Ladder Logic

Figure 2.4: Comparison between ST language and Ladder Logic

PLC Security PLCs were originally designed to operate as closed systems, not connected and exposed to the outside world via communication networks: the question of the safety of these systems, therefore, was not a primary aspect. The advent of Internet has brought undoubted advantages, but has introduced problems relating to the safety and protection of PLCs from external attacks and vulnerabilities.

Indeed, a variety of different communication protocols used in ICSs are designed to be efficient in communications, but do not provide any security measure i.e. confidentiality, authentication and data integrity, which makes these protocols vulnerable against many of the IT classic attacks such as *Replay Attack* or *Man in the Middle Attack*.

Countermeasures to enhance security in PLC systems may include [10]:

- 190 • protocol modifications implementing **data integrity, authentication**
191 and **protection** against *Replay Attacks*
- 192 • use of *Intrusion Detection and Prevention Systems* (IDP)
- 193 • creation of *Demilitarized Zones* (DMZ) on the network

194 In addition to this, keeping the process network and Internet separ-
195 ated, limiting the use of USB devices among users to reduce the risks of
196 infections, and using strong account management and maintenance poli-
197 cies are best practices to prevent attacks and threats and to avoid potential
198 damages.

199 2.2.2.2 Remote Terminal Units

200 *Remote Terminal Units* (RTUs) are computers with radio interfacing sim-
201 ilar to PLCs: they transmit telemetry data to the control center or to the
202 PLCs and use messages from the master supervisory system to control
203 connected objects [11].

204 The purpose of RTUs is to operate efficiently in remote and isolated
205 locations by utilizing wireless connections. In contrast, PLCs are designed
206 for local use and rely on high-speed wired connections. This key difference
207 allows RTUs to conserve energy by operating in low-power mode for ex-
208 tended periods using batteries or solar panels. As a result, RTUs consume
209 less energy than PLCs, making them a more sustainable and cost-effective
210 option for remote operations.

211 Industries that require RTUs often operate in areas without reliable ac-
212 cess to the power grid or require monitoring and control substations in re-
213 mote locations. These include telecommunications, railways, and utilities
214 that manage critical infrastructure such as power grids, pipelines, and wa-
215 ter treatment facilities. The advanced technology of RTUs allows these in-
216 dustries to maintain essential services, even in challenging environments
217 or under adverse weather conditions.

2.2.3 Area Control Layer

The Area Control layer encompasses hardware and software systems useful for supervising, monitoring and controlling the physical process, driving the behavior of the entire infrastructure. The layer includes systems such as *Supervisory Control and Data Acquisition* (SCADA), *Distributed Control Systems* (DCSs), that perform SCADA functions but are usually deployed locally, and engineer workstations.

2.2.3.1 Supervisory Control And Data Acquisition

Supervisory Control And Data Acquisition (SCADA) is a system of software and hardware elements that allows industrial organizations to [12]:

- Control industrial processes locally or at remote locations;
- Monitor, gather, and process real-time data;
- Directly interact with devices such as sensors, valves, pumps, motors, and more through human-machine interface (HMI) software;
- Record and aggregate events to send to historian server.

The SCADA software processes, distributes, and displays the data, helping operators and other employees analyze the data and make important decisions.

2.2.3.2 Human-Machine Interface

The *Human-Machine Interface* (HMI) is the hardware and software interface that operators use to monitor the processes and interact with the ICS. A HMI shows the operator and authorized users information about system status and history; it also allows them to configure parameters on the ICS such as set points and, send commands and make control decisions [13].

The HMI can be in the form of a physical panel, with buttons and indicator lights, or PC software as shown in Figure 2.5.

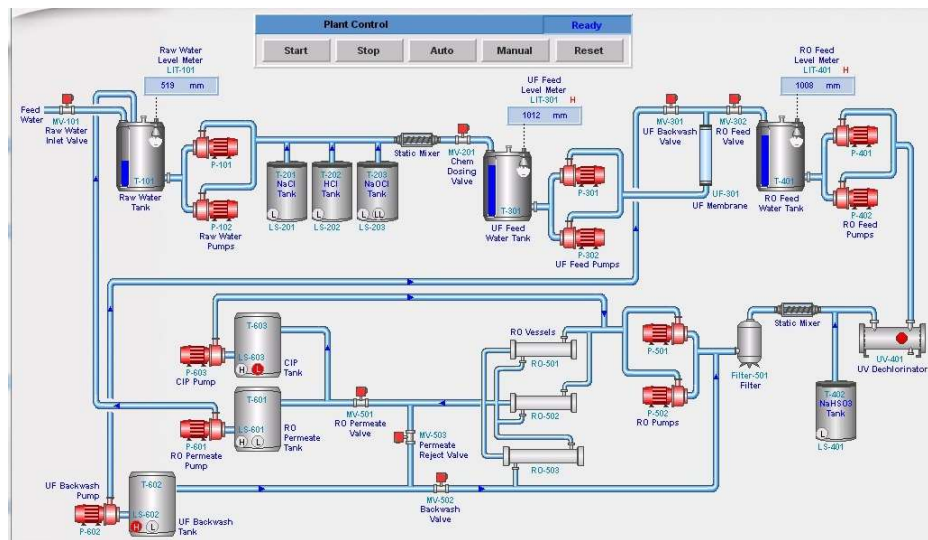


Figure 2.5: Example of HMI for a water treatment plant

2.2.4 Operations/Control Layer

Within this zone, there are specialized OT devices that are utilized to manage production workflows on the shop floor [14]. These devices include:

- *Manufacturing Operations Management (MOM)* systems, which are responsible for overseeing production operations.
- *Manufacturing Execution Systems (MES)*, which collect real-time data to optimize production processes.
- *Data Historians*, which store process data and, in modern solutions, analyze it within its contextual framework.

2.2.5 Demilitarized Zone

This zone comprises security systems like firewalls, proxies, *Intrusion Detection and Prevention systems (IDP)* and *Security Information and Event Management (SIEM)* systems which are implemented to mitigate the risk of lateral threat movement between IT and OT domains. With the rise

of automation, the need for bidirectional data flows between OT and IT systems has increased. The convergence of IT and OT in this layer can offer organizations a competitive edge. However, it's important to note that adopting a flat network approach in this context can potentially heighten cyber risks for the organization.

2.2.6 Industrial Protocols

Industrial Protocols are the networks that are used to connect the different components of the ICS and allow them to communicate with each other. Industrial Protocols can include wired and wireless networks, such as Ethernet/IP, Modbus, DNP3, Profinet and others.

As mentioned at the beginning of this Chapter, industrial systems differ from classical IT systems in the purpose for which they are designed: controlling physical processes the former, processing and storing data the latter. For this reason, ICSs require different communication protocols than traditional IT systems for real time communications and data transfer.

A wide variety of industrial protocols exists: this is because originally each vendor developed and used its own proprietary protocol. However, these protocols were often incompatible with each other, resulting in devices from different vendors being unable to communicate with each other.

To solve this problem, standards were defined with a view to allowing these otherwise incompatible device to intercommunicates.

Among all the various protocols, some have risen to prominence as widely accepted standards. These *de facto* protocols are commonly utilized in industrial systems due to their proven reliability and effectiveness. In the following sections, we will provide a brief overview of some of the most prevalent and widely used protocols in the industry.

2.2.6.1 Modbus

Modbus is a serial communication protocol developed by Modicon (now Schneider Electric) in 1979 for use with its PLCs [15] and designed expressly for industrial use: it facilitates interoperability of different devices connected to the same network (sensors, PLCs, HMIs, ...) and it is also often used to connect RTUs to SCADA acquisition systems.

Modbus is the most widely used communication protocol among industrial systems because it has several advantages:

- simplicity of implementation and debugging
- it moves raw bits and words, letting the individual vendor to represent the data as it prefers
- it is, nowadays, an **open** and *royalty-free* protocol: there is no need to sustain licensing costs for implementation and use by industrial device vendors

Modbus is a **request/response** (or *master/slave*) protocol: this makes it independent of the transport layer used.

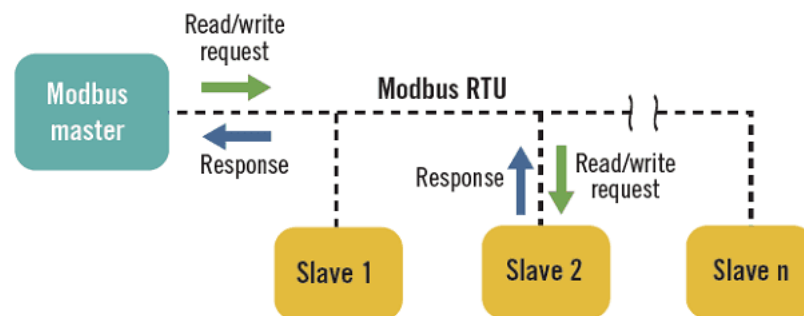


Figure 2.6: Modbus Request/Response schema

In this kind of architecture, a single device (master) can send requests to other devices (slaves), either individually or in broadcast: these slave devices (usually peripherals such as actuators) will respond to the master

by providing data or performing the action requested by the master using the Modbus protocol. Slave devices cannot generate requests to the master [16].

There are several variants of Modbus, of which the most popular and widely used are Modbus RTU (used in serial port connections) and Modbus TCP (which instead uses TCP/IP as the transport layer). Modbus TCP embeds a standard Modbus frame in a TCP frame (see Figure 2.7): both masters and slaves listen and receive data via TCP port 502.

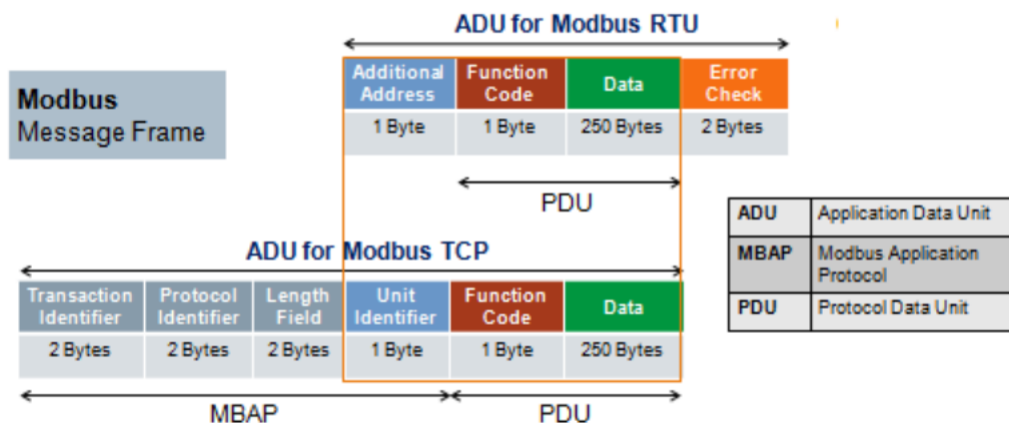


Figure 2.7: Modbus RTU frame and Modbus TCP frame

Modbus registers Modbus provides four object types, which map the data accessed by master and slave to the PLC memory:

- *Coil*: binary type, read/write accessible by both masters and slaves
- *Discrete Input*: binary type, accessible in read-only mode by masters and in read/write mode by slaves
- *Analog Input*: 16 bits in size (word), are accessible in read-only mode by masters and in read/write mode by slaves
- *Holding Register*: 16 bits in size (word), accessible in read/write mode by both masters and slaves. Holding Registers are the most commonly used registers for output and as general memory registers.

Modbus Function Codes *Modbus Function Codes* are specific codes used by the Modbus master within a request frame (see Figure 2.7) to tell the Modbus slave device which register type to access and which action to perform on it.

Two types of Function Codes exists: for data access and for diagnostic Function Codes list for data access are listed in Table 2.2:

Function Code	Description
FC01	Read Coils
FC02	Read Discrete Input
FC03	Read Holding Registers
FC04	Read Analog Input Registers
FC05	Write/Force Single Coil
FC06	Write/Force Single Holding Register
FC15	Write/Force Multiple Coils
FC16	Write/Force Multiple Holding Registers

Table 2.2: Modbus Function Codes list

Modbus Security Issues Despite its simplicity and widespread use, the Modbus protocol does not have any security feature, which exposes it to vulnerabilities and attacks.

Data in Modbus are transmitted unencrypted (*lack of confidentiality*), with no data integrity controls (*lack of integrity*) and authentication checks (*lack of authentication*), in addition to the *lack of session*. Hence, the protocol is vulnerable to a variety of attacks, such as Denial of Services (DoS), buffer overflows and reconnaissance activities.

The easiest attack to bring to the Modbus protocol, however, is **packet sniffing**: since, as mentioned earlier, network traffic is unencrypted and the data transmitted is in cleartext, it is sufficient to use a packet sniffer to capture the network traffic, read the packets and thus gather informations

about the system such as ip addresses, function codes of requests and to modify the operation of the devices.

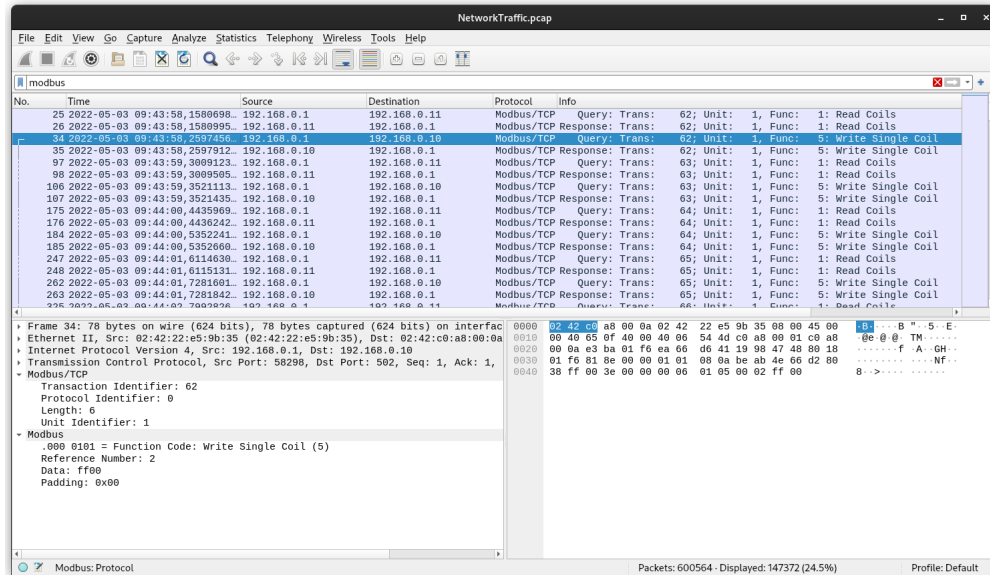


Figure 2.8: Example of packet sniffing on the Modbus protocol

To make the Modbus protocol more secure, an encapsulated version was developed within the *Transport Security Layer* (TLS) cryptographic protocol, also using mutual authentication. This version of the Modbus protocol is called **Secure Modbus** or **Modbus TLS**. In addition to this, Secure Modbus also includes X.509-type certificates to define permissions and authorisations [17].

2.2.6.2 EtherNet/IP

EtherNet/IP (where IP stands for *Industrial Protocol*) is an open industrial protocol that allows the *Common Industrial Protocol* (CIP) to run on a typical Ethernet network [18]. It is supported by ODVA [19].

EtherNet/IP uses the major Ethernet standards, such as IEEE 802.3 and the TCP/IP suite, and implements the CIP protocol stack at the upper layers of the OSI stack (see Figure 2.9). It is furthermore compatible with the main Internet standard protocols, such as SNMP, HTTP, FTP and DHCP,




Figure 2.9: OSI model for EtherNet/IP stack

and other industrial protocols for data access and exchange such as *Open Platform Communication* (OPC).

Physical and Data Link layer The use of the IEEE 802.3 standard allows EtherNet/IP to flexibly adopt different network topologies (star, linear, ring, etc.) over different connections (copper, fibre optic, wireless, etc.), as well as the possibility to choose the speed of network devices. IEEE 802.3 in addition defines at Data Link layer the *Carrier Sense Multiple Access - Collision Detection* (CSMA/CD) protocol, which controls access to the communication channel and prevents collisions.

Transport layer At the transport level, EtherNet/IP encapsulates messages from the CIP stack into an Ethernet message, so that messages can be transmitted from one node to another on the network using the TCP/IP protocol. EtherNet/IP uses two forms of messaging, as defined by CIP standard [18][20]:

- **unconnected messaging:** used during the connection establishment phase and for infrequent, low priority, explicit messages. Uncon-

nected messaging uses TCP/IP to transmit messages across the network asking for connection resource each time from the *Unconnected Message Manager* (UCMM).

- **connected messaging:** used for frequent message transactions or for real-time I/O data transfers. Connection resources are reserved and configured using communications services available via the UCMM.

EtherNet/IP has two types of message connection [18]:

- **explicit messaging:** *point-to-point* connections to facilitate *request-response* transactions between two nodes. These connections use TCP/IP service on port 44818 to transmit messages over Ethernet.
- **implicit messaging:** this kind of connection moves application-specific **real-time I/O data** at regular intervals. It uses multicast *producer-consumer* model in contrast to the traditional *source-destination* model and UDP/IP service (which has lower protocol overhead and smaller packet size than TCP/IP) on port 2222 to transfer data over Ethernet.

Session, Presentation and Application layer At the upper layers, Ethernet/IP implements the CIP protocol stack. We will discuss this protocol more in detail in Section 2.2.6.3.

2.2.6.3 Common Industrial Protocol (CIP)

The *Common Industrial Protocol* (CIP) is an open industrial automation protocol supported by ODVA. It is a **media independent** (or *transport independent*) protocol using a *producer-consumer* communication model and providing a **unified architecture** throughout the manufacturing enterprise [21][22].

CIP has been adapted in different types of network:

- 401 • **EtherNet/IP**, adaptation to *Transmission Control Protocol* (TCP) tech-
402 nologies
- 403 • **ControlNet**, adaptation to *Concurrent Time Domain Multiple Access*
404 (CTDMA) technologies
- 405 • **DeviceNet**, adaptation to *Controller Area Network* (CAN) technolo-
406 gies
- 407 • **CompoNet**, adaptation to *Time Division Multiple Access* (TDMA) tech-
408 nologies

409 **CIP objects** CIP is a *strictly object oriented* protocol at the upper layers:
410 each object of CIP has **attributes** (data), **services** (commands), **connec-**
411 **tions**, and **behaviors** (relationship between values and services of attributes)
412 which are defined in the **CIP object library**. The object library supports
413 many common automation devices and functions, such as analog and dig-
414 ital I/O, valves, motion systems, sensors, and actuators. So if the same
415 object is implemented in two or more devices, it will behave the same way
416 in each device [23].

417 **Security** [24] In EtherNet/IP implementation, security issues are the same
418 as in traditional Ethernet, such as network traffic sniffing and spoofing.
419 The use of the UDP protocol also exposes CIP to transmission route ma-
420 nipulation attacks using the *Internet Group Management Protocol* (IGMP)
421 and malicious traffic injection.

422 Regardless of the implementation used, it is recommended that certain
423 basic measures be implemented on the CIP network to ensure a high level
424 of security, such as *integrity, authentication and authorization*.

Chapter 3

State of the Art

IN CONVENTIONAL IT, the objective of an attacker is to comprehend the behavior of a program using diverse techniques in order to launch attacks that alter its execution flow, functionalities, or bypass limitations imposed by software licensing. These attack techniques involve an initial examination of the program, consisting of *static analysis* (i.e., analyzing the software without running it) and *dynamic analysis* (i.e., analyzing the program while it is running).

The outcome of these two investigative techniques is the reverse engineering of the software, which serves the purpose of identifying vulnerabilities or bugs and subsequently strategizing an attack.

In the context of OT, the notion of *reverse engineering* is not limited to its conventional definition, but also includes the concept of **process comprehension**. This term, introduced by Green et al. [25], refers to gaining a comprehensive understanding of the system's characteristics and the physical components responsible for its proper functioning.

There is limited literature available concerning the gathering and analysis of information related to the comprehension and operation of an Industrial Control System (ICS). In Section 3.1, we will provide a brief overview of the existing literature on this topic, and in the subsequent sections, we will specifically focus on one of the presented papers.

3.1 Literature on Process Comprehension

Keliris and Maniatikos The first approach presented in this section is by Keliris and Maniatikos [26]: they present a methodology for automating the reverse engineering of ICS binaries based on a *modular framework* (called ICSREF) that can reverse binaries compiled with CODESYS, one of the most popular and widely used PLC compilers, irrespective of the language used.

Yuan et al. Yuan et al. [27] propose a *data-driven* approach to discovering cyber-physical systems from data directly: to achieve this goal, they have implemented a framework whose purpose is to identify physical systems and transition logic inference, and to seek to understand the mechanisms underlying these cyber-physical systems, making furthermore predictions concerning their state trajectories based on the discovered models.

Feng et al. Feng et al. [28] developed a framework that can generate system *invariant rules* based on machine learning and data mining techniques from ICS operational data log. These invariants are then selected by systems engineers to derive IDS systems from them.

The experiment results on two different testbeds, the *Water Distribution system* (WaDi) and the *Secure Water Treatment system* (SWaT), both located at the iTrust - Center for Research in Cyber Security at the University of Singapore [29], show that under the same false positive rate invariant-based IDSs have a higher efficiency in detecting anomalies than IDS systems based on a residual error-based model.

Pal et al. Pal et al. [30] work is somewhat related to Feng et al.'s: this paper describes a data-driven approach to identifying invariants automatically using *association rules mining* [31] with the aim of generate invariants sometimes hidden from the design layout. The study has the same objective of Feng et al.'s and uses too the iTrust SwaT System as testbed.

Currently this technique is limited to only pair wise sensors and actuators: for more accurate invariants generation, the technique adopted must be capable of deriving valid constraints across multiple sensors and actuators.

Winnicki et al. Winnicki et al. [32] instead propose a different approach to process comprehension based on the **attacker's perspective** and not limited to mere *Denial of Service* (DoS): their approach is to discover the dynamic behavior of the system, in a semi-automated and process-aware way, through *probing*, that is, slightly perturbing the cyber physical system and observing how it reacts to changes and how it returns to its original state. The difficulty and challenge for the attacker is to perturb the system in such a way as to achieve an observable change, but at the same time avoid this change being seen as a system anomaly by the IDSs.

Green et al. Green et al. [25] also adopt an approach based on the attacker's perspective: this approach consists of two practical examples in a *Man in the Middle* (MitM) scenario to obtain, correlate, and understand all the types of information an attacker might need to plan an attack to alter the process while avoiding detection.

The paper shows *step-by-step* how to perform a ICS **reconnaissance**, which is fundamental to process comprehension and thus to the execution of MitM attacks.

Ceccato et al. Ceccato et al. [9] propose a methodology based on a *black box dynamic analysis* of an ICS using a reverse engineering tool to derive from the scans performed on the memory registers of the exposed PLCs and network scans an approximate model of the physical process. This model is obtained by inferring statistical properties, business process and system invariants from data logs.

The proposed methodology was tested on a non-trivial case study, using a testbed inspired by an industrial water treatment plant.

505 In the next section I will examine this latest work in more detail,
506 which will be the basis for my work and thus the subsequent chap-
507 ters of this thesis.

508 3.2 Ceccato et al.'s methodology for analyzing water- 509 tank systems

510 As previously mentioned, the paper introduces a methodology that re-
511 lies on black box dynamic analysis of an Industrial Control System (ICS).
512 This methodology involves identifying potential Programmable Logic Con-
513 trollers (PLCs) within the network and scanning the memory registers of
514 these identified controllers. The purpose of this process is to obtain an
515 approximate model of the controlled physical process.

516 The primary goal of this black box analysis is to establish a correlation
517 between the different memory registers of the targeted PLCs and funda-
518 mental concepts of an ICS such as sensors (i.e., measurements), actuators,
519 setpoints (i.e., range of values of a physical variable), network communi-
520 cations, among others.

521 To accomplish this, the various types of memory registers associated with
522 the Modbus protocol are analyzed, and attempts are made to determine
523 the nature of the data they might contain.

524 The second goal is to establish a relationship between the dynamic evolu-
525 tion of these fundamental concepts.

526 To accomplish this, Ceccato et al. have developed a prototype tool [33]
527 that facilitates the reverse engineering of the physical system. This tool
528 goes through four distinct phases:

- 529 1. **scanning of the system and data pre-processing:** this phase involves
530 gathering data to generate data logs for the registers of PLCs.
- 531 2. **graphs and statistical analysis:** The collected data is utilized to pro-
532 vide insights into the memory registers by leveraging graphs and

statistical data. This analysis approach offers valuable information about the characteristics and patterns of the memory registers.

3. **invariants inference and analysis:** generates system invariants, which are used to identify specific patterns and regularities within the system. Additionally, this phase provides users with the capability to view invariants related to a particular sensor or actuator.
4. **business process mining and analysis:** Using event logs, this phase involves reconstructing the business process that depicts how a process is executed. This step enables a thorough understanding of the sequence of events that occur in the system and how they are interrelated, ultimately leading to a comprehensive overview of the business process.

Figure 3.1 presents a schematic representation of the workflow associated with this work. In the subsequent sections of this chapter, we will provide a detailed exploration of each of these phases, offering a comprehensive understanding of the entire process.

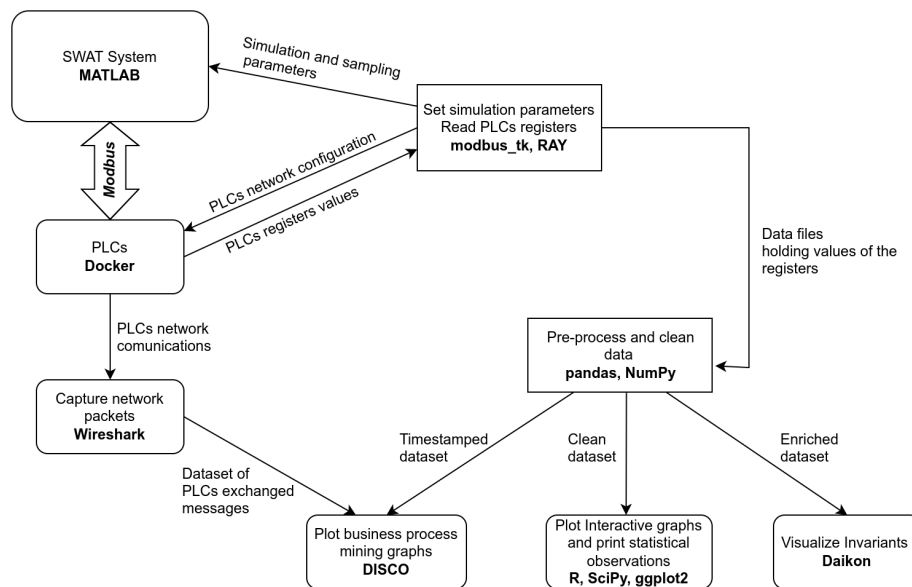


Figure 3.1: Overview

549 **3.2.1 Testbed**

550 Before delving into the description of the methodology's different phases,
 551 let's first examine the testbed utilized to evaluate this approach. The testbed
 552 employed for testing purposes is a (very) simplified rendition of the iTrust
 553 SWaT system [34], as implemented by Lanotte et al. [35]. Figure 3.2 pro-
 554 vides a graphical representation of the testbed. This simplified version
 555 comprises three stages, each governed by a dedicated PLC.

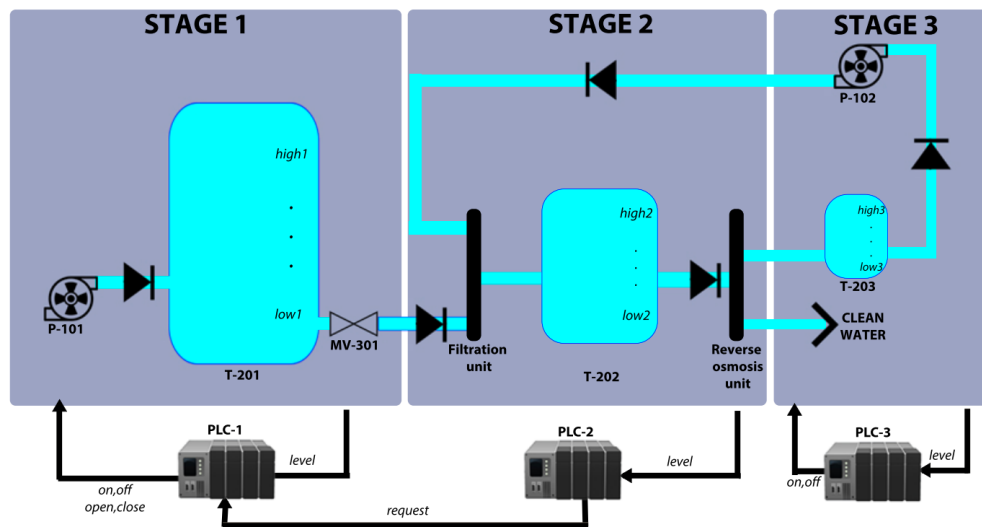


Figure 3.2: The simplified SWaT system used for running Ceccato et al. methodology

556 **Stage 1** During the initial stage, a **tank** referred to as T-201 with a capacity
 557 of 80 gallons is filled with raw water using the P-101 pump. Con-
 558 nected to the T-201 tank, the MV-301 motorized valve flushes out the
 559 accumulated water from the tank, directing it to the next stage. Ini-
 560 tially, the water flows from the T-201 tank to the *filtration unit* (which
 561 is not specifically identified by any sensor), and subsequently to a
 562 **second tank** denoted as T-202, with a capacity of 20 gallons.

563 **Stage 2** At the second stage, the water stored in tank T-202 flows into the
 564 *reverse osmosis unit* (RO), which serves as both a valve and a contin-
 565 uous water extractor. The purpose of the RO unit is to reduce organic

impurities present in the water. Subsequently, the water flows from the *RO unit* to the third and final stage of the system.

Stage 3 At the third stage, the water coming from the *RO unit* undergoes division based on whether it meets the required standards. If the water is deemed clean and meets the standards, it is directed into the distribution system. However, if the water fails to meet the standards, it is redirected to a *backwash tank* identified as T-203, which has a capacity of one gallon. The water stored in this tank is then pumped back to the stage 2 *filtration unit* using pump P-102.

As previously mentioned, each stage of the system corresponds to a dedicated PLC, namely PLC1, PLC2, and PLC3, which are responsible for controlling their respective stages. Let's briefly explore the behavior of each PLC:

PLC1 PLC1 monitors the level of tank T-201 and distinguishes three different cases based on the level readings:

1. when the level of tank T-201 reaches the defined *low setpoint low1* (which is hardcoded in the memory registers), PLC1 **opens pump P-101** and **closes valve MV-301**. This configuration allows the tank to be filled with water.
2. if the level of T-201 reaches the *high setpoint high1* (which is also hardcoded in the memory registers), then the pump **P-101 is closed**.
3. in cases where the level of T-201 is between the *low setpoint low1* and the *high setpoint high1*, PLC1 waits for a request from PLC2 to open or close the valve MV-301. If a request to open the valve MV-301 is received, water will flow from T-201 to T-202. However, if no request is received, the valve remains closed. In both situations, the pump P-101 remains closed.

594 **PLC2** PLC2 monitors the level of tank T-202 and adjusts its behavior based
595 on the water level. There are three cases to consider:

- 596 1. when the water level in tank T-202 reaches *low setpoint low2* (also
597 hardcoded in the memory registers), PLC2 sends a request to
598 PLC1 through a Modbus channel to **open valve MV-301**. This
599 request is made in order to allow the water to flow from tank T-
600 201 to tank T-202. The transmission channel between the PLCs
601 is established by copying a boolean value from a memory reg-
602 ister of PLC2 to a corresponding register of PLC1.
- 603 2. when the water level in tank T-202 reaches the *high setpoint high2*
604 value (also hardcoded in the memory registers), PLC2 sends a
605 **close request to PLC1 for valve MV-301**. This request prompts
606 PLC1 to close the valve, stopping the flow of water from tank
607 T-201 to tank T-202.
- 608 3. In cases where the water level in tank T-202 is between the low
609 and high setpoints, the valve MV-301 remains in its current state
610 (open or closed) while the tank is either filling or emptying.

611 **PLC3** PLC3 monitors the level of the T-203 backwash tank and adjusts its
612 behavior accordingly. There are two cases to consider:

- 613 1. If the water level in the backwash tank reaches the *low setpoint*
614 *low3*, **PLC3 sets pump P103 to off**. This allows the backwash
615 tank to be filled.
- 616 2. If the water level in the backwash tank reaches the *high setpoint*
617 *high3*, **PLC3 opens pump P103**. This action triggers the pump-
618 ing of the entire content of the backwash tank back to the filter
619 unit of T-202.

620 3.2.2 Scanning of the System and Data Pre-processing

621 **Scanning tool** The Ceccato et al. scanning tool is closely derived from
622 a project I did [36] for the "Network Security" and "Cyber Security for IoT"

courses taught by Professors Massimo Merro and Mariano Ceccato, respectively, in the 2020/21 academic year. The original project involved, in its first part, the recognition within a network of potential PLCs listening on the standard Modbus TCP port 502 using the Nmap module for Python, obtaining the corresponding IP addresses: then a (sequential) scan of a given range of the memory registers of the found PLCs was performed to collect the register data. The data thus collected were saved to a file in *JavaScript Object Notation* (JSON) format for later use in the second part of my project.

The scanning tool by Ceccato et. al works in a similar way, but extends what I originally did by trying to discover other ports on which the Modbus protocol might be listening (since in many realities Modbus runs on different ports than the standard one, according to the concept of *security by obscurity*) and, most importantly, by **parallelizing and distributing the scan** of PLC memory registers through the Ray module [37], specifying moreover the desired granularity of the capture. An example of raw data capture can be seen at Listing 3.1:

```
"127.0.0.1/8502/2022-05-03 12_10_00.591": {  
  "DiscreteInputRegisters": {"%IX0.0": "0"},  
  "InputRegisters": {"%IW0": "53"},  
  "HoldingOutputRegisters": {"%QW0": "0"},  
  "MemoryRegisters": {"%MW0": "40", "%MW1": "80"},  
  "Coils": {"%QX0.0": "0"}}
```

Listing 3.1: Example of registers capture

The captured data includes PLC's IP address, Modbus port and timestamp (first line), type and name of registers with their values read from the scan (subsequent lines).

The tool furthermore offers the possibility, in parallel to the memory registers scan, of **sniffing network traffic** related to the Modbus protocol using the *Man in the Middle* (MitM) technique on the supervisory control network using a Python wrapper for tshark/Wireshark [38] [39]. An example of raw data obtained with this sniffing can be seen in Listing 3.2:

```

654 Time,Source,Destination,Protocol,Length,Function Code,
655     ↳ Destination Port,Source Port,Data,Frame length on the
656     ↳ wire,Bit Value,Request Frame,Reference Number,Info
657 2022-05-03 11:43:58.158,IP_PLC1,IP_PLC2,Modbus/TCP,76,Read
658     ↳ Coils,46106,502,,76,TRUE,25,,,"Response: Trans: 62;
659     ↳ Unit: 1, Func: 1: Read Coils"

```

Listing 3.2: Example of raw network capture

Data Pre-processing The data collected by scanning the memory registers of the PLCs are then reprocessed by a Python script and converted in order to create a distinct raw dataset in *Comma Separated Value* format (CSV) for each PLC, containing the memory register values associated with the corresponding controller registers. These datasets are reprocessed again through the Python modules for **pandas** [40] and **NumPy** [41] by another script to first perform a **data cleanup**, removing all those memory registers that do not take values and are therefore useless within the system, **merged** into a single dataset, and finally **enriched** with additional data¹.

This process leads to the creation of two copies of the full dataset: one enriched with the additional data, but not timestamped, which will be used for the invariant analysis; the other unenriched, but timestamped, which will be used for business process mining.

3.2.3 Graphs and Statistical Analysis

The paper mentions the presence of a *mild graph analysis*, performed with **R** [42] at the time of data gathering to find any uncovered patterns, trends and identify measurements and/or actuator commands through the analysis of registers holding mutable values.

There is actually no trace of this within the tool: *graph analysis* and *statistical analysis* of the data contained in the PLC memory registers are in-

¹Not all additional data are calculated and entered automatically by the tool: some are manually inserted.

stead performed using the **matplotlib** libraries and statistical algorithms made available by the **SciPy** libraries [43], through two separate Python scripts (see Figure 3.3).

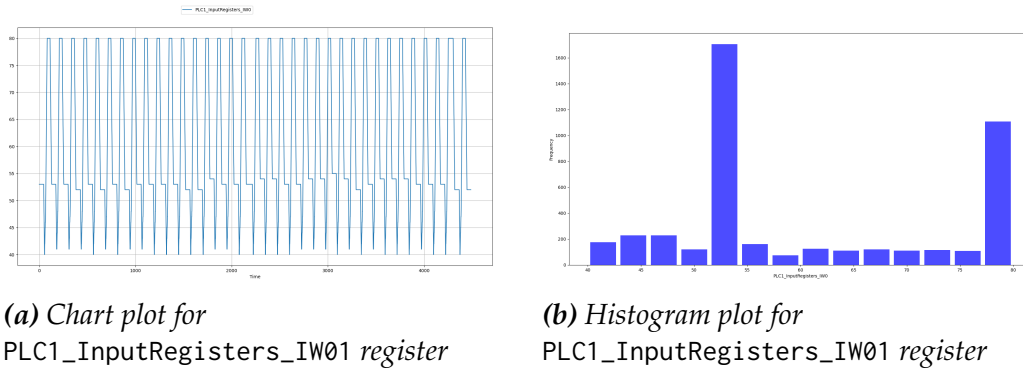


Figure 3.3: Output graphs from graph analysis

The first script plots the charts, one at the time, of certain registers entered by the user from the command line, plots in which one can see the trend of the data and get a first basic idea of what that particular register contains (a measurement, an actuation, a hardcoded setpoint, ...) and possibly the trend; the second script, instead, shows **a histogram and statistical informations** about the register entered as command-line input. These informations include:

- the mean, median, standard deviation, maximum value and minimum value
- two tests for the statistical distribution: *Chi-squared* test for uniformity and *Shapiro-Wilk* test for normality, as shown in Listing 3.3:

```

695 Chi-squared test for uniformity
696 Distance      pvalue    Uniform?
697 12488.340     0.00000000    NO
698
699 Shapiro-Wilk test for normality
700 Test statistic pvalue    Normal?
701 0.844        0.00000000    NO
702

```

```

703 Stats of PLC1_InputRegisters_IW0
704 Sample mean = 60.8881; Stddev = 13.0164; max = 80; min =
705 ↪ 40 for 4488 values

```

Listing 3.3: Statistical data for PLC1_InputRegisters_IW0 register

706 3.2.4 Invariants Inference and Analysis

707 For invariant analysis Ceccato et al. rely on **Daikon** [44], a framework
 708 to **dynamically detect likely invariants** within a program. An *invariant*
 709 is a property that holds at one or more points in a program, properties
 710 that are not normally made explicit in the code, but within assert state-
 711 ments, documentation and formal specifications: invariants are useful in
 712 understanding the behavior of a program (in our case, of the cyber physi-
 713 cal system).

714 Daikon uses *machine learning* techniques applied to arbitrary data with
 715 the possibility of setting custom conditions for analysis by using a spe-
 716 cific file [45] with a *.spinfo* extension (see Listing 3.4). The framework is
 717 designed to find the invariants of a program, with various supported pro-
 718 gramming languages, starting from the direct execution of the program
 719 itself or passing as input the execution run (typically a file in CSV format):
 720 the authors of the paper tried to apply it by analogy also to the execution
 721 runs of a cyber physical system, to extract the invariants of this system.

```

722 PPT_NAME aprogram.point:::POINT
723 VAR1 > VAR2
724 VAR1 == VAR3 && VAR1 != VAR4

```

Listing 3.4: Generic example of a .spinfo file for customizing rules in Daikon

725 Therefore, Daikon is fed with the no-timestamp enriched dataset ob-
 726 tained in the pre-processing phase (in the paper, the timestamped dataset
 727 is erroneously mentioned as input): a simple bash script launches Daikon
 728 (optionally specifying the desired condition for analysis in the *.spinfo* file),
 729 which output is simply redirected to a text file containing the general in-
 730 variants of the system (i.e., valid regardless of any custom condition speci-

fied), those generated based on the custom condition in the *.spinfo* file, and those generated based on the negation of the condition. When the analysis is finished, the user is asked to enter the name of a registry to view its related invariants.

Some examples of invariants derived from the enriched dataset may be:

- measurements bounded by some setpoint
- Actuators state changes occurred in the proximity of setpoints or, vice versa, proximity of setpoints upon the occurrence of a regular actuator state change
- state invariants of some actuator correspond to a specific trend in the evolution of the measurement (ascending, descending, or stable) or, vice versa, the measurement trend corresponds to a specific state invariant of some actuator

3.2.5 Business Process Mining and Analysis

Process mining is the analysis of operational processes based on the event log [46]: the aim of this analysis is to **extract useful informations** from the event data to **reconstruct and understand the behavior** of the business process and how it was actually performed.

Process mining for the system under consideration starts from the event logs obtained from scanning the memory registers of the PLCs and sniffing the network communications related to the Modbus protocol, described in Subsection 3.2.2 and representing the *execution trace* of the system: through a Java program, information is extracted and combined from these event logs, and the result saved in a CSV format file.

This file is fed to **Disco** [47], a commercial process mining tool, which generates an *activity diagram* similar to UML Activity Diagram and whose nodes represent the activities while the edges represent the relations be-

759 tween these activities: in Figure 3.4 we can see an example of this diagram
 760 referred to PLC2 of the testbed.

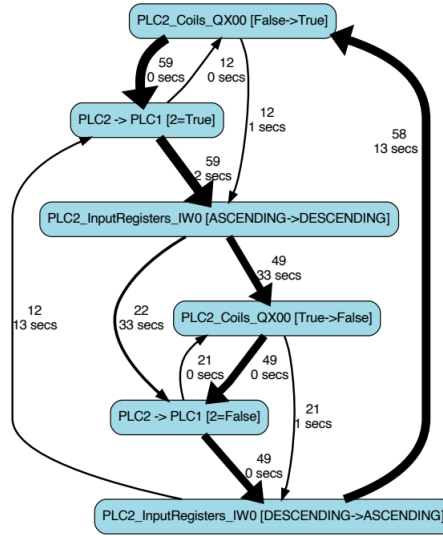


Figure 3.4: An example of Disco generated activity diagram for PLC2

761 The *business process* obtained in this way provides an **overview of the**
 762 **system** and makes it possible to **make conjectures** about its behavior, par-
 763 ticularly between changes in actuator state and measurement trends (i.e.,
 764 a given change in state of some actuators corresponds to a specific mea-
 765 surement trend and vice versa), and with the possibility of **establishing**
 766 **causality** between Modbus communications and state changes within the
 767 physical system.

768 3.2.6 Application

769 In this section we will see how the black box analysis presented above
 770 in its various phases is applied in practice, using the testbed described in
 771 Subsection 3.2.1. The methodology supports a **top-down approach**: that
 772 is, we start with an overview of the industrial process and then gradually
 773 refine our understanding of the process by descending to a higher and
 774 higher level of detail based on the results of the previous analyses and

focusing on the most interesting parts of the system for further in-depth analysis.

Data Collection and Pre-processing According to what is described in the paper, the data gathering process lasted six hours, with a granularity of one data point per second (a full system cycle takes approximately 30 minutes). Each datapoint consists of 168 attributes (55 registers plus a special register concerning the tank slope of each PLC) after the enrichment. In addition, IP addresses are automatically replaced by an abstract name identified by the prefix PLC followed by a progressive integer (PLC1, PLC2, PLC3), in order to make reading easier.

Graphs and Statistical Analysis It is unclear from the paper where exactly the information that follows was derived (graph analysis? Statistical analysis? Human reading of the dataset?), however, three properties about the contents of the registers were discovered:

Property 1: PLC1_MemoryRegisters_MW0, PLC1_MemoryRegisters_MW1, PLC2_MemoryRegisters_MW0, PLC2_MemoryRegisters_MW1, PLC3_MemoryRegisters_MW0 and PLC3_MemoryRegisters_MW1

registers contain constant integer values (40, 80, 10, 20, 0, 10 respectively)². We may speculate that they may be (relative) hardcoded **setpoints**.

Property 2: PLC1_Coils_QX01, PLC1_Coils_QX02, PLC2_Coils_QX01, PLC2_Coils_QX02, PLC3_Coils_QX01 and PLC3_Coils_QX03 contain mutable binary (Boolean) values. We can assume that these registers can be associated with the **actuators** of the system.

Property 3: PLC1_InputRegisters_IW0, PLC2_InputRegisters_IW0 and PLC3_InputRegisters_IW0 registers contain mutable values.

²From my tests on the original tool and dataset, the PLC3_MemoryRegisters_MW0 register is deleted during the *pre-processing* phase, as it is recognized as an unused register because of the constant value "0" it takes on. This leads me to assume that the properties are derived from a human read of the dataset prior to the *pre-processing* phase.

801 *Property 3* suggests that those registers might contain **values related to**
 802 **measurements**: it is therefore necessary to investigate further to see if the
 803 conjecture (referred to as *Conjecture 1* in the paper) is correct.




Figure 3.5: Execution traces of InputRegisters_IW0 on the three PLCs

804 The graph analysis of the InputRegisters_IW0 registers of the three
 805 PLCs (summarized in Figure 3.5 with a single plot) not only seems to con-
 806 firm the conjecture, but also allows the measurements to be correlated with
 807 the contents of the MemoryRegisters_MW0 and MemoryRegisters_MW1 regis-
 808 ters to the measurements, which represent the **relative setpoints of the**
 809 **measurements**.

810 Hence, we have *Conjecture 2* described in the paper referring to the relative
 811 setpoints:

812

813 *Conjecture 2:*

- 814 - 40 and 80 are the relative setpoints for PLC1_InputRegisters_IW0
- 815 - 10 and 20 are the relative setpoints for PLC2_InputRegisters_IW0
- 816 - 0 and 9 are the relative setpoints for PLC3_InputRegisters_IW0

817 Further confirmation of this conjecture may come from statistical anal-
 818 ysis. Indeed, in the example in Listing 3.1, some statistical data are given
 819 for the register PLC1_InputRegisters_IW0, including the maximum value

and the minimum value: these values are, in fact, 80 and 40 respectively.

Business Process Mining and Analysis With Business Process Mining, the authors aim to **visualize and highlight relevant system behaviors** by relating PLC states and Modbus commands.

Through analysis of the activity diagrams shown in Figure 3.6, drawn through Disco, we derive the following properties and conjectures:

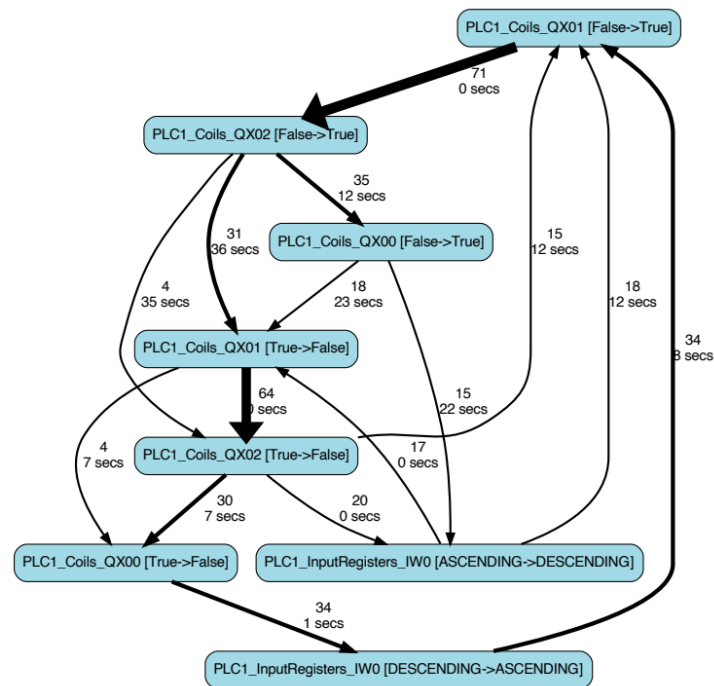
Property 4: PLC2 sends messages to PLC1 (see Figure 3.6b) which are recorded to PLC1_Coils_QX02.

Conjecture 3: PLC2_Coils_QX00 determines the trend in tank T-202 (Figure 3.6b).

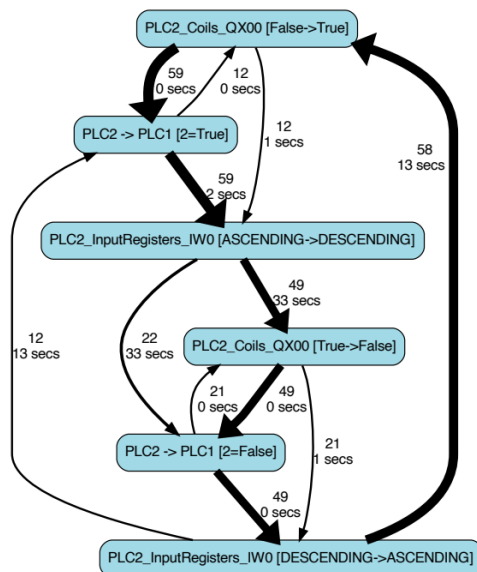
When this register is set to *True*, the input register PLC2_InputRegisters_IW0 related to the tank controlled by PLC2 starts an **ascending trend**; vice versa, when the coil register is set to *False*, the input register starts a **descending trend**.

Conjecture 4: If PLC1_Coils_QX00 change his value to True, trend in tank T-201, related to PLC1_InputRegisters_IW0 and controlled by PLC1, become **ascending** (see Figure 3.6a)

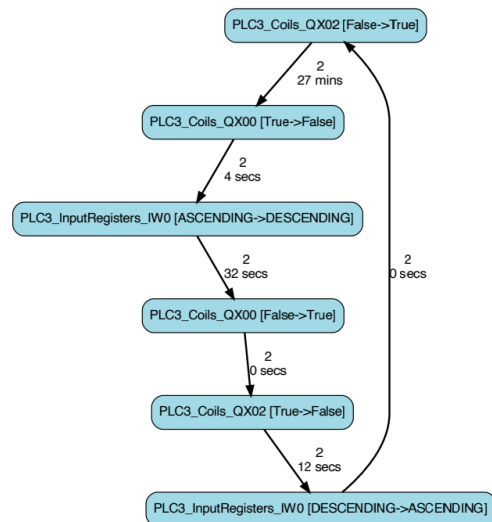
Conjecture 5: PLC3_Coils_QX00 starts a **decreasing trend** in tank T-203, related to PLC3_InputRegisters_IW0 and controlled by PLC3, whereas PLC3_Coils_QX02 starts an **increasing trend** on the tank (see Figure 3.6c)



(a) States in PLC1



(b) States and Modbus command in PLC2



(c) States in PLC3

Figure 3.6: Business process with states and Modbus commands for the three PLCs

Invariants Inference and Analysis The last phase of the analysis of the example industrial system is invariant analysis, performed through Daikon framework. At this stage, an attempt will be made to confirm what has been seen previously and to derive new properties of the system based on the results of the Daikon analysis.

To get gradually more and more accurate results, the authors presumably performed more than one analysis with Daikon, including certain rules within the *splitter information file* (see Section 3.2.4 and Listing 3.4) based on specific conditions placed on the measurements, for example, the level of water contained in a tank. Given moreover the massive amount of invariants generated by Daikon's output, it is not easy to identify and correlate those that are actually useful for analysis: this must be done manually.

However, it was possible to have confirmation of the conjectures made in the previous stages of the analysis: starting with the setpoints, analyzing the output of the invariants returned by Daikon³ reveals that

```

PLC1_InputRegisters_IW0 >= PLC1_MemoryRegisters_MW0 == 40.0
PLC1_InputRegisters_IW0 <= PLC1_MemoryRegisters_MW1 == 80.0
PLC2_InputRegisters_IW0 >= PLC2_MemoryRegisters_MW0 == 10.0
PLC2_InputRegisters_IW0 <= PLC2_MemoryRegisters_MW1 == 20.0
PLC3_InputRegisters_IW0 >= PLC3_MemoryRegisters_MW0 == 0.0
PLC3_InputRegisters_IW0 <= PLC3_MemoryRegisters_MW1 == 9.0

```

i.e., that the MemoryRegisters_MW0 and MemoryRegisters_MW1 registers of each PLC contain the **absolute minimum and maximum setpoints**, respectively (*Property 5*).

There is also a confirmation regarding *Property 4*: from the computed invariants it can be seen that

³The invariants shown here are a manual summary and derivation of those actually returned in output by Daikon. I will discuss this more in Section 3.2.7

870

871 PLC1_Coils_QX01 == PLC1_Coils_QX02 == PLC2_Coils_QX00

872

873 and from this derive that there is a **communication channel between PLC2**
 874 **and PLC1**, where the value of PLC2_Coils_QX00 is copied to PLC1_Coils_QX01
 875 and PLC1_Coils_QX02 (*Property 6*).

876 Regarding the **relationships between actuator state changes and mea-**
 877 **surement trends**, invariant analysis yields the results summarized in the
 878 following rules:

879 *Property 7:* Tank T-202 level *increases* iif PLC1_Coils_QX01 == True. Oth-
 880 erwise, if PLC1_Coils_QX01 == False will be *non-increasing*.

881 This is because if the coil is *True* the condition

882 PLC2_InputRegisters_IW0 == PLC2_MemoryRegisters_MW0 == 20.0 && PLC2_slope > 0

883 is verified. On the opposite hand, if the coil is *False*, the condition

884 PLC2_InputRegisters_IW0 == PLC2_MemoryRegisters_MW0 == 20.0 && PLC2_slope <= 0 is verified. The
 885 *slope* is an auxiliary attribute indicating the trend of the measurement: in-
 886 creasing if > 0, decreasing if < 0, stable otherwise.

887 *Property 8:* Tank T-201 level *increases* iif PLC1_Coils_QX00 == True. On the
 888 other hand, if PLC1_Coils_QX00 == False and if PLC1_Coils_QX01 ==
 889 True the level will be *non-decreasing*.

890 *Property 9:* Tank T-203 level *decreases* iif PLC3_Coils_QX00 == True. It will
 891 be *non-decreasing* if PLC1_Coils_QX00 == False.

892 The last two properties concern the **relationship between actuator state**
 893 **changes and the setpoints**: it is intended to check what happens to the
 894 actuators when the water level reaches one of these setpoints. From the
 895 analysis of the relevant invariants, the following properties are derived:

896 *Property 10:* Tank T-201 reaches the upper absolute setpoint when
 897 PLC1_Coils_QX00 changes its state from *True* to *False*. If the coil changes
 898 from *False* to *True*, the tank reaches its absolute lower setpoint.

Property 11: Tank T-203 reaches the upper absolute setpoint when PLC3_Coils_QX00 changes its state from *True* to *False*. If the coil changes from *False* to *True*, the tank reaches its absolute lower setpoint.

3.2.7 Limitations

The methodology proposed by Ceccato et al. is certainly valid and offers a good starting point for approaching the reverse engineering of an industrial control system from the attacker's perspective, while also providing a tool to perform this task.

The limitations of this approach, however, all lie in the tool mentioned above and also in the testbed described in Section 3.2.1. In this section I will explain which are the criticisms of each phase, while in Chapter 4 I will formulate proposals to improve and make this methodology more efficient.

General Criticism The general critical aspects of the application of this approach are many: the primary one concerns the fact that the proposed tool seems to be built specifically for the testbed used and that it is not applicable to other contexts, even to the same type of industrial control system (water treatment systems, in this case).

What severely limits the analysis performed with the tool implemented by Ceccato et al. is the use of *ad hoc* solutions and *a posteriori* interventions done manually on the datasets after the data gathering process: I will discuss this last aspect in more detail later.

Moreover, there is the presence of many *hardcoded* variables and conditions within the scripts: this makes the system unconfigurable and unable to properly perform the various stages of the analysis as errors can occur due to incorrect data and mismatches with the system under analysis.

Having considered, furthermore, only the Modbus protocol for network communications between the PLCs is another major limiting factor and does not help the methodology to be adaptable to different systems

928 communicating with different protocols (sometimes even multiple ones
929 on the same system).

930 Let us now look at the limitations and critical aspects of each phase.

931 **Testbed** The testbed environment used by Ceccato et. al is entirely simu-
932 lated, from the physical system to the control system. The PLCs were built
933 with **OpenPLC** [48] in a Docker environment [49], while the physics part
934 was built through **Simulink** [50].

935 OpenPLC is an open source cross-platform software that simulates the
936 hardware and software functionality of a physical PLC and also offers a
937 complete editor for PLC program development with support for all stan-
938 dard languages: *Ladder Logic* (LD), *Function Block Diagram* (FBD), *Instruc-*
939 *tion List* (IL), *Structured Text* (ST), and *Sequential Function Chart* (SFC).

940 It is for sure an excellent choice for creating a zero-cost industrial or home
941 automation and *Internet of Things* (IoT) system that is easy to manage via a
942 dedicated, comprehensive and functional web interface. In spite of these
943 undoubted merits, however, there are (at the moment) **very few supported**
944 **protocols**: the main one and also referred to in the official documentation
945 is **Modbus**, while the other protocol is DNP3.

946 The biggest problem with the testbed, however, is not with the con-
947 troller part, but with the **physical part**: first of all, it must be said that
948 although this is something purely demonstrative even though it is fully
949 functional, the implemented Simulink model is really **oversimplified** com-
950 pared to the iTrust SWaT system, which itself is a scaled-down version of a
951 real water treatment plant. In fact, in the entire system there are only three
952 actuators, two of which are connected to the same tank and controlled by
953 the same PLC, and sensors related only to the water level in the system's
954 tanks: in a real system there are many more *field devices*, which can mon-
955 itor and control other aspects of the system beyond the mere contents of
956 the tanks. Consider, for example, measuring and controlling the chemicals
957 in the water, the pressure of the liquid in the filter unit, or more simply the


amount of water flow at a given point or time.

All these must be considered and represent a number of additional variables that makes analysis and consequently reverse engineering of the system more difficult.


The second critical aspect concerns the **simulation of the physics of the liquid** inside the tanks: Simulink does not consider the fact that inside a tank that is filling (emptying) the liquid in it undergoes **fluctuations** which cause the level sensor not to see the water level constantly increasing (decreasing) or at most being stable at each point of detection. Figure 3.7 exemplifies more clearly with an example the concept just expressed: these oscillations cause a **perturbation** in the data.

This issue leads to the difficulty, on a real physical system, of **correctly calculating the trend of a measurement** by using the slope attribute: if this was obtained with a too low granularity, the trend will be oscillating between increasing and decreasing even when in reality this would be in general increasing (decreasing) or stable; on the other hand, if the slope was obtained with a too high granularity there is a loss of information and the trend may be "flattened" with respect to reality.

In the present case, the slope in the Simulink model was calculated statically *point-to-point*, thus with a granularity of one second: an averagely careful reader will have already guessed that this granularity is inapplicable to the real system in Figure 3.7b. As we will later see, we need to **operate on the data perturbations** to be able to obtain a suitable granularity and a correct calculation of the slope and consequently of the measurement trend.



(a)



(b)

Figure 3.7: Water physics compared: simulated physics in the Simulink model (a) and physics in a real system (iTrust SWaT) (b). Fluctuations in the tank level in (b), almost completely absent in (a), can be appreciated.

983 **Pre-processing** In the pre-processing phase, the authors make use of a
 984 Python script to merge all the datasets of the individual PLCs into a single
 985 dataset, remove the (supposedly) unused registers, and finally enrich the
 986 obtained dataset with additional attributes. These attributes are:

- 987 • the previous value of all registers

- some additional relative setpoints named `PLCx_Max_safety` and `PLCx_Min_safety` (where x is the PLC number), which represent a kind of alert on reaching the maximum and minimum water levels of the tanks
- the measurement slope, that is, the trend of the water level in the system tanks along the system cycles.

Merging the datasets of all individual PLCs into a single dataset representing the entire system can be a sound practice if the system to be analyzed is (very) small as is the testbed analyzed here, consisting of a few PLCs and especially a few registers. If, however, the complexity of the system increases, this type of merging can become counterproductive and make it difficult to analyze and understand the data obtained in subsequent steps.

In short, there is no possibility to analyze only a subsystem and thus make the analysis faster and more understandable. Moreover, a data gathering can take up to days, and the analyst/attacker may need to make an analysis of the system isolating a precise time range, ignoring everything that happens before and/or after: all of this, with the tool we have seen, cannot be done.

Regarding the additional attributes, looking at the code of the script that performs the enrichment, I observed that **some attributes were manually inserted** after the merging phase: I am referring in particular to the attributes `PLCx_Max_safety` and `PLCx_Min_safety`, whose references were moreover hardcoded into the script, and the *slope* whose calculation method I mentioned in the previous paragraph about the testbed limitations.

In the end, only the attribute *prev* related to the value at the previous point of the detection is inserted automatically for all registers, moreover without the possibility to choose whether this attribute should be extended to all registers or only to a part.

1018 **Graphs and Statistical Analysis** Describing the behavior of graphical
1019 analysis in Section 3.2.3 I had already mentioned that only one register plot
1020 at a time was shown and not, for example, a single window containing the
1021 charts of all registers entered by the user as input from the command line,
1022 such as in Figure 3.5.

1023 While displaying charts for individual registers still provides useful in-
1024 formation about the system such as the distinction between actuators and
1025 measurements and the general trend of the latter, single display does not
1026 allow one to catch, or at least makes it difficult, the relationship that exists
1027 between actuators and measurements, where it exists, because a view of
1028 the system as a whole is missing.

1029 In this way, the risk is to make conjectures about the behavior of the system
1030 that may prove to be at least imprecises, if not inaccurates.

1031 On the other hand, regarding the statistical analysis, two observations
1032 need to be made: the first is that for the given system, I personally was
1033 unable to appreciate the usefulness of the generated histogram, as it does
1034 not provide any particular new information that has not already been ob-
1035 tained from the graphical analysis (except maybe something marginal);
1036 the second observation is that precisely the plot of the histogram "hides"
1037 the statistical informations obtained: these are in fact shown on the ter-
1038 minal from which the script is launched, but to an uncared eye or one
1039 unfamiliar with the script's behavior they can easily be interpreted as sim-
1040 ple debugging output, since at the same time the window containing the
1041 histogram plot is shown. In general, however, little statistical information
1042 is provided.

1043 **Business Process Mining and Analysis** Concerning the data mining,
1044 this is a purely *ad hoc solution*, designed to work under special conditions:
1045 first, the timestamped dataset of the physical process and the one obtained
1046 after the packet sniffing operation of Modbus traffic on the network need
1047 to be synchronized and have the same granularity, in this case one event
1048 per second.

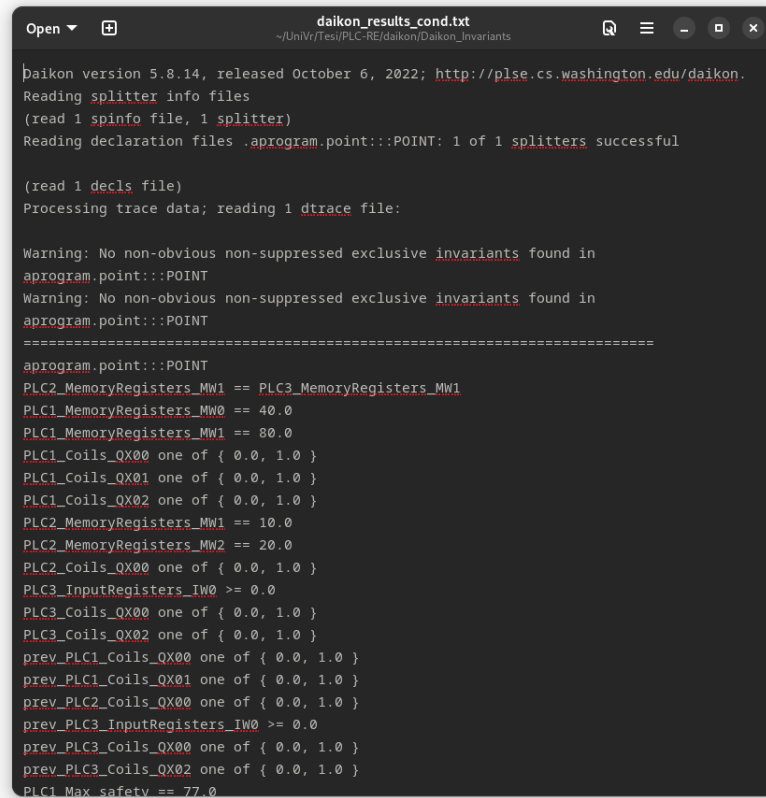
It is relatively easy, therefore, to find correspondences between Modbus commands sent over the network and events occurring on the physical system, such as state changes in actuators, due in part to the fact that the number of communications over the network is really small (see Section 3.2.1). In a real system, network communications are much more numerous and involve many more devices even in the same second: finding the exact correspondence with what is happening in the cyber physical system becomes much more difficult.

Since this is, as mentioned, an *ad hoc* solution, only the Modbus protocol is being considered: as widely used as this industrial protocol is, other protocols that are widely used such as EtherNet/IP (see Section 2.2.6.2) should be considered in order to extend the analysis to other industrial systems that use a different communication network.

The other limiting aspect of the business process mining phase is the **process mining software** used to generate the activity diagram. As mentioned in Section 3.2.5, the process mining software used by Ceccato et al. is **Disco**: this is commercial software, with an academic license lasting only 30 days (although free of charge), released for Windows and MacOS operating systems only, which makes its use under Linux systems impossible except by using emulation environments such as Wine. For what is my vision and training as a computer scientist, it would have been preferable to use a *cross-platform, freely licensed open source* software alternative to Disco: one such software could have been **ProM Tools** [51], a framework for process mining very similar to Disco in functionality, but fitting the criteria just described, or use Python libraries such as **PM4PY** [52], which offer ready-to-use algorithms suitable for various process mining needs.

Invariants Inference and Analysis The limitation in this case is principally Daikon: this software is designed to compute the invariants of a software from its live execution or from a file containing its execution flow, not to find the invariants of a cyber physical system. Since there are currently

no better consolidated alternatives for inferring invariants, however, an attempt was still made to use Daikon as best as possible.



```

daikon version 5.8.14, released October 6, 2022; http://plse.cs.washington.edu/daikon.
Reading splitter info files
(read 1 spinfo file, 1 splitter)
Reading declaration files .aprogram.point::POINT: 1 of 1 splitters successful

(read 1 decls file)
Processing trace data; reading 1 dtrace file:

Warning: No non-obvious non-suppressed exclusive invariants found in
aprogram.point::POINT
Warning: No non-obvious non-suppressed exclusive invariants found in
aprogram.point::POINT
=====
aprogram.point::POINT
PLC2_MemoryRegisters_MW1 == PLC3_MemoryRegisters_MW1
PLC1_MemoryRegisters_MW0 == 40.0
PLC1_MemoryRegisters_MW1 == 80.0
PLC1_Coils_QX00 one of { 0.0, 1.0 }
PLC1_Coils_QX01 one of { 0.0, 1.0 }
PLC1_Coils_QX02 one of { 0.0, 1.0 }
PLC2_MemoryRegisters_MW1 == 10.0
PLC2_MemoryRegisters_MW2 == 20.0
PLC2_Coils_QX00 one of { 0.0, 1.0 }
PLC3_InputRegisters_IW0 >= 0.0
PLC3_Coils_QX00 one of { 0.0, 1.0 }
PLC3_Coils_QX02 one of { 0.0, 1.0 }
prev_PLC1_Coils_QX00 one of { 0.0, 1.0 }
prev_PLC1_Coils_QX01 one of { 0.0, 1.0 }
prev_PLC2_Coils_QX00 one of { 0.0, 1.0 }
prev_PLC3_InputRegisters_IW0 >= 0.0
prev_PLC3_Coils_QX00 one of { 0.0, 1.0 }
prev_PLC3_Coils_QX02 one of { 0.0, 1.0 }
PLC1_Max_safety == 77.0

```

Figure 3.8: Example of Daikon's output

The biggest problem with Daikon applied to the computation of invariants of an industrial system is the difficult reading of the resulting output: the software in fact returns a very long list of invariants, one invariant per line, many of no use and without correlating invariants that may have common features or deriving additional information from them. The process of screening and recognizing the significant invariants, as well as the correlation between them, must be done by a human: certainly not an easy task given the volume of invariants one could theoretically be faced with (hundreds and hundreds of invariants). An example of Daikon's output

1091 can be seen in Figure 3.8.

1092 The bash script used in this phase of the analysis does not help at all
1093 in deriving significant invariants more easily: it merely launches Daikon
1094 and saves its output to a text file by simply redirecting the stdout to file.
1095 No data reprocessing is done during this step. In addition, if a condition
1096 is to be specified to Daikon before performing the analysis, it is necessary
1097 each time to edit the .spinfo file by manually entering the desired rule, an
1098 inconvenient operation when multiple analyses are to be performed with
1099 different conditions each time.

Chapter 4

A Framework to Improve Ceccato et al.'s Work.

1100 **I**N CHAPTER 3, I presented the state of the art of *process comprehension* of
1101 an Industrial Control System (ICS) focusing later on the methodology
1102 proposed by Ceccato et al. [9][Section 3.2], explaining what it consists of,
1103 its practical application on a testbed, and most importantly highlighting
1104 its limitations and critical issues (see Section 3.2.7).

1105 In this chapter I will present **my proposals to improve the methodol-**
1106 **ogy** presented in the previous chapter, overcoming (or at least trying to do
1107 so) the criticalities mentioned above by almost completely rewriting the
1108 original framework, enhancing its functionalities and inserting new ones
1109 where possible, while keeping its general structure and approach: the sys-
1110 tem analysis will in fact consist of the same four steps as in the original
1111 methodology (Data Pre-processing, Graph and Statistical Analysis, Busi-
1112 ness Process Mining and Invariants Inference), but each of them will be
1113 deeply revised in order to provide a richer, clearer and more complete
1114 process comprehension of the industrial system to be analyzed and its be-
1115 havior.

1116 As it may have already been noted, my proposals do not involve im-
1117 proving the data gathering phase: this is due simply to the fact that the
1118 novel framework will not be tested on the same case study used by Cec-

1119 cato et, al. (Section 3.2.1), but on a different case study, the ITrust SWaT
1120 system [34], of which (some) datasets containing the execution trace of
1121 the physical system and the network traffic scan are already provided by
1122 iTrust itself. For more details about this case study, see Chapter 5.

1123 4.1 The novel Framework

1124 The implementation of the novel framework for ICSs analysis starts
1125 from several assumptions:

- 1126 1. it must be implemented in a **single programming language**
- 1127 2. it must be **independent of the system** to be analyzed
- 1128 3. It must provide greater **flexibility and ease of use** for the user at
1129 every stage

1130 In the following, these three points will be discussed in more detail.

1131 **Single Programming Language** The original tool was implemented us-
1132 ing various programming languages in each of the different phases:
1133 from Python up to Java, passing through Bash scripting.
1134 In my opinion, this heterogeneity makes it more difficult and less
1135 intuitive for the user to operate on the tool: moreover, the use of
1136 multiple technologies makes it more difficult to maintain the code
1137 and add new features, particularly if only a single person is manag-
1138 ing the code (he/she might be proficient in one language, but little
1139 of the others).

1140 For these reasons, I decided to use a single programming language,
1141 to ensure homogeneity to the framework and ease of use and main-
1142 tenance of the code for anyone who wants to manage it in the future:
1143 I chose to use Python, because of its simplicity and easy readability
1144 combined with its versatility and powerfulness: moreover, Python
1145 can count on a massive number of available libraries and packages
1146 that meet all kinds of needs.

System Independence One of the biggest limitations of Ceccato et al.'s tool that I highlighted in Section 3.2.7 is the fact that it is **highly dependent on the testbed used**: that is, it is *not* possible to configure any of the tool's parameters to analyze different industrial systems. To overcome this issue and make my framework independent of the system to be analyzed, also eliminating all references to hardcoded variables and values present in the previous tool, I decided to use a **general configuration file**, named *config.ini*, in which the user can, at will, customize all the parameters necessary to perform the analysis of the targeted system.

Flexibility and Ease on Use The lack of flexibility and ease of use in a tool can be a significant disadvantage, limiting its effectiveness and making it challenging for the user to get the desired outcomes. The original tool suffered from these limitations, with users having to run scripts from the command line, with little to no options or parameters available to customize the analysis. As a result, the tool was not user-friendly and lacked the flexibility to adapt to specific user needs.

To settle these issues, I enhanced the command-line interface in the novel framework by adding new options and parameters. These new features provide the user with greater flexibility, enabling to specify parameters and options that allow for more in-depth analysis and focused results analyzing data more effectively and efficiently. With these enhancements, the framework has become more user-friendly, reducing the learning curve and making it more accessible to a wider range of users.

This, in turn, makes the framework more valuable and useful, increasing its adoption and effectiveness across a range of industries and applications.

Moreover, with new options and parameters users no longer have to rely solely on the command line interface, which can be challenging and intimidating for those with limited technical expertise. Instead,

1179 users can now access a range of customizable options and paramete-
1180 ters, making the tool more intuitive and user-friendly.

1181 Overall, the enhancements made to the framework represent a significant
1182 step forward in making it more effective, efficient, and user-friendly.

1183 4.1.1 Framework Structure

1184 The structure of the novel framework mostly follows the structure of
1185 the original tool: it is divided into four main directories each representing
1186 the different phases of the analysis (data pre-processing, graphs and sta-
1187 tistical analysis, process mining, and invariant analysis), and containing
1188 the relevant Python scripts that perform the analysis, as well as subdirec-
1189 tories and any input/output files necessary for the proper behavior of the
1190 framework.

```
1191 .  
1192 |-- config.ini  
1193 |-- daikon  
1194 |   |-- Daikon_Invariants  
1195 |   |-- daikonAnalysis.py  
1196 |   |-- runDaikon.py  
1197 |-- network-analysis  
1198 |   |-- data  
1199 |   |-- export_pcap_data.py  
1200 |   |-- swat_csv_extractor.py  
1201 |-- pre-processing  
1202 |   |-- mergeDatasets.py  
1203 |   |-- system_info.py  
1204 |-- process-mining  
1205 |   |-- data  
1206 |   |-- process_mining.py  
1207 |-- statistical-graphs  
1208 |   |-- histPlots_Stats.py  
1209 |   |-- runChartSubPlots.py
```

Listing 4.1: Novel Framework structure and Python scripts

Ahead of these directories there is the most important part, that allows the framework to be independent of the industrial control system being analyzed: the *config.ini* file. Here the user can configure general parameters and options, such as paths to read from or write files to, or related to individual analysis phases.

The file is divided into sections, each covering a different aspect of the configuration: each section contains user-customizable parameters that will then be called within the Python scripts that constitute the framework. Sections of *config.ini* are:

- **[PATHS]**: defines general paths such as the project root directory and some source directories for datasets
- **[PREPROC]**: contains some parameters needed for the pre-processing phase
- **[DAIKON]**: defines parameters needed for invariant analysis with Daikon
- **[DATASET]**: defines settings and parameters used during the dataset enrichment stage and possibly in further phases
- **[MINING]**: contains parameters used during the process mining phase
- **[NETWORK]**: Contains specific settings for extracting the data obtained from the packet sniffing phase on the ICS network and converting it to CSV format. It also defines the network protocols that are to be analyzed

4.1.2 Python Libraries and External Tools

Since the framework has been entirely developed in Python, I have tried to make use of external tools as little as possible, with the idea of integrating all the various functionalities within the framework and making it independent of further software: the only external tool remaining from the old Ceccato et al. tool is Daikon, precisely because there is currently

no better alternative or Python packages that performs the same functionalities.

Instead, large use of Python libraries is made for handling functionality and input data: the fundamental libraries upon which the framework is based are:

- **Pandas**, also used in the previous tool for dataset management, but whose use here has been deepened and extended
- **NumPy**, often used together with Pandas to perform some operations to support it
- **Matplotlib**, for managing and plotting graphical analysis
- other scientific libraries such as **SciPy**, **StatsModel** [53] and **NetworkX** [54], for mathematical, statistical and analysis operations on the data
- **GraphViz**, for the creation of activity diagrams in the process mining phase

Having now seen the structure of the framework, in the next sections we will go into more detail describing my proposals and what I have done to improve the various stages of the analysis.

4.2 Analysis Phases

4.2.1 Phase 1: Data Pre-processing

Data Pre-processing phase is probably the most delicate and significant one: depending on how large the industrial system to be analyzed is, the data collected, and how it is enriched using the additional attributes, the subsequent system analysis will provide more or less accurate outcomes.

The previous tool has many limitations, especially at this stage: it is not possible to isolate a subsystem (either on a temporal basis or on the number of PLCs to be analyzed - the system is considered in its whole), and many of the additional attributes were actually added manually: moreover, for those automatically entered, there is no way to specify which register type to associate the additional attribute with.

All this, combined with the fact that in the tool code many references to attributes and registers are hardcoded, makes the analysis of the system much more difficult and the obtained results less accurate in terms of quantity and quality.

In the novel framework these problems have been overcome by introducing the possibility, starting from the datasets of individual PLCs obtained from data gathering process, to select a subsystem from the command line both on a temporal basis and of the PLCs to be considered; I have also redesigned the whole process of enrichment of the resulting dataset, eliminating the manual entry of additional attributes and giving the user the possibility to be able to decide which type of additional attribute to associate with a given register. In addition to this, at the end of the pre-processing operation, it is possible to perform a brief preliminary analysis of the obtained dataset in order to estimate which registers are connected to actuators, which to measurements, and which represent hardcoded relative setpoints or constants: this operation also makes it possible to be able to refine the enrichment step by setting the relevant parameters in the *config.ini* file

In the next sections we will look in more detail at what has been accomplished.

1289 4.2.1.1 Subsystem Selection

1290 In the previous tool, the datasets in CSV format referring to each single
1291 PLC are placed in a fixed directory (hardcoded in the script) from which
1292 the dedicated script later perform merge and enrichment of them all, re-
1293 sulting in a single dataset representing the entire process trace of the in-
1294 dustrial system as an output. As mentioned, the script makes no provision
1295 to choose the individual PLCs to be analyzed, nor to decide on a temporal
1296 range over which to perform the analysis: in fact, it may happen that dur-
1297 ing the period of scanning and data gathering there is a so-called *transient*,
1298 i.e., a general state in which the industrial system is still initializing before
1299 actually reaching full operation; or, more simply, there is the need to ana-
1300 lyze the process of only a specific part of the industrial system in a certain
1301 period of interest: whatever the motivation, the lack of elasticity and op-
1302 tions to provide to the user makes the analysis much more complex than
1303 it might be, affecting even the later phases, as the number of variables to
1304 be analyzed becomes enormously higher.

1305 In addition, it is not possible to specify an output CSV file where to save
1306 the resulting dataset: at each dataset creation and enrichment operation,
1307 therefore, the resultant file will be overwritten. This is very awkward
1308 when making comparisons between two different execution traces, for ex-
1309 ample, unless the files are renamed manually.

1310 Let's see how all these issues were solved in the novel framework I
1311 developed: first of all, in the general *config.ini* file there are some general
1312 default settings about paths, and among them the one concerning the di-
1313 rectory where to place the datasets of the individual PLCs to be processed.
1314 In addition to this option, there are other ones that define further aspects
1315 related to the operations performed in this phase. Listing 4.2 shows the
1316 settings in question:

```
1317 [PATHS]  
1318 root_dir = /home/marcuzzo/Univr/Tesi  
1319 project_dir = %(root_dir)s/PLC-RE  
1320 input_dataset_directory = %(root_dir)s/datasets_SWaT/2015
```

```

1321 net_csv_path = %(root_dir)s/datasets_SWaT/2015/Network_CSV
1322
1323 [DEFAULTS]
1324 dataset_file = PLC_SWaT_Dataset.csv # Default output
1325     ↪ dataset
1326 granularity = 10 # slope granularity
1327 number_of_rows = 20000 # Seconds to consider
1328 skip_rows = 100000 # Skip seconds from beginning

```

Listing 4.2: Paths and parameters for the Pre-processing phase in config.ini file

1329 Concurrently, the same options can be specified by the user via the
 1330 command line of the new Python script (named mergeDatasets.py and
 1331 contained in the directory pre-processing of the project) and will override
 1332 the default ones found in *config.ini*. These options are:

- 1333 • **-s or --skiprows:** seconds to jump from the beginning of the file. This
 1334 option is useful in case the system has an initial transient or to start
 1335 the analysis from a certain point in the dataset
- 1336 • **-n or --nrows:** reference temporal period in seconds (rows) for the
 1337 analysis from the beginning of the dataset or from the point specified
 1338 in the -s option or in the corresponding setting in *config.ini*.
 1339 This option makes a **selection** on the data of the dataset.
- 1340 • **-p or --plcs:** PLCs to be merged and enriched. The user can specify
 1341 the desired PLCs by indicating the CSV file names of the associated
 1342 datasets with no limitations on number.
 1343 This option makes a **projection** on the data of the dataset.
- 1344 • **-d or --directory:** performs the merge and enrichment of all CSV files
 1345 contained in the directory specified by user, overriding the default
 1346 setting in *config.ini*. It is in fact the old functionality of the previous
 1347 tool, maintained here to give the user more flexibility and convenience
 1348 in case he wants to perform the analysis on the whole system.
 1349 This is also the default behavior in case the -p option is not specified.

- **-o or --output:** specifies the name of the file in which the obtained dataset will be saved. It must necessarily be a file in CSV format.
- **-g or --granularity:** specifies a granularity that will be used to calculate the measurement slope during the dataset enrichment phase. We will discuss this later in Section 4.2.1.2.

4.2.1.2 Dataset Enrichment

After a step in which a function is applied to each PLC-related dataset that eliminates its registers that have not been used within the system (this is especially true if the Modbus register scan has been performed, in which ranges of registers are scanned: it is assumed that unused registers have constant value zero), the **dataset enrichment operation** is performed.

This operation differs from the previous version not only in the fact that it is performed on each individual dataset and not on the resulting dataset, but also in the additional attributes: not only are they greater in number, but they are automatically calculated and inserted by the `mergeDatasets.py` script into the dataset and, most importantly, it is possible to decide through the parameters in the `config.ini` configuration file under the `[DATASET]` section to which registers these attributes should be assigned. In Listing 4.3 we can see the list of additional attributes and how they should be associated with the registers of the dataset:

```
[DATASET]
timestamp_col = Timestamp
max_prefix = max_
min_prefix = min_
max_min_cols_list = lit|ait|dpit
prev_cols_prefix = prev_
prev_cols_list = mv[0-9]{3}|p[0-9]{3}
trend_cols_prefix = trend_
trend_cols_list = lit
trend_period = 150
slope_cols_prefix = slope_
```



```
slope_cols_list = lit
```

Listing 4.3: config.ini parameters for dataset enriching

Following is a brief explanation of the parameters just seen:

timestap_col indicates the name of the column that contains the data timestamps. This parameter is used not only in this phase, but is also referred to in the Process Mining phase. In the previous work, this parameter was hardcoded and not configurable (and thus causing errors if the system being analyzed changed)

max_prefix, min_prefix, max_min_cols_list refer to any relative maximum or minimum values (*relative setpoints*) of one or more measures and that can be found and inserted as new columns within the dataset. The first two parameters indicate the prefix to be used in the column names affected by this additional attribute, while the third specifies of which type of registers we want to know the maximum and/or minimum value reached (several options can be specified using the logical operator | - or).

If, for example, we want to know the maximum value of the registers associated with the tanks, indicated in the iTrust SWaT system by the prefix LIT, we only need to specify the necessary parameter in the *config.ini* file, so `max_min_cols_list = lit`.

The result will be to have in the dataset thus enriched a new column named `max_P1_LIT101`.

prev_cols_prefix, prev_cols_list refer to the values at the previous step of the registers specified in `prev_cols_list`. It is possible to specify registers using *regex*, as in the example shown. It may be useful in some cases to have this value available to check, for example, when a change of state of a single given actuator occurs. The behavior of these parameters is the same as described in the point above.

slope_cols_prefix, slope_cols_list are related to the calculation of the slope of a specific register (usually a measure), that is, its trend. The

slope can be ascending (if its value is greater than zero), descending (if less than zero) or stable (if approximately equal to zero). I will discuss the slope calculation in more detail in the next paragraph, as it is related to the attributes `trend_cols_prefix`, `trend_cols_list` and `trend_period`

Initially, the parameters for registers to be associated with each additional attribute may be left blank, assuming that we do not know the system at all and therefore do not know which registers may be actuators, which measures, and which further. This information can be obtained from the **brief analysis** following the datasets merging operation: this analysis, performed at the user's choice, indicates which may be likely sensors, which actuators, and further information of various kinds: these indications allow the user to be able to set the desired values in `config.ini` file and hence refine the enrichment process by re-launching the `mergeDatasets.py` script again.

Slope Calculation The *slope* is an attribute that indicates the trend of the measurement we are considering and is useful, in our context, during the inference and invariant analysis phase in order to derive information about this trend given specific conditions: this trend can be, in general, increasing (slope > 0), decreasing (slope < 0) or stable (slope = 0). Normally, the slope is calculated through a simple mathematical formula: given an interval a, b relative to the measurement l , the slope is given by the difference of these two values divided by the amount of time t that the measurement takes to reach b from a :

$$slope = \frac{l(b) - l(a)}{t(b) - t(a)}$$

In the novel framework as in the old tool, this time interval (also called **granularity**) can be either long or short, depending on the accuracy desired on the slope: the lower the granularity, the more the slope will reflect the actual measurement trend; the higher the granularity, on the other hand, the more the slope data will be flattened. Each time interval into


which the measure is divided corresponds to a slope, the set of which inserted as an additional attribute in the dataset will later be used to define the trend of the measure itself in specific situations.

Calculating the slope directly from the raw measurement data may be an acceptable solution for those systems whose measurements are not significantly affected by **perturbations** (such as the oscillations of the liquid inside a tank during the filling and emptying phases, leading to fluctuating level readings): in this case, granularity can be kept low and thus obtain a very accurate overall trend calculation close to the actual measurement trend. This is the case, for example, with the tanks of the testbed used by Ceccato et al.


However, in the case where these perturbations significantly afflict the detections on the measurement, the slope calculation on the individual time intervals of the measurement may lead to an erroneous result in trend definition, regardless of the granularity used.

Figure 4.1 demonstrates this assertion: the measurement, in blue, refers to the LIT101 tank of the iTrust SWaT system; in red, the slope calculation related to the measurement with three different granularities: 30 (Figure 4.1a), 60 (Figure 4.1b) and 120 seconds (Figure 4.1c). It can be seen that in addition to the flattening of slope values as the granularity increases, in the time interval between seconds 1800 and 4200 the level of LIT101 has a generally increasing trend, but the slope values vary from positive to negative: the result is that in the invariant analysis the general increasing trend will not be detected thus losing the information.


The possibility of a having (strongly) perturbed data was not considered in the previous tool, so I was faced with this issue to solve.



(a)



(b)



(c)

Figure 4.1: Slope comparison with granularity 30 (a), 60 (b) and 120 seconds (c)

The solution to this problem is trying to remove as much "noise" as possible from the data, in order to get a more linear trend in the curve representing the measurement and consequently be able to calculate slopes more accurately.

There are several ways to smooth out the noise, but two of them seemed to me to be the most suitable and that I considered and evaluated: using **polynomial regression**, thus creating a filter on the noise, or a **seasonal decomposition**, and more specifically the part concerning the **trending**.

With regard to polynomial regression, I evaluated the *Savitzky-Golay* algorithm [55], and with regard to seasonal decomposition, I evaluated *Seasonal-Trend decomposition using LOESS* (STL) [56].

For reasons of the length of this paper I cannot describe these two solutions in detail (for this, I refer to the bibliographical notes): Figure 4.2, however, shows a quick graphical comparison of them compared with the original data. The solution chosen is the STL decomposition, which is more effective in attenuating noise than the Savitzky-Golay filter although at the cost of more delay (still present in all algorithms of this kind) in some parts.




Figure 4.2: Savitzky-Golay filter (blue line) and STL decomposition (green) comparison

1482 The application of the STL decomposition results in a significant im-
 1483 provement in slope calculation even when using low granularity: in Fig-
 1484 ure 4.3 it can be seen that, with the same granularity used for the example
 1485 in Figure 4.1a, the slope values, although with unavoidable fluctuations,
 1486 remain within the same trend, corresponding to the trend of the data curve
 1487 net of the introduced lag caused by the periodicity set for the decomposi-
 1488 tion.




Figure 4.3: Slope after the application of the STL decomposition

1489 This periodicity, which indicates the sampling time window for de-
 1490 composition and thus the level of noise smoothing, can be set in the *con-*
 1491 *fig.ini* configuration file in the *trend_period* directive.

1492 The slope calculation will then be performed on the data from the addi-
 1493 tional measurement trend attributes specified in the *trend_cols_list* di-
 1494 rective in the configuration file, and no longer on the original unfiltered
 1495 data.

1496 Finally, to enable Daikon to correctly interpret the slope data, the deci-
 1497 mal values corresponding to each calculated slope are converted into three
 1498 numerical values -1, 0, and 1, which correspond to the decreasing (if the
 1499 slope is less than zero), stable (if it is equal to zero), and increasing (if it is

greater than zero) trends, respectively. In Figure 4.4, the new slope can be seen, along with the curve obtained from the STL decomposition:




Figure 4.4: The new slope representation (green line) and the smoothed measurement data obtained with the STL decomposition (red)

4.2.1.3 Datasets Merging

In this step the individual datasets are merged, obtaining two new distinct datasets: one without the enrichment and that will be use in the process mining phase, and the other, containing the additional attributes but with the timestamp column removed, intended for inference and invariant analysis.

The first of the two datasets obtained will be saved in CSV format by default in the `$(project-dir)/process-mining/data` directory while the second one in the `$(project-dir)/daikon/Daikon_Invariants` directory (both paths are however configurable via `config.ini`).

The dataset name is specified in `config.ini` or via the `-o` command-line option: in the case of the dataset intended for process mining, the script will automatically add a `_TS` suffix to the filename, indicating the fact that the dataset includes the timestamp.

1516 The opportunity for the user to specify a different filename for the output
1517 each time allows the user to save the execution trace of the selected sub-
1518 system without overwriting the previous ones and thus to use all of them
1519 in the subsequent phases of the analysis.

1520 4.2.1.4 Brief Analysis of the Obtained Subsystem

1521 At the end of the datasets merging and saving step, the user is asked
1522 whether to perform a brief optional analysis of the final resulting dataset
1523 to extract preliminary data, with the purpose of obtaining basic informa-
1524 tion about the (sub)system and possibly refining the enrichment.
1525 If the user responds affirmatively to the request, `mergeDatasets.py` launches
1526 within it a further Python script located in the same
1527 `$(project-dir)/pre-processing` directory, called `system_info.py`. This
1528 script, using a combination of Daikon and Pandas analysis, performs a
1529 quick analysis of the dataset contents trying to **recognize**, however roughly,
1530 **the register types**, with possible maximum and minimum values and hard-
1531 coded setpoints. In addition, using the additional attribute `prev_`, it is ca-
1532 pable of deriving measurement values in correspondence with state changes
1533 of individual actuators.
1534 Listing 4.4 shows an example of this brief analysis elated to PLC1 of the
1535 iTrust SWaT system (for brevity, only one measurement is reported in the
1536 analysis of actuator state changes):

```
1537 Do you want to perform a brief analysis of the dataset? [y
1538 ↪ /n]: y
1539
1540 Actuators:
1541 P1_MV101 [0.0, 1.0, 2.0]
1542 P1_P101 [1.0, 2.0]
1543
1544 Sensors:
1545 P1_FIT101 {'max_lvl': 2.7, 'min_lvl': 0.0}
1546 P1_LIT101 {'max_lvl': 815.1, 'min_lvl': 489.6}
1547
1548 Hardcoded setpoints or spare actuators:
```



```

1549   P1_P102 [1.0]
1550
1551   Actuator state changes:
1552           P1_LIT101  P1_MV101  prev_P1_MV101
1553   669      800.7170      0          2
1554   1850     499.0203      0          1
1555   4876     800.5992      0          2
1556   6052     498.9026      0          1
1557   9071     800.7170      0          2
1558   10260    499.1381      0          1
1559   13268    801.3058      0          2
1560   14435    498.4315      0          1
1561   17423    801.4628      0          2
1562   18603    498.1567      0          1
1563
1564   P1_LIT101  P1_MV101  prev_P1_MV101
1565   677      805.0741      1          0
1566   4885     805.7414      1          0
1567   9079     805.7806      1          0
1568   13276    805.1133      1          0
1569   17432    804.4068      1          0
1570
1571   P1_LIT101  P1_MV101  prev_P1_MV101
1572   1858     495.4483      2          0
1573   6060     497.9998      2          0
1574   10269    495.9586      2          0
1575   14443    495.8016      2          0
1576   18611    494.5847      2          0
1577
1578           P1_LIT101  P1_P101  prev_P1_P101
1579   118      536.0356      1          2
1580   4322     533.3272      1          2
1581   8537     542.1591      1          2
1582   12721    534.8581      1          2
1583   16883    540.5890      1          2
1584
1585   P1_LIT101  P1_P101  prev_P1_P101
1586   1190     813.0031      2          1
1587   5395     813.0031      2          1

```

```

1588      9597      811.8256      2      1
1589      13776     812.7283      2      1
1590      17938     813.3171      2      1
1591
1592      Actuator state durations:
1593      P1_MV101 == 0.0
1594      9  9  10  9  9  10  9  9  10  9
1595
1596      P1_MV101 == 1.0
1597      1174  1168  1182  1160  1172
1598
1599      P1_MV101 == 2.0
1600      669  3019  3012  3000  2981
1601
1602      P1_P101 == 1.0
1603      1073  1074  1061  1056  1056
1604
1605      P1_P101 == 2.0
1606      118  3133  3143  3125  3108

```

Listing 4.4: Example of preliminar system analysis

From these results we can see that:

- the probable actuators are P1_MV101, which assumes three states identified by the values 0, 1 and 2, and P1_P101, which instead assumes two states identified by the values 1 and 2
- there are two probable measures: P1_FIT101 whose values range from 2.7 to 0.0, and P1_LIT101 whose values range from 815.1 to 489.6. One conjecture could already be made about the topology of the system: P1_LIT101 represents a tank
- apparently there are no related *hardcoded setpoints*, but a probable spare actuator, P1_P102, whose value is always 1. From this data, another conjecture can be made: the value 1 is the close state for that particular type of actuators, while 2 represents the open state
- from the analysis of state changes, in summary, we derive some *relative setpoints*: for example, we know that P1_P101 changes state from

value 1 to value 2 when the level of P1_LIT101 is about 813, while it changes from value 2 to 1 when the level of P1_LIT101 is about 535. We can deduce that P1_P101 is responsible for emptying the tank

- as the last information we have duration of actuator states for each cycle of the system: this information can be useful for making assumptions and conjectures about the behavior of an actuator in a specific state or, by observing the duration values of each cycle, highlighting anomalies in the system.
- In the specific case, the very short duration of P1_MV101 in state 0 is observable: since we are unable, at this preliminary stage, to make assumptions about this, we will try to understand the behavior of the system in this actuator state in the next stages of the analysis

The information obtained here can be used, as mentioned above, to refine the enrichment of the dataset by setting directives in the [DATASET] section of the *config.ini* file, should this be empty or only partially set, or to make the first conjectures about the system, as we have just seen.

The *system_info.py* file can also run in standalone mode if needed: it takes as command-line arguments the dataset to be analyzed, a list of actuators, and a list of sensors. For analysis related to state changes, the dataset must mandatorily be of the enriched type.

4.2.2 Phase 2: Graphs and Statistical Analysis

The new *graph analysis* arises from the need to give the user an overview of the (sub)system obtained in the previous pre-processing phase, identifying more easily the typology of the registers and grasping more effectively the relationships and the dynamics that may exist between the registers controlled by one or more PLCs, confirming the initial conjectures if the brief analysis described in the previous section has been performed, or making new ones thanks to the visual graph support.

1649 In the previous tool, as already pointed out in Section 3.2.7, it is pos-
 1650 sible to view the chart of one only register at a time: this certainly makes
 1651 it possible to identify, or at least to hypothesize, the type of that register,
 1652 but it makes it very complicated to be able to relate it to the other compo-
 1653 nents of the system and thus to derive conjectures about the behavior of
 1654 the latter, conjectures to be possibly confirmed in the later phases. Hence
 1655 the need to create a tool that was better than the previous one and that
 1656 provided more information in an easier way.

1657 The first thing I thought was that an approach similar to Figure 3.5,
 1658 with all the graphs contained within a single plot, might be the most suit-
 1659 able one: unfortunately, I quickly realized that this solution presented sev-
 1660 eral issues, which are highlighted in figure 4.5.





Figure 4.5: Plotting registers on the same y-axis


1661 In figure 4.5a it can be seen that the biggest issue is the use of the same
 1662 y-axis, related to the values of the individual registers, for all charts: if the
 1663 range between register values is wide, the risk of some charts resembling a
 1664 single flat line, or at any rate being in fact indistinguishable and very diffi-
 1665 cult to read immediately; furthermore, as shown in Figure 4.5b, if registers
 1666 have similar values to each other, the respective graphs may be confusing
 1667 to each other, making them more difficult to understand.

1668 The solution to this issue is simple and effective: the use of **subplots**.
 1669 Basically, each register corresponds to a subplot of the graph that shares
 1670 the time axis (the x-axis) with the other subplots, but keeps the y-axis of

the values of each register independent. This maintains the readability and comprehensibility of the charts, while simultaneously being able to immediately grasp the relationships between them. In addition, by sharing the time axis, it is possible to zoom in on a particular area of one of the charts and automatically the other ones will be zoomed in as well, thus not losing any information and no connection between registers. Figure 4.6 illustrates more clearly what has just been explained: the charts refer to the PLC1 registers of the iTrust SWaT system.



(a) Example of plotting charts of a PLC registers using subplots



(b) Zooming on a particular zone of the charts

Figure 4.6: Example of the new graph analysis

1679 To demonstrate the behavior and effectiveness of the new graph anal-
1680 ysis process, we observe in particular, Figure 4.6b: we can already have
1681 some validation on the conjectures made in the brief analysis from the
1682 previous step:

- 1683 • the chart of P1_FIT101 seems to confirm that this register can be as-
1684 sociated with a **measurement**; moreover, we can see that the trend
1685 of this register is closely related to the periods in which P1_LIT101
1686 has an increasing trend and to the evolution of P1_MV101. Its values,
1687 between about 0 and 2.5, are too narrow to say that this register rep-
1688 represents the tank, and the general trend also leads in that direction:
1689 we can therefore assume that it is a **pressure or flow sensor**
- 1690 • P1_LIT101, because of the above, would therefore seem to **represent**
1691 **a tank**, as already assumed in the brief analysis
- 1692 • P1_MV101 assumes value 2 at the beginning of the ascending trend of
1693 P1_LIT101, while it has value 1 in correspondence with the *plateau*
1694 that is seen when the measurement is about 800 and when the trend
1695 of it is descending. Thus, we can have a confirmation of the ini-
1696 tial conjecture: 1 represents the OFF state of the sensor, 2 the ON
1697 state, and P1_MV101 is the **actuator responsible for the rising level**
1698 of P1_LIT101
- 1699 • in the brief analysis we observed the short duration of P1_MV101 in
1700 state 0, but we were not able to speculate on the reason for this.
1701 Graph analysis shows that P1_MV101 switches and stays in the 0
1702 state acting as a kind of "transient" between states 1 and 2. It can
1703 be assumed that that is the period of time it actually takes for the
1704 actuator to change state
- 1705 • P1_P101 assumes value 2 at the beginning of the descending trend of
1706 P1_LIT101, after the *plateau*, and returns to value 1 at the change of
1707 slope of the (likely) tank increase: taken by itself this fact is not very
1708 clear, so it needs to be compared with P1_LIT101 and P1_MV101 to un-
1709 derstand its behavior: it can be seen that when P1_101 and P1_MV101

are in state 1 the water level remains stable, whereas, when P1_P101 changes to state 2 the level of the measurement drops rapidly; when P1_MV101 then also returns to state 2, the level of the measurement slowly starts rising again and then has a sudden rise at the change of state of P1_P101 from 2 to 1. We can therefore infer that P1_P101 is the **actuator responsible for emptying** the (presumed) tank: state 1 represents the OFF state, while state 2 represents the ON state

- P1_P102 seems to play no role in the system, since its state remains at 1 all the time. It seems improbable that it could be a relative setpoint, so it can be assumed, according to the brief analysis, that it may be a **spare actuator**

As you have probably already noticed, the vast majority of the graphs shown in the previous sections and chapters were generated precisely using the new graph analysis script.

The script that performs the new graph analysis is `runChartsSubPlots.py`, located in the `$(project-dir)/statistical-graphs` directory, which is written in Python and makes use, like the previous one, of the *matplotlib* libraries to generate the graph plots.

The script accepts the following command-line parameters:

- **-f** or **--filename**: specifies the dataset, in CSV format, from which to read the data. The dataset must be located within the directory containing the enriched datasets for the invariant analysis phase. If subsequent parameters are not specified, the script will show all registers in the dataset, except for additional attributes
- **-r** or **--registers**: specifies one or more specific registers to be displayed
- **-a** or **--addregisters**: adds one or more registers to the default visualization. This option is useful in case an additional attribute such as slope is to be analyzed

- **-e** or **--excluderegisters**: excludes one or more specific registers from the default visualization. This option is useful to avoid displaying hardcoded setpoints or spare registers

This script, like the previous ones, is designed to provide the maximum flexibility and ease of use for the user, combined with greater power and effectiveness in deriving useful information about the analyzed system.

Statistical Analysis A quick mention of the statistical analysis part: after careful evaluation, I decided **not to include this aspect** of the previous tool within the framework, because I saw no real practical use for it. The actual statistical information can be easily integrated into the brief analysis immediately following the pre-processing phase, while the histogram has relative utility and, in my opinion, is outdated by the new graph analysis I implemented.

However, the Python `histPlots_Stats.py` script remains in the directory, essentially unchanged from the one developed by Ceccato et al., in case future use is needed.

4.2.3 Phase 3: Invariant Inference and Analysis

The phase of invariant inference and analysis has undergone a redesign and improvement to offer the user a more comprehensive and effortless approach to identify invariants. This has been achieved through the application of new criteria to analyze and reorganize the Daikon analysis results. The outcome of this is a more compact presentation of information that highlights the possible relationships among invariants.

The new design not only enables the identification of undiscovered aspects of the system behavior but also confirms the hypotheses made during the earlier stages of analysis. This step is now semi-automated, unlike before, and allows for the analysis of invariants on individual actuator states and their combinations.

4.2.3.1 Revised Daikon Output

To streamline the process of identifying invariants quickly and efficiently, it is necessary to revise the output generated by standard Daikon analysis. The goal is to create a more compact and readable format for the output.

The current Daikon results consist of three sections:

1. the first section containing generic invariants, i.e., valid regardless of whether a condition is specified for the analysis
2. the second section containing invariants obtained by specifying a condition for the analysis in the *.spinfo* file
3. a third section containing the invariants that are obtained from the negations of the condition possibly specified in the *.spinfo* file

In each section only a single invariant per row is shown, without relating it in any way to the others: this makes it difficult to identify significant invariants and any invariant chains that might provide much more information about the behavior of the system than the single invariant.

A brief example of the structure and format of this output related to PLC1 of the iTrust SWaT system is shown in Listing 4.5, where a condition was specified on the measurement P1_LIT101 and on actuator P1_MV101:

```

aprogram.point:::POINT
P1_P102 == prev_P1_P102
P1_FIT101 >= 0.0
P1_MV101 one of { 0.0, 1.0, 2.0 }
P1_P101 one of { 1.0, 2.0 }
P1_P102 == 1.0
max_P1_LIT101 == 816.0
min_P1_LIT101 == 489.0
slope_P1_LIT101 one of { -1.0, 0.0, 1.0 }
[...]
P1_LIT101 > P1_MV101
P1_LIT101 > P1_P101

```

```

1799     P1_LIT101 > P1_P102
1800     P1_LIT101 < max_P1_LIT101
1801     P1_LIT101 > min_P1_LIT101
1802     [...]
1803     P1_MV101 < min_P1_LIT101
1804     P1_MV101 < trend_P1_LIT101
1805     P1_P101 >= P1_P102
1806     P1_P101 < max_P1_LIT101
1807     [...]
1808     =====
1809     aprogram.point:::POINT;condition="P1_MV101 == 2.0 &&
1810     ↪ P1_LIT101 < max_P1_LIT101 - 16 && P1_LIT101 >
1811     ↪ min_P1_LIT101 + 15"
1812     P1_MV101 == prev_P1_MV101
1813     P1_P102 == slope_P1_LIT101
1814     P1_MV101 == 2.0
1815     P1_FIT101 > P1_MV101
1816     P1_FIT101 > P1_P101
1817     P1_FIT101 > P1_P102
1818     P1_FIT101 > prev_P1_P101
1819     P1_MV101 >= P1_P101
1820     P1_MV101 >= prev_P1_P101
1821     P1_P101 <= prev_P1_P101
1822     =====
1823     aprogram.point:::POINT;condition="not(P1_MV101 == 2.0 &&
1824     ↪ P1_LIT101 < max_P1_LIT101 - 16 && P1_LIT101 >
1825     ↪ min_P1_LIT101 + 15)"
1826     P1_P101 >= prev_P1_P101
1827     Exiting Daikon.


```

Listing 4.5: Standard Daikon output for PLC1 of the iTrust SWaT system


In the framework I am presenting, this output is reduced to two sections, the general one and the related to the user-specified condition, eliminating the one related to the negated condition because it is not useful in this context and a source of potential misinterpretation; moreover, the relationships between invariants will be highlighted instead by using **transitive closures**: transitive closure of a relation R is another relation, typically denoted R^+ that adds to R all those elements that, while not necessarily

related directly to each other, can be reached by a *chain* of elements related to each other. In other words, the transitive closure of R is the smallest (in set theory sense) transitive relation R such that $R \subset R^+$ [57].


To implement the transitive closure of the invariants generated by Daikon, I first divided the invariants in each section by type according to their equivalence relation, not considering all those invariants that refer to additional attributes except for the slope, and for each of them I generated a **graph** using the NetworkX libraries, connecting with an arc the registers (represented by nodes) that have a common endpoint in the invariant through the transitive property. To reconstruct the individual invariant chains, then, I used a simple *Depth-first Search* (DFS) on each of the graphs, thus obtaining the desired outcome. Figure 4.7 shows some examples of these graphs:




(a) Transitive closure graph for "="



(b) Transitive closure graph for ">"



(c) Transitive closure graph for "<"



(d) Transitive closure graph for ">" cond.

Figure 4.7: Example of transitive closure graphs for invariants in PLC1 of iTrust SWaT system

At the end of this process, still applied to PLC1 of the iTrust SWaT system and with the same analysis condition, we get the following complete output:

```

1851 1  =====
1852 2  Generic
1853 3  =====
1854 4  P1_MV101 one of { 0.0, 1.0, 2.0 }
1855 5  P1_P101 one of { 1.0, 2.0 }
1856 6  slope_P1_LIT101 one of { -1.0, 0.0, 1.0 }
1857 7  P1_FIT101 != P1_P101, P1_P102
1858 8  P1_P102 == 1.0
1859 9  max_P1_LIT101 == 816.0
1860 10 min_P1_LIT101 == 489.0
1861 11 P1_LIT101 > P1_MV101
1862 12 P1_LIT101 > P1_P101
1863 13 P1_LIT101 > P1_P102
1864 14 P1_LIT101 > min_P1_LIT101 > slope_P1_LIT101
1865 15 P1_FIT101 >= 0.0
1866 16 P1_P101 >= P1_P102 >= slope_P1_LIT101
1867 17
1868 18  =====
1869 19  P1_MV101 == 2.0 && P1_LIT101 < max_P1_LIT101 - 16 &&
1870 20  ↪ P1_LIT101 > min_P1_LIT101 + 15
1871 21  =====
1872 22  slope_P1_LIT101 == P1_P102
1873 23  P1_MV101 == 2.0
1874 24  P1_FIT101 > P1_MV101
1875 25  P1_FIT101 > P1_P101
1876 26  P1_FIT101 > P1_P102
1877 27  P1_MV101 >= P1_P101

```

Listing 4.6: Revised Daikon output with transitive closures for PLC1 of the iTrust SWaT system

Transitive closures can be appreciated in lines 7, 14 and 16 of Listing 4.6. In general, the output has been reduced in the number of effective rows (19 versus the 61 in the output of Listing 4.5, making it certainly better to read and identify significant invariants) and the invariant chains make it more immediate to grasp the relationships between registers.

4.2.3.2 Types of Analysis

Compared to Ceccato et al.'s solution, in which individual analyses are performed manually, my proposal is to implement two types of semi-automated analysis: the first performs an analysis of all states for each individual register, while the second performs the analysis on the current system configuration based on the actual states of the actuators.

Both analyses refer to a specific measurement, selected by the user.

These two types of analysis will be handled by the Python script `daikonAnalysis.py`, contained in the default directory `$(project_dir)/daikon`. The script accepts the following command-line arguments:

- **-f** or **--filename**: specifies the enriched dataset, in CSV format, from which to read the data. The dataset must be located within the directory containing the enriched datasets specified in the `config.ini` file (by default `$(project_dir)/daikon/Daikon_Invariants`).
- **-s** or **--simpleanalysis**: performs the analysis on the states of individual actuators
- **-c** or **--customanalysis**: performs the analysis on combinations of actual states of the actuators
- **-u** or **--uppermargin**: defines a percentage margin on the maximum value of the measurement
- **-l** or **--lowermargin**: defines a percentage margin on the minimum value of the measurement

One or both types of analysis can be selected. The last two parameters set a condition on the value of the measurement that is meant to bypass the transient periods between the actuator state change and the actual trend change at the maximum and minimum values: this expedient is especially useful for the first type of analysis, allowing for generally more accurate data on measurement trends.

1911 **Analysis on single actuator states** Analysis on the states of individual
 1912 actuators is the simplest: after the user is prompted to input the measure-
 1913 ment, chosen from a list of likely available measurements, the script rec-
 1914 ognizes the likely actuators and the relative states of each, using the same
 1915 mixed Daikon/Pandas technique adopted in the brief analysis during the
 1916 pre-processing phase.
 1917 For each actuator and each state it assumes, a single Daikon analysis is
 1918 performed, eventually placing the condition on the maximum and mini-
 1919 mum level of the measurement.
 1920 The result of these analyses are saved in the form of text files in a directory
 1921 having the name corresponding to the analyzed actuator and contained in
 1922 the default parent directory $\$(project_dir)/daikon/Daikon_Invariants/results$:
 1923 each file generated by the analysis is identified by the name of the actuator,
 1924 the state and the condition, if any, on the measurement (see Figure 4.8).




Figure 4.8: Directory (a) and outcome files (b) for the single actuator states analysis

1925 An example of the outcome of this analysis is Listing 4.6, where the
 1926 actuator analyzed is P1_MV101 in state 2: from the generic invariants we
 1927 can observe that P1_P101 is a likely actuator (line 5), assuming binary val-
 1928 ues; from line 8, however, we note that P1_P102 is permanently equal to 1,
 1929 so it can be confirmed that it is a spare actuator rather than a setpoint;
 1930 furthermore, lines 9 and 10 show the maximum and minimum values
 1931 reached by P1_LIT101, which we have assumed to be the register con-
 1932 nected to the tank. The most interesting information, however, comes

from the condition-generated invariants: at line 21 we notice that the slope of P1_LIT101 is 1 (thus increasing) when P1_MV101 is set to 2: this confirms what we assumed in the previous steps, namely, that P1_MV101 represents the ON state of the actuator and is responsible for filling the tank P1_LIT101; moreover, at line 23 we see that P1_FIT101 is greater than P1_MV101: this means that when the actuator assumes the value 2 the sensor P1_FIT101 detects data (in contrast, if P1_MV101 is 1 P1_FIT101 is 0, detecting nothing), confirming the original hypothesis made that this may be a pressure or flow sensor.

Analysis of the Current System Configuration The analysis of the current system configuration based on the actual states of the actuators is more complex, but at the same time offers more interesting outcomes as it provides better evidence of actuator behavior in relation to the selected measurement.




Figure 4.9: Daikon outcome files for system configuration analysis. Each file represents a single system state

For this analysis, the script automatically recognizes the actual configurations of the actuators (e.g., P1_MV101 == 2, P1_P101 == 1) that represent a system state, and for each of these configurations Daikon analysis is automatically run, after prompting the user for the measurement and

1951 eventually the actuators whose configurations are to be studied (If the user
 1952 does not select any actuators, all those previously detected by the Python
 1953 script will be considered): the analysis outcomes are saved, still in the text
 1954 format, within a specific directory under the same parent directory of the
 1955 previous type of analysis and, as before, their name is characterized by the
 1956 rule used for the analysis (see Figure 4.9).

1957

1958 An example of the obtained outcomes can be seen in Listing 4.7:

```

1959 1  =====
1960 2  Generic
1961 3  =====
1962 4  P1_MV101 <= P1_P101 ==> P1_FIT101 >= 0.0
1963 5  P1_MV101 <= P1_P101 ==> P1_MV101 one of { 0.0, 1.0, 2.0
1964 6  ↪ }
1965 7  P1_MV101 <= P1_P101 ==> P1_P101 one of { 1.0, 2.0 }
1966 8  P1_MV101 <= P1_P101 ==> slope_P1_LIT101 one of { -1.0,
1967 9  ↪ 0.0, 1.0 }
1968 10 P1_MV101 > P1_P101 ==> P1_FIT101 > P1_MV101
1969 11 P1_MV101 > P1_P101 ==> P1_FIT101 > P1_P101
1970 12 P1_MV101 > P1_P101 ==> P1_FIT101 > P1_P102
1971 13 P1_MV101 > P1_P101 ==> P1_FIT101 > slope_P1_LIT101
1972 14 P1_MV101 > P1_P101 ==> P1_MV101 == 2.0
1973 15 P1_MV101 > P1_P101 ==> P1_MV101 > P1_P102
1974 16 P1_MV101 > P1_P101 ==> P1_MV101 > slope_P1_LIT101
1975 17 P1_MV101 > P1_P101 ==> P1_P101 == 1.0
1976 18 P1_MV101 > P1_P101 ==> P1_P101 == P1_P102
1977 19 P1_MV101 > P1_P101 ==> P1_P101 == slope_P1_LIT101
1978 20 P1_MV101 > P1_P101 ==> slope_P1_LIT101 == 1.0
1979 21 [...]
1980 22
1981 23 =====
1982 24 P1_MV101 == 2 && P1_P101 == 1 && P1_LIT101 < max_P1_LIT101
1983 25 ↪ && P1_LIT101 > min_P1_LIT101
1984 26 =====
1985 27 slope_P1_LIT101 == P1_P102 == P1_P101 == 1.0
1986 28 P1_MV101 == 2.0
1987 29 P1_FIT101 > P1_MV101

```



```

1988 27 P1_FIT101 > P1_P101

```

Listing 4.7: *Daikon outcomes for the system configuration $P1_MV101 == 2$, $P1_P101 == 1$ on $P1_LIT101$*

Compared to Listing 4.6, this time we can observe in the general invariants section the presence of implications, which were previously absent (the remaining generic invariants have been omitted for reasons of space). Such implications can provide very useful information, as in this case: for example, the invariant on line 18 tells us that if the value of $P1_MV101$ is greater than that of $P1_P101$ then the value of the slope is 1, that is, the tank level is increasing. Consequently, we have further confirmation that the state $P1_MV101 == 2$ is the ON state for that actuator and that it is the actuator responsible for filling the tank $P1_LIT101$. Comparing the other analysis outcomes, we will discover that $P1_P101$ is instead responsible for emptying the tank and that its ON and OFF states are 2 and 1, respectively.

In addition, the invariant at line 8 also indicates that when the actuator $P1_MV101$ is in state 2 and $P1_P101$ is in state 1 then the value of $P1_FIT101$ is greater than 2: consequently, it follows that the corresponding sensor is measuring something relative to the tank $P1_LIT101$.

The above is confirmed by the invariants related to the analysis condition: at line 24 we have that the slope is indeed equal to 1 (thus increasing) and that $P1_FIT101$, at line 26, takes values greater than 2 when $P1_MV101$ is equal to 2.

Refining the Analysis In some cases, it may happen that the outcomes provided by the semi-automated analyses are not satisfactory to the user (e.g., the value of the slope may not emerge clearly) or simply the user wants to investigate a particular aspect of the system in more detail by trying to discover additional invariants that did not emerge previously: in this case, it is possible to run a new and more specific invariant analysis using the Python script `runDaikon.py`, which allows for more punctual analyses of the system.

2017 The script, also contained in the default directory `$(project_dir)/daikon`,
2018 accepts three command-line parameters:

- 2019 • **-f** or **--filename**: specifies the enriched dataset, in CSV format, from
2020 which to read the data. Even in this case, the dataset must be located
2021 within the directory containing the enriched datasets
- 2022 • **-c** or **--condition**: specifies the condition for the analysis, which will
2023 be automatically overwritten in the *.spinfo* file. It is possible to spec-
2024 ify more than one condition, but it is recommended to use the logical
2025 operator `&&`
- 2026 • **-r** or **--register**: specifies the directory where to save the text file with
2027 the outcomes

2028 Performing this single analysis will produce a single output file contain-
2029 ing the invariants discovered, file that will be saved in the directory spec-
2030 ified by the user via the `-r` command-line option (by default the file will be
2031 saved in the directory `$(project_dir)/daikon/Daikon_Invariants/results`)
2032 to be examined later by the user.

2033 In conclusion, the two types of analysis presented, along with the pos-
2034 sibility of allowing for more refined analysis at a later date and Daikon's
2035 redefined output, make this stage much more complete, clear, and power-
2036 ful than before

2037 4.2.4 Phase 4: Business Process Analysis

Chapter 5

Case study: the iTrust SWaT System

HAVING introduced the innovative framework and highlighted its potential in the preceding chapter, we now turn our attention to the case study where we will apply this framework. As previously mentioned in Chapter 4 and demonstrated through various examples in the same chapter, our focus will be on the **iTrust SWaT system** [34], developed by the iTrust – Center for Research in Cyber Security of University of Singapore for Technology and Design [29]. The acronym SWaT represents *Secure Water Treatment*.

The iTrust SWaT system is a testbed that replicates on a small scale a real water treatment plant arises to support research in the area of cyber security of industrial control systems and has been operational since March 2015: it is still being used by students at the University of Singapore for educational and training purposes and is available to organizations to train their operators on cyber physical incidents.

5.1 Architecture

In contrast to the virtualized testbed discussed in Section 3.2.1 by Cecato et al., the iTrust SWaT system is composed entirely of physical hardware components. It encompasses various elements, starting from field

2056 devices and extending to PLCs, HMI, SCADA workstations, and the SCADA
2057 server (also referred to as the *historian*). The historian is responsible for
2058 recording data from field devices for further analysis. In the upcoming
2059 sections, we will delve deeper into the architecture of the physical process
2060 and the communication network.

2061 5.1.1 Physical Process

2062 The physical process of the SWaT consists of six stages, denoted P1
2063 through P6. These stages are [58][59]:

- 2064 1. **taking in raw water:** feeds unfiltered water into the system
- 2065 2. **chemical dosing:** adds chemicals to water useful for initial pretreat-
2066 ment
- 2067 3. **Ultra Filtration (UF) system:** the water is filtered through a semi-
2068 permeable membrane (ultrafiltration membrane) using the liquid pres-
2069 sure, effectively capturing impurities and suspended solids, as well
2070 as removing bacteria, viruses, and other pathogens present in the
2071 water.
- 2072 4. **dechlorination:** removes residual chlorine from disinfected water
2073 using ultraviolet lamps
- 2074 5. **Reverse Osmosis (RO):** performs further filtration of the water
- 2075 6. **backwash process:** cleans the membranes in UF using the water pro-
2076 duced by RO

2077 Figure 5.1 shows a graphical representation of the architecture and the six
2078 stages of the SWaT system.

2079 The SWaT system incorporates an array of sensors that play a crucial
2080 role in monitoring the system's operations and ensuring their safe. These
2081 sensors are responsible for continuously collecting data and providing
2082 valuable insights into the functioning of the system. These sensors are:




Figure 5.1: SWaT architecture

- Level Indication Transmitter (measured in mm)
- Flow Indication Transmitter (m³/hr)
- Analyser Indicator Transmitter
 - o Conductivity (μS/cm)
 - o pH
 - o Oxidation Reduction Potential (mV)
- Differential Pressure Indicator Transmitter (kPa)
- Pressure Indicator Transmitter (kPa)

The sensors and actuators associated with each PLC are shown in Figure 5.2.

Sensors and actuators are mapped to tags by the communication protocol used (see 5.1.2): a tag can be addressed via string descriptor defined by the system designer (e.g. MV101, to indicate motorized valve number 1 at stage 1) or by referring directly to the analog/digital pins of the PLC I/O unit [59].

Raw Water	Pre-Treatment	Ultra-Filtration	De-Chlorination	Reverse Osmosis	RO Product
P-101 Stopped	P-201 Stopped	P-301 Stopped	P-401 Stopped	P-501 Stopped	P-601 Stopped
P-102 Stopped	P-202 Stopped	P-302 Stopped	P-402 Stopped	P-502 Stopped	P-602 Stopped
MV-101 Closed	P-203 Stopped	MV-301 Closed	P-403 Stopped	MV-501 Closed	P-603 Stopped
LIT-101 520 mm	P-204 Stopped	MV-302 Closed	P-404 Stopped	MV-502 Closed	LS-601 Normal
FIT-101 0.00 m³/h LL	P-205 Stopped	MV-303 Closed	UV-401 Stopped	MV-503 Closed	LS-602 HIGH
FIT-201 0.00 m³/h LL	P-206 Stopped	MV-304 Closed	LS-401 Normal	MV-504 Closed	LS-603 LOW
	P-207 Stopped	PSH-301 Normal	LIT-401 1008 mm H	PSL-501 Normal	FIT-601 0.00 m³/h LL
	P-208 Stopped	DP SH-301 Normal	FIT-401 0.00 m³/h LL	PSH-501 Normal	
	MV-201 Closed	LIT-301 1012 mm H	AIT-401 0.17 ppm	AIT-501 6.89	
	LS-201 Normal	FIT-301 0.00 m³/h	AIT-402 275.70 mV	AIT-502 204.20 mV	
	LS-202 Normal	DP IT-301 0.95 kPa LL		AIT-503 264.23 µS/cm H	
	LS-203 Normal			AIT-504 14.27 µS/cm H	
	AIT-201 142.18 µS/cm L			FIT-501 0.00 m³/h L	
	AIT-202 7.20 H			FIT-502 0.00 m³/h HH	
	AIT-203 293.59 mV L			FIT-503 0.00 m³/h HH	
				FIT-504 0.00 m³/h LL	
				PIT-501 2.64 kPa LL	
				PIT-502 0.00 kPa H	
				PIT-503 0.00 kPa	

Figure 5.2: Sensors and actuators associated with each PLC

5.1.2 Control and Communication Network

The SWaT system’s network architecture follows the principles of layering and zoning, which enable segmentation and control of traffic within the network.

Five layers are present starting from the highest to the lowest:

- Layer 3.5 – Demilitarized Zone (DMZ)
- Layer 3 – Operation Management (Historian)
- Layer 2 – Supervisory Control (Touch Panel, Engineering Workstation, HMI Control Clients)
- Layer 1 – Plant Control Network (PLCs) (Star Network)
- Layer 0 – Process (Actuator/Sensors and Input/output modules) (Ring Network)

PLCs at Layer 1 communicate with their respective sensors and actuators at Layer 0 through a conventional ring network topology based on Ethernet/IP, to ensure that the system can tolerate the loss of a single link

without any adverse impact on data or control functionality.

PLCs between the different process stages at Layer 1 communicate with each other through a star network topology using the CIP protocol on Ethernet/IP, previously discussed in Section 2.2.6.3.

Regarding zoning, the SWaT system is divided into three zones, each containing one or more layers. These zones are, in descending order of security level:

- **Plant Control Network, or Control System:** includes layers from 0 to 2
- **DMZ:** includes Layer 3.5
- **Plant Network:** includes Layer 3

Figure 5.3 provides a clearer visualization of the zoning and layer division within the network architecture of the SWaT system. This diagram highlights the distinct zones and their corresponding layers, offering a comprehensive overview of the system's network structure.




Figure 5.3: SWaT network architecture

A specific IP address is associated with each device: in Table 5.1 we report the addresses for the PLCs, historian, and Touch Panel in the *Plant Control Network* (PCN) zone:

IP Address	Device
192.168.1.10	PLC1
192.168.1.20	PLC2
192.168.1.30	PLC3
192.168.1.40	PLC4
192.168.1.50	PLC5
192.168.1.60	PLC6
192.168.1.100	Touch Panel (PCN)
192.168.1.200	Historian Server

Table 5.1: main IP addresses of the six PLCs and SCADA in SWaT

5.2 Datasets

To facilitate the study and testing of technologies related to cyber security in Industrial Control Systems and critical infrastructure, iTrust offers researchers worldwide the opportunity to access a range of datasets [60]. These datasets consist of a collection of data obtained from the SWaT system, encompassing information on both the physical processes and network communications. The data is organized into different years, and researchers can request access to these datasets for their analysis and experimentation purposes.

Physical Process Datasets The datasets containing information about the physical processes are provided in CSV format files. These files encompass data collected during different time intervals, which can vary from a few hours to entire days. The granularity of the data is typically at a one-second interval, although there may be some exceptions. The collected data primarily consists of timestamped sensor measurements and

actuator status values for each PLC, describing the physical properties of the testbed in operational mode.

Network Communications Datasets Network communications are typically available in the form of *Packet Capture* format (PCAP) files. These files contain captures of communication network traffic, allowing researchers to analyze and examine the network interactions. In some instances, CSV files are provided instead of PCAP files, featuring different characteristics for the collected data.

5.2.1 Our Case Study: the 2015 Dataset

The dataset selected as a case study to apply the framework discussed in the previous chapter is specifically the dataset from the year 2015 [61]. The main reason for this choice is the unique characteristics found in the physical process dataset that are not present in datasets from subsequent years.

Physical Process Data The data collection process lasted 11 consecutive days, 24 hours per day. During the first 7 days, the system operated normally without any recorded attacks. However, attacks were observed during the remaining 4 days. The collected data reflects the impact of these attacks, leading to the creation of two separate CSV files: one containing the recorded data of SWaT during the system's regular operations, and the other containing data recorded during the days of the attacks. To ensure accurate information about the system, the dataset pertaining to the normal operations, which spans seven days, was chosen for analysis.

Data collection occurs at a frequency of one data point per second, with the assumption that significant attacks cannot occur within a shorter time frame. Additionally, the firmware of the PLCs remains unchanged throughout the data collection period.

At the beginning of data gathering the tanks are empty and the system must be initialized in order to then reach full operation: it typically takes

2176 around five hours for all tanks to be fully filled and for the system to sta-
2177 bilize and reach the appropriate operational state (see Figure 5.4).

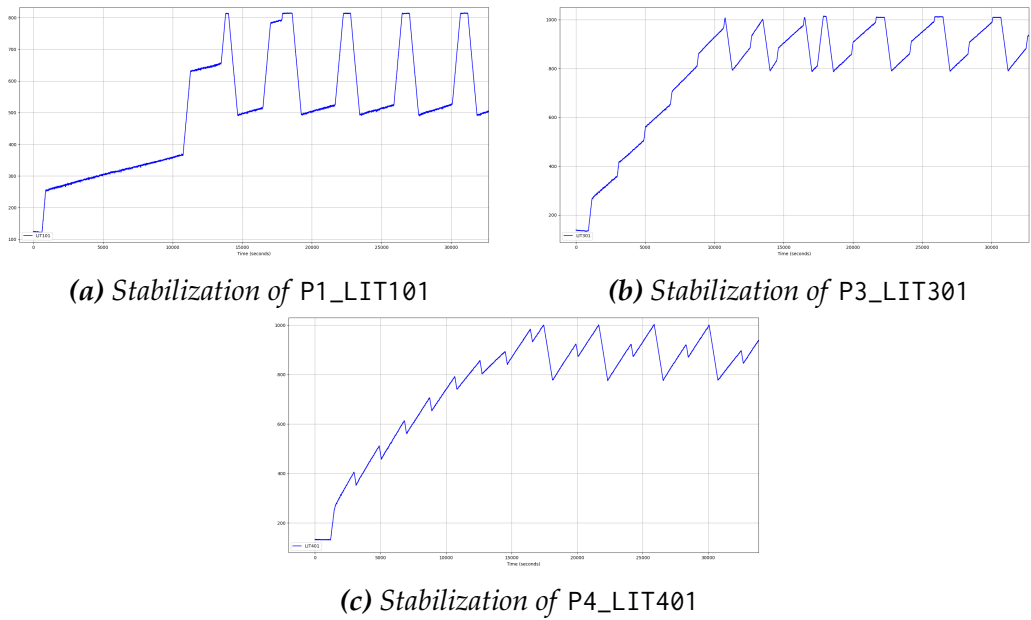


Figure 5.4: SWaT stabilization

2178 In total, the dataset consists of thousand and thousand of samples re-
2179 lated to 51 attributes.

2180 **Network Traffic** The network traffic was collected using an appliance
2181 from a well-known network hardware manufacturer and was made avail-
2182 able only in CSV format and not PCAP format. Table 5.2 shows some of
2183 the main data captured:

Category	Description
Date	Date of Log
Time	Time of Log
Origin	IP of server
Source IP	IP address of source
Destination IP	IP address of destination
Protocol	Network protocol

Application Name	Name of application
Modbus Function Code	Function Code
Modbus Function Description	Description of function
Modbus Transaction ID	Transaction ID
SCADA Tag	Sensor or actuator ID
Modbus Value	Value transmitted
Service / Destination Port	Port number of destination IP
Source Port	Port number of source IP

Table 5.2: SWaT network traffic data

Unfortunately, the data provided are partial as it only includes readings from sensors and does not include information on actuator status.

Do not mislead by the word *Modbus* in the captured data fields: this is most likely a feature of the sniffing appliance, which may have encapsulated EtherNet/IP-CIP protocol packets in Modbus frames, as described by T. A. Snide in [62].

To provide evidence that the protocol being used is CIP over EtherNet/IP, we can examine the "Source / Destination port" field in the captured data. The presence of TCP port 44818 as the destination port aligns with the TCP port used for transporting **explicit messages** in the EtherNet/IP protocol, as explained in Section 2.2.6.2. Additionally, the "Application name", "Modbus Function Code" and "Modbus Function Description" fields further support this conclusion. The "Application name" explicitly states "CIP_read_tag_service," while the "Modbus Function Code" field contains a decimal value of 76, which does not correspond to any known Modbus Function Calls. When converted to hexadecimal, this value is 4C, matching the CIP protocol read tag request code as indicated in the "Modbus Function Description" field.

In summary, the datasets for subsequent years were not selected due to certain limitations. In the case of the year 2017, the physical system

2205 data had a wide granularity, resulting in significant information loss. Ad-
2206 ditionally, in some cases, the datasets were considered "spurious" as there
2207 was no clear differentiation between data obtained during normal SWaT
2208 operations and data associated with attacks.

2209 Regarding network traffic data, there were instances where the data was
2210 either missing or fragmented into large PCAP files weighing several gi-
2211 gabytes (more than 6 GB each!), despite containing only a few minutes'
2212 worth of data. This made it impractical to manage such files given the
2213 available resources.

2214 Despite their large quantity, the CSV files associated with network traffic
2215 for the year 2015 are comparatively smaller in size, just slightly over 100
2216 MB each. These files cover a more extensive time period, which facilitates
2217 easier management and analysis. Prioritizing the physical process dataset
2218 as crucial, I have selected the year 2015 as a case study, even though the
2219 network data may be incomplete.

Chapter **6**

Our framework at work: reverse engineering of
the SWaT system

6.1 Pre-processing

6.2 Graph Analysis

6.2.1 Conjectures About the System

6.3 Invariants Analysis

6.3.1 Actuators Detection

6.3.2 Daikon Analysis and Results Comparing

6.4 Extra information on the Physics

6.5 Business Process Analysis

Chapter 7

Conclusions

7.1 Discussions

7.2 Guidelines

7.3 Future work

List of Figures

2.1	ICS architecture schema	6
2.2	PLC architecture	9
2.3	PLC communication schema	10
2.4	Comparison between ST language and Ladder Logic	11
2.5	Example of HMI for a water treatment plant	14
2.6	Modbus Request/Response schema	16
2.7	Modbus RTU frame and Modbus TCP frame	17
2.8	Example of packet sniffing on the Modbus protocol	19
2.9	OSI model for EtherNet/IP stack	20
3.1	Overview	27
3.2	The simplified SWaT system used for running Ceccato et al. methodology	28
3.3	Output graphs from graph analysis	33
3.4	An example of Disco generated activity diagram for PLC2	36
3.5	Execution traces of InputRegisters_IW0 on the three PLCs	38
3.6	Business process with states and Modbus commands for the three PLCs	40
3.7	Water physics compared: simulated physics in the Simulink model (a) and physics in a real system (iTrust SWaT) (b). Fluctuations in the tank level in (b), almost completely ab- sent in (a), can be appreciated.	46
3.8	Example of Daikon's output	50

4.1	Slope comparison with granularity 30 (a), 60 (b) and 120 seconds (c)	66
4.2	Savitzky-Golay filter (blue line) and STL decomposition (green) comparison	67
4.3	Slope after the application of the STL decomposition	68
4.4	The new slope representation (green line) and the smoothed measurement data obtained with the STL decomposition (red)	69
4.5	Plotting registers on the same y-axis	74
4.6	Example of the new graph analysis	75
4.7	Example of transitive closure graphs for invariants in PLC1 of iTrust SWaT system	81
4.8	Directory (a) and outcome files (b) for the single actuator states analysis	84
4.9	Daikon outcome files for system configuration analysis. Each file represents a single system state	85
5.1	SWaT architecture	91
5.2	Sensors and actuators associated with each PLC	92
5.3	SWaT network architecture	93
5.4	SWaT stabilization	96

List of Tables

2.1	differences between Information Tecnology (IT) and Industrial Control Systems (ICSs)	4
2.2	Modbus Function Codes list	18
5.1	main IP addresses of the six PLCs and SCADA in SWaT . . .	94
5.2	SWaT network traffic data	97

Listings

3.1	Example of registers capture	31
3.2	Example of raw network capture	32
3.3	Statistical data for PLC1_InputRegisters_IW0 register	33
3.4	Generic example of a .spinfo file for customizing rules in Daikon	34
4.1	Novel Framework structure and Python scripts	56
4.2	Paths and parameters for the Pre-processing phase in <i>con-</i> <i>fig.ini</i> file	60
4.3	<i>config.ini</i> parameters for dataset enriching	62
4.4	Example of preliminar system analysis	70
4.5	Standard Daikon output for PLC1 of the iTrust SWaT system	79
4.6	Revised Daikon output with transitive closures for PLC1 of the iTrust SWaT system	82
4.7	Daikon outcomes for the system configuration P1_MV101 == 2, P1_P101 == 1 on P1_LIT101	86

References

- [1] NIST. *ICS definition*. URL: https://csrc.nist.gov/glossary/term/industrial_control_system (visited on 2023-04-06).
- [2] Wikipedia. *Stuxnet*. URL: <https://en.wikipedia.org/wiki/Stuxnet> (visited on 2023-05-18).
- [3] Wikipedia. *Industroyer*. URL: <https://en.wikipedia.org/wiki/Industroyer> (visited on 2023-05-18).
- [4] F. Robles and N. Perlroth. “‘Dangerous Stuff’: Hackers Tried to Poison Water Supply of Florida Town”. In: *The New York Times* (2021-02-08). URL: <https://www.nytimes.com/2021/02/08/us/oldsmar-florida-water-supply-hack.html> (visited on 2023-05-18).
- [5] A. Akbarzadeh. “Dependency based risk analysis in Cyber-Physical Systems”. PhD thesis. Norwegian University of Science, Technology - Faculty of Information Technology, and Electrical Engineering, 2023. ISBN: 978-82-326-6986-8 (electronic version).
- [6] G. M. Makrakis, C. Kolias, G. Kambourakis, C. Rieger, and J. Benjamin. “Industrial and Critical Infrastructure Security: Technical Analysis of Real-Life Security Incidents”. In: *IEEE Access* (2021-12-06). DOI: <https://dx.doi.org/10.1109/ACCESS.2021.3133348>.
- [7] NIST. *PLC definition*. URL: https://csrc.nist.gov/glossary/term/programmable_logic_controller (visited on 2023-04-06).

- [8] W. Bolton. *Programmable Logic Controllers, 6th edition*. Newnes, 2015, pp. 7–9.
- [9] M. Ceccato, Y. Driouich, R. Lanotte, M. Lucchese, and M. Merro. “Towards Reverse Engineering of Industrial Physical Processes”. In: CPS4CIP@ESORICSAt 2022 (Copenhagen, Denmark, Sept. 30, 2022). 2022, 273–290. DOI: https://doi.org/10.1007/978-3-031-25460-4_15.
- [10] H. S. G. Pussewalage, P. S. Ranaweera, and V. Oleshchuk. “PLC security and critical infrastructure protection”. In: 2013 IEEE 8th International Conference on Industrial and Information Systems (Dec. 2013). 2013. DOI: <https://dx.doi.org/10.1109/ICIIInfS.2013.6731959>.
- [11] Wikipedia. *Remote Terminal Unit*. URL: https://en.wikipedia.org/wiki/Remote_terminal_unit (visited on 2023-04-08).
- [12] *What is SCADA?* URL: <https://www.inductiveautomation.com/resources/article/what-is-scada> (visited on 2023-04-05).
- [13] NIST. *HMI definition*. URL: https://csrc.nist.gov/glossary/term/human_machine_interface (visited on 2023-04-11).
- [14] Zscaler. *What Is the Purdue Model for ICS Security?* URL: <https://www.zscaler.com/resources/security-terms-glossary/what-is-purdue-model-ics-security> (visited on 2023-05-17).
- [15] Schneider Electric. *What is Modbus and how does it work?* URL: <https://www.se.com/us/en/faqs/FA168406/> (visited on 2023-04-13).
- [16] O. Morando. *Hacking: Modbus - Cyber security OT/ICS*. URL: <https://blog.omarmorando.com/hacking/ot-ics-hacking/hacking-modbus> (visited on 2023-04-15).
- [17] Modbus.org. “MODBUS/TCP Security”. In: pp. 7–8. URL: https://modbus.org/docs/MB-TCP-Security-v21_2018-07-24.pdf (visited on 2023-04-16).

- [18] ODVA Inc. "EtherNet/IP - CIP on Ethernet Technology". In: URL: https://www.odva.org/wp-content/uploads/2021/05/PUB00138R7_Tech-Series-EtherNetIP.pdf (visited on 2023-04-18).
- [19] *Odva, Inc.* URL: <https://www.odva.org> (visited on 2023-04-21).
- [20] *Introduction to EtherNet/IP Technology*. URL: https://scadahacker.com/library/Documents/ICS_Protocols/Intro%20to%20EthernetIP%20Technology.pdf (visited on 2023-04-19).
- [21] ODVA Inc. *Common Industrial Protocol*. URL: <https://www.odva.org/technology-standards/key-technologies/common-industrial-protocol-cip/> (visited on 2022-12-26).
- [22] Wikipedia. *Common Industrial Protocol*. URL: https://en.wikipedia.org/wiki/Common_Industrial_Protocol (visited on 2023-04-21).
- [23] *What is the Common Industrial Protocol (CIP)?* URL: <https://www.motioncontroltips.com/what-is-the-common-industrial-protocol-cip/> (visited on 2023-04-21).
- [24] *CIP (Common Industrial Protocol): CIP messages, device types, implementation and security in CIP*. URL: <https://resources.infosecinstitute.com/topic/cip/> (visited on 2023-04-21).
- [25] B. Green, M. Krotofil, and A. Abbasi. "On the Significance of Process Comprehension for Conducting Targeted ICS Attacks". In: Workshop on Cyber-Physical Systems Security and PrivaCy (Nov. 2017). ACM, 2017, pp. 57–67. DOI: <https://doi.org/10.1145/3140241.3140254>.
- [26] A. Keliris and M. Maniatakos. "ICSREF: A Framework for Automated Reverse Engineering of Industrial Control Systems Binaries". In: 26th Annual Network and Distributed System Security Symposium (NDSS) (Feb. 2019). 2019. DOI: <https://doi.org/10.48550/arXiv.1812.03478>.

- [27] Y. Yuan, X. Tang, W. Zhou, W. Pan, X. Li, H. T. Zhang, H. Ding, and J. Goncalves. "Data driven discovery of cyber physical systems". In: *Nature Communications* 10 (2019-10-25). URL: <https://www.nature.com/articles/s41467-019-12490-1> (visited on 2023-04-20).
- [28] C. Feng, V.R. Palleti, A. Mathur, and D. Chana. "A Sysematic Framework to Generate Invariants for Anomaly Detection in Industrial Control Systems". In: NDSS Symposium 2019 (San Diego, CA, USA, Feb. 24–27, 2019). 2019. DOI: <https://dx.doi.org/10.14722/ndss.2019.23265>.
- [29] Singapore University of Technology and Design. *iTrust - Center for Research in Cyber Security*. URL: <https://itrust.sutd.edu.sg/> (visited on 2023-05-15).
- [30] K. Pal, A. Adepu, and J. Goh. "Effectiveness of Association Rules Mining for Invariants Generation in Cyber-Physical Systems". In: 2017 IEEE 18th International Symposium on High Assurance Systems Engineering (HASE) (Jan. 2017). 2017. DOI: <https://doi.org/10.1109/HASE.2017.21>.
- [31] *Association rules mining*. URL: https://en.wikipedia.org/wiki/Association_rule_learning (visited on 2023-04-21).
- [32] K. Pal, A. Adepu, and J. Goh. "Cyber-Physical System Discovery: Reverse Engineering Physical Processes". In: 3rd ACM Workshop on Cyber-Physical System Security (Apr. 2017). 2017, pp. 3–14. DOI: <https://doi.org/10.1145/3055186.3055195>.
- [33] *Ceccato et al. reverse engineering prototype tool*. URL: <https://github.com/SeUniVr/PLC-RE> (visited on 2023-04-23).
- [34] iTrust Singapore University of Technology and Design. *Secure Water Treatment*. URL: <https://itrust.sutd.edu.sg/testbeds/secure-water-treatment-swat/> (visited on 2023-05-15).

- [35] R. Lanotte, M. Merro, and A. Munteanu. “Industrial Control Systems Security via Runtime Enforcement”. In: *ACM Transactions on Privacy and Security* 26.4 (2023-02), pp. 1–41. DOI: <https://doi.org/10.1145/3546579>.
- [36] M. Oliani. *AttackPLC - Project for Network Security teaching*. URL: <https://github.com/marcooliani/AttackPLC> (visited on 2023-04-23).
- [37] Ray - *Productionizing and scaling Python ML workloads simply*. URL: <https://www.ray.io/> (visited on 2023-04-25).
- [38] Tshark. URL: <https://tshark.dev/> (visited on 2023-04-25).
- [39] Wireshark. URL: <https://www.wireshark.org/> (visited on 2023-04-25).
- [40] pandas - *Python Data Analysis library*. URL: <https://pandas.pydata.org/> (visited on 2023-04-25).
- [41] NumPy - *The fundamental package for scientific computing with Python*. URL: <https://numpy.org/> (visited on 2023-04-25).
- [42] R project for statistical computing. URL: <https://www.r-project.org/> (visited on 2023-04-25).
- [43] SciPy - *Fundamental alghorithms for scientific computing with Python*. URL: <https://scipy.org/> (visited on 2023-04-25).
- [44] *The Daikon dynamic invariant detector*. URL: <http://plse.cs.washington.edu/daikon/> (visited on 2023-04-25).
- [45] *Enhancing Daikon output*. URL: <https://plse.cs.washington.edu/daikon/download/doc/daikon/Enhancing-Daikon-output.html#Splitter-info-file-format> (visited on 2023-04-25).
- [46] *Process mining*. URL: https://en.wikipedia.org/wiki/Process_mining (visited on 2023-04-26).
- [47] *Fluxicon Disco*. URL: <https://fluxicon.com/disco/> (visited on 2023-04-26).
- [48] *OpenPLC - Open source PLC software*. URL: <https://openplcproject.com/> (visited on 2023-04-23).

- [49] Docker. URL: <https://www.docker.com/> (visited on 2023-04-26).
- [50] MathWorks. *Simulink*. URL: <https://it.mathworks.com/products/simulink.html> (visited on 2023-04-26).
- [51] *ProM Tools - The Process Mining Framework*. URL: <https://promtools.org/> (visited on 2023-04-26).
- [52] Fraunhofer Institute for Applied Information Technology. *pm4py - Process Mining for Python*. URL: <https://promtools.org/> (visited on 2023-04-26).
- [53] statsmodels. *statsmodels - statistical models, hypothesis tests, and data exploration*. URL: <https://www.statsmodels.org/stable/index.html> (visited on 2023-05-05).
- [54] NetworkX. *NetworkX - Network Analysis in Python*. URL: <https://networkx.org/> (visited on 2023-05-05).
- [55] Wikipedia. *Savitzky–Golay filter*. URL: https://en.wikipedia.org/wiki/Savitzky%E2%80%93Golay_filter (visited on 2023-05-06).
- [56] Otext. *STL decomposition*. URL: <https://otexts.com/fpp2/stl.html> (visited on 2023-05-06).
- [57] Wikipedia. *Transitive closures*. URL: https://en.wikipedia.org/wiki/Transitive_closure (visited on 2023-05-11).
- [58] iTrust Singapore University of Technology and Design. *Secure Water Treatment (SWaT) Testbed*. URL: https://itrust.sutd.edu.sg/wp-content/uploads/2021/07/SWaT_technical_details-160720-v4.4.pdf (visited on 2023-05-15).
- [59] A. Mathur and N. O. Tippenhauer. “SWaT: a water treatment testbed for research and training on ICS security”. In: 2016 International Workshop on Cyber-physical Systems for Smart Water Networks (CySWater) (Apr. 2016). 2016. DOI: <http://dx.doi.org/10.1109/CySWater.2016.7469060>.

-
- [60] iTrust Singapore University of Technology and Design. *Datasets*. URL: https://itrust.sutd.edu.sg/itrust-labs_datasets/ (visited on 2023-05-15).
 - [61] S. Adepu, J. Goh, K.N. Junejo, and A. Mathur. "A Dataset to Support Research in the Design of Secure Water Treatment Systems". In: The 11th International Conference on Critical Information Infrastructures Security (France, Oct. 2016). 2016.
 - [62] T. A. Snide. *CIP-Modbus Integration*. 2008. URL: https://www.modbus.org/docs/CIP%20Modbus%20Integration%20Hanover%20Fair_0408.pdf (visited on 2023-05-17).